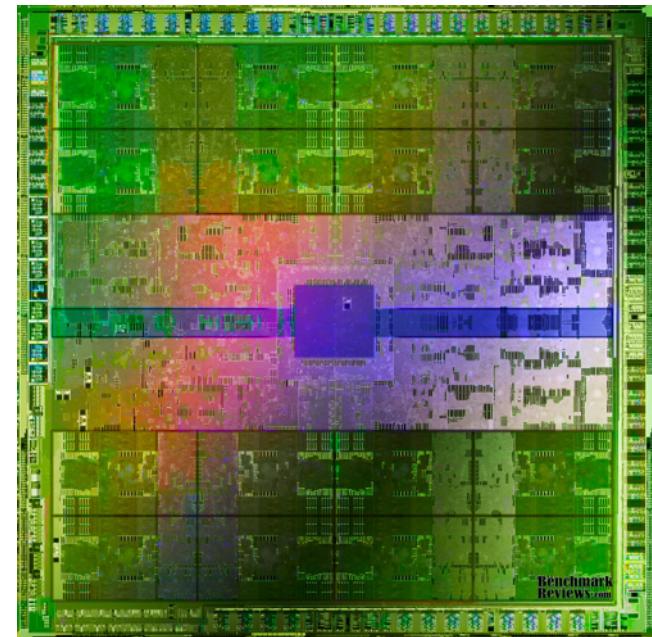


Large Scale Data Management on the GPU



Tim Kaldewey

Research Staff Member

IBM TJ Watson Research Center

tkaldew@us.ibm.com

Disclaimer

The author's views expressed in this presentation do not necessarily reflect the views of IBM.

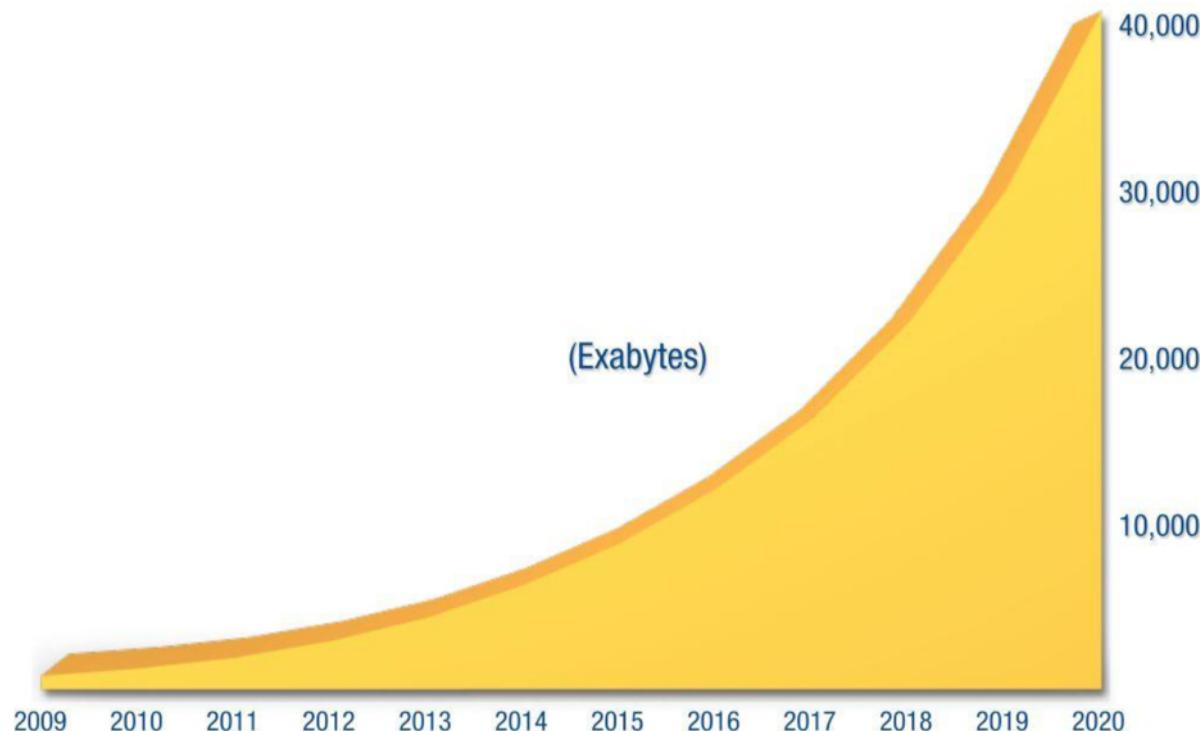
Acknowledgements

I would like to thank all my co-authors from IBM and my prior positions at Oracle and UCSC whose work I am showing in this presentation.

I would also like to thank Patrick Cozzi for inviting me to teach this class many years in a row.

Managing Big (!) Data

The Digital Universe: 50-fold Growth from the Beginning of 2010 to the End of 2020



Source: IDC's Digital Universe Study, sponsored by EMC, December 2012

- Supercomputing is tackling the Exaflop
- Data Management got past the Exabyte range already

Where does all the data come from ?

Longer Data Retention Regulations

+

Unstructured Data: image, text, audio, video

+

Data maintenance for business analytics

+

Increased number of transactions

= DATA GROWTH

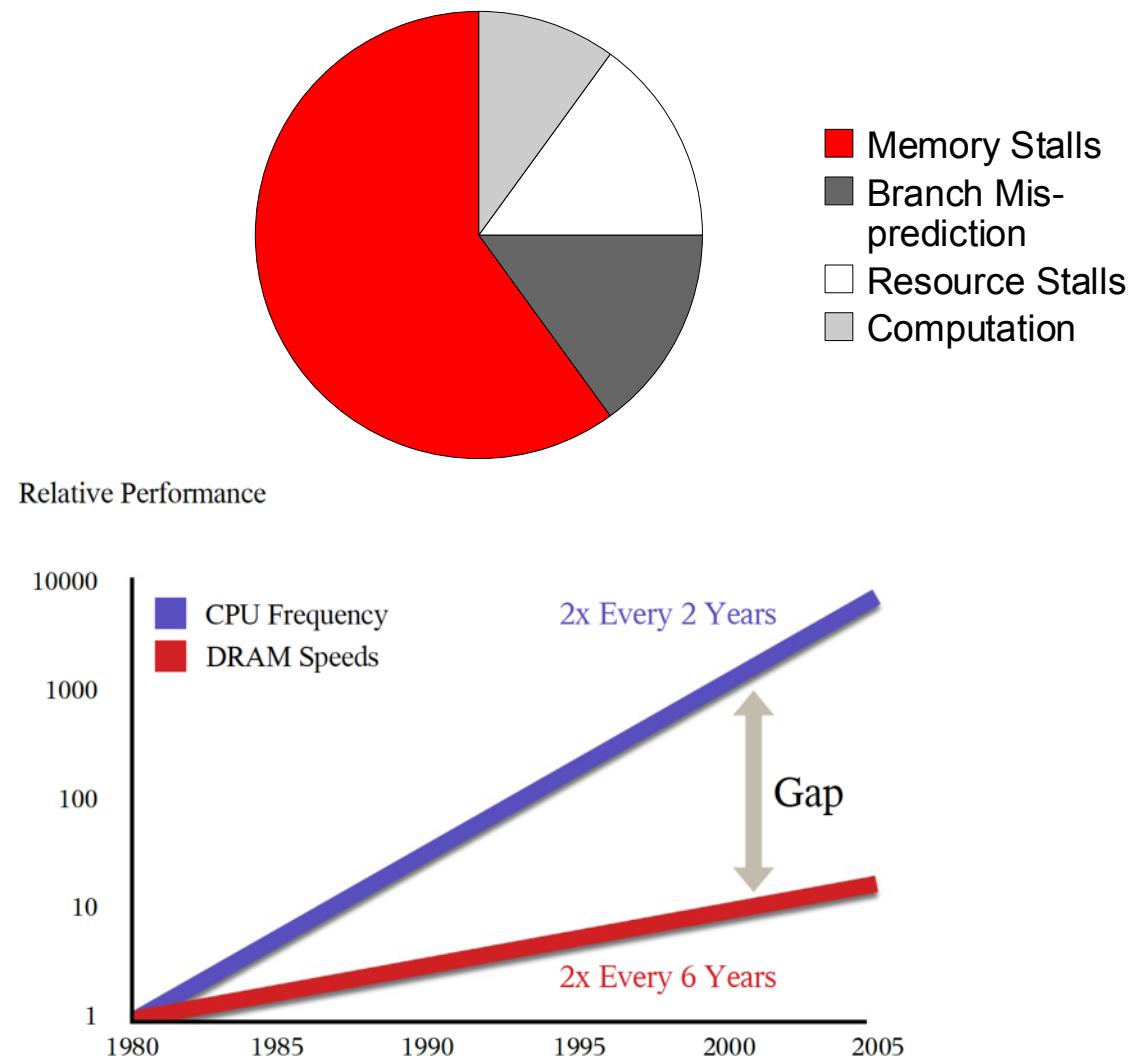


Source: Sybase Adaptive Server Enterprise Data Compression, SAP Sybase White paper 2012

- How to efficiently analyze/process rapidly increasing volumes of (business data) ?
- What is the problem ?
 - Aren't processors getting faster every day ?

DB Performance – Where does time go ?

- CPU? I/O? Memory?
 - Select 10% of rows based on an index ¹
- It's getting worse ²

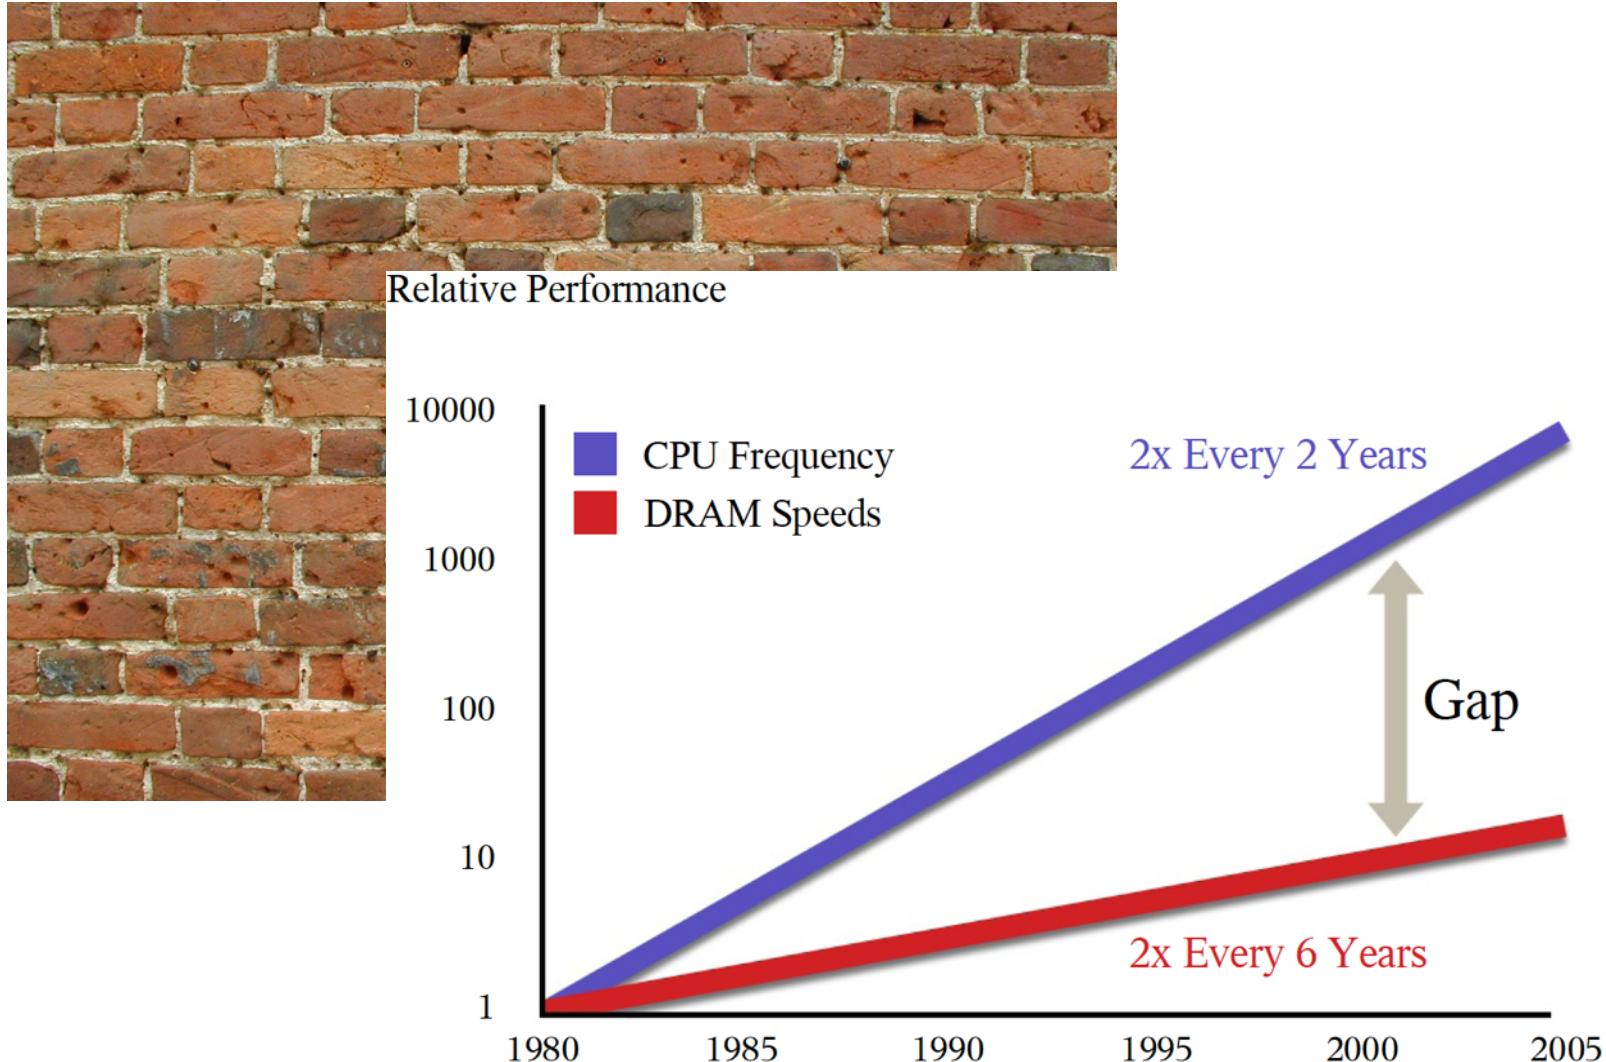


Sources:

¹ A. Ailamaki, et al. DBMSs on a modern processor: Where does time go? VLDB'99

² David Yen. Opening Doors to the MultiCore Era. MultiCore Expo 2006

The memory wall ¹

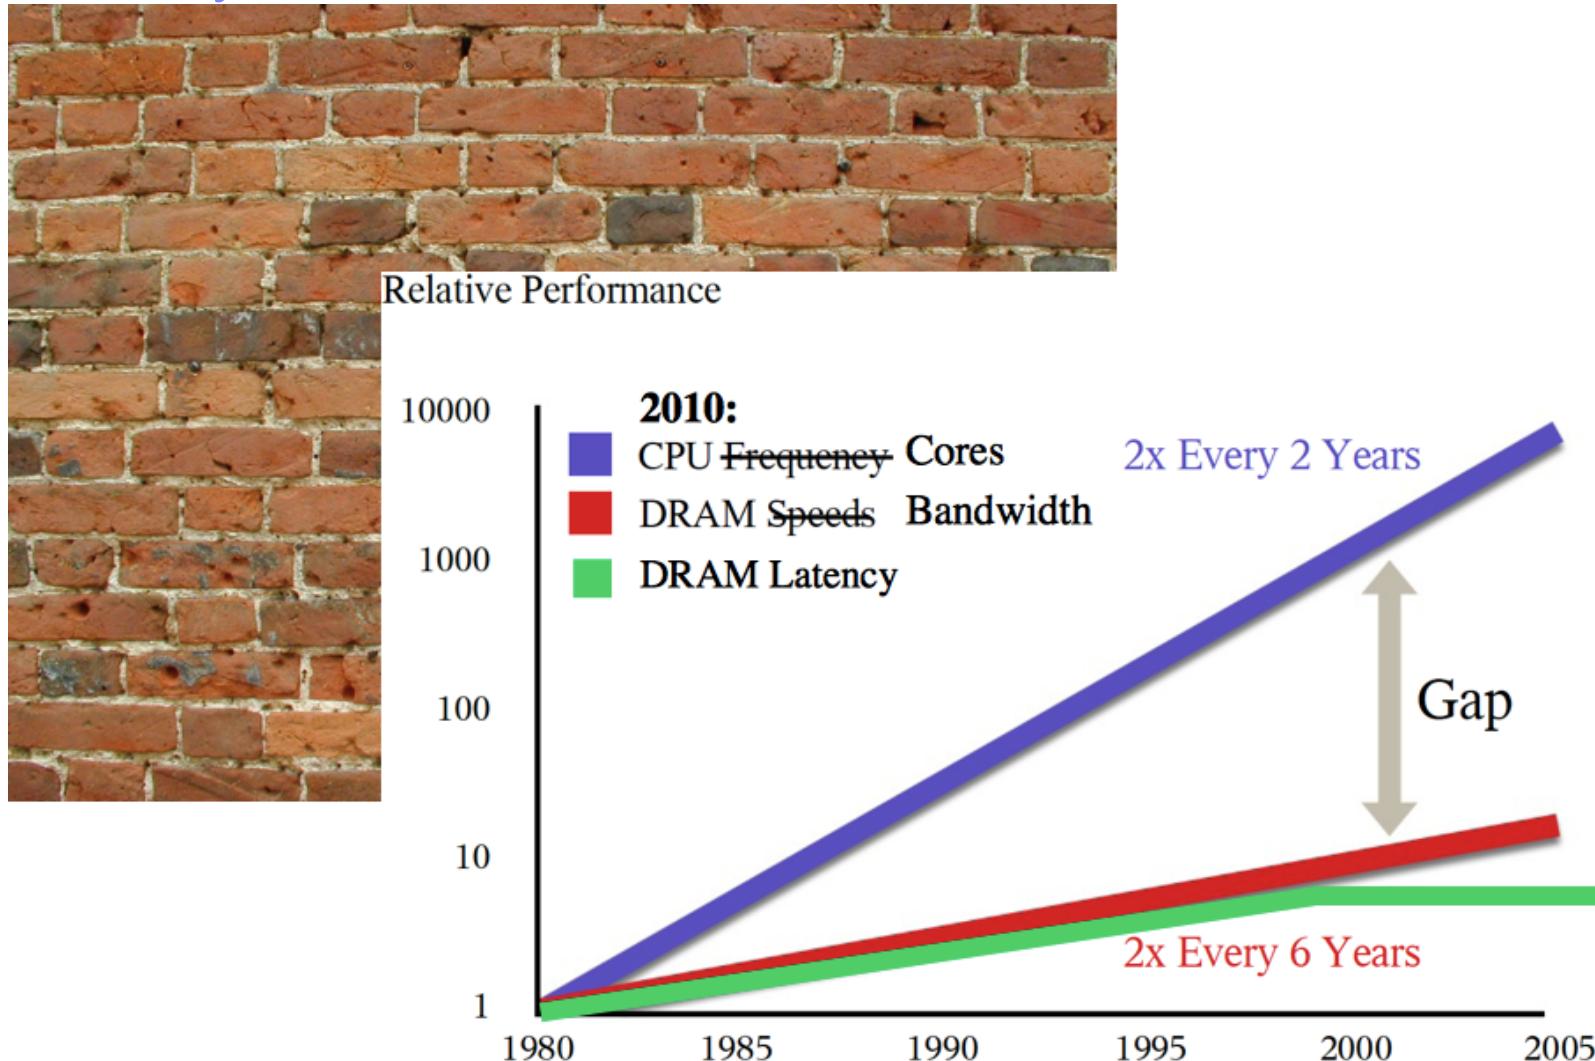


Sources:

¹ W.A.Wulf et al. Hitting the memory wall: implications of the obvious. SIGARCH - Computer Architecture News'95

² David Yen. Opening Doors to the MultiCore Era. MultiCore Expo 2006

The memory wall ¹

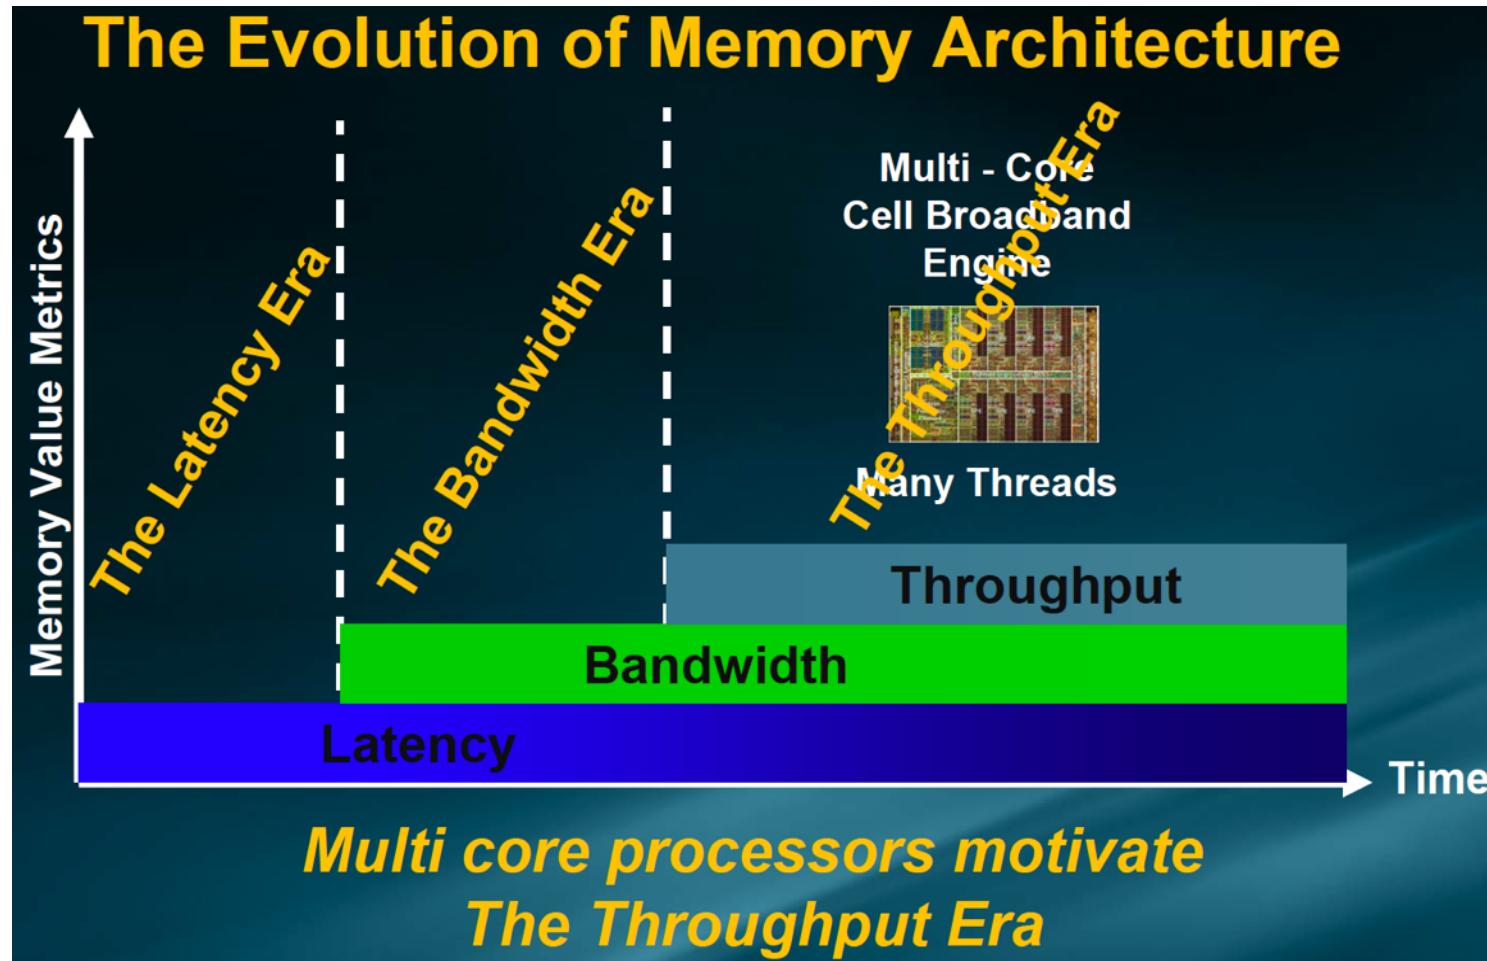


Sources:

¹ W.A.Wulf et al. Hitting the memory wall: implications of the obvious. SIGARCH - Computer Architecture News'95

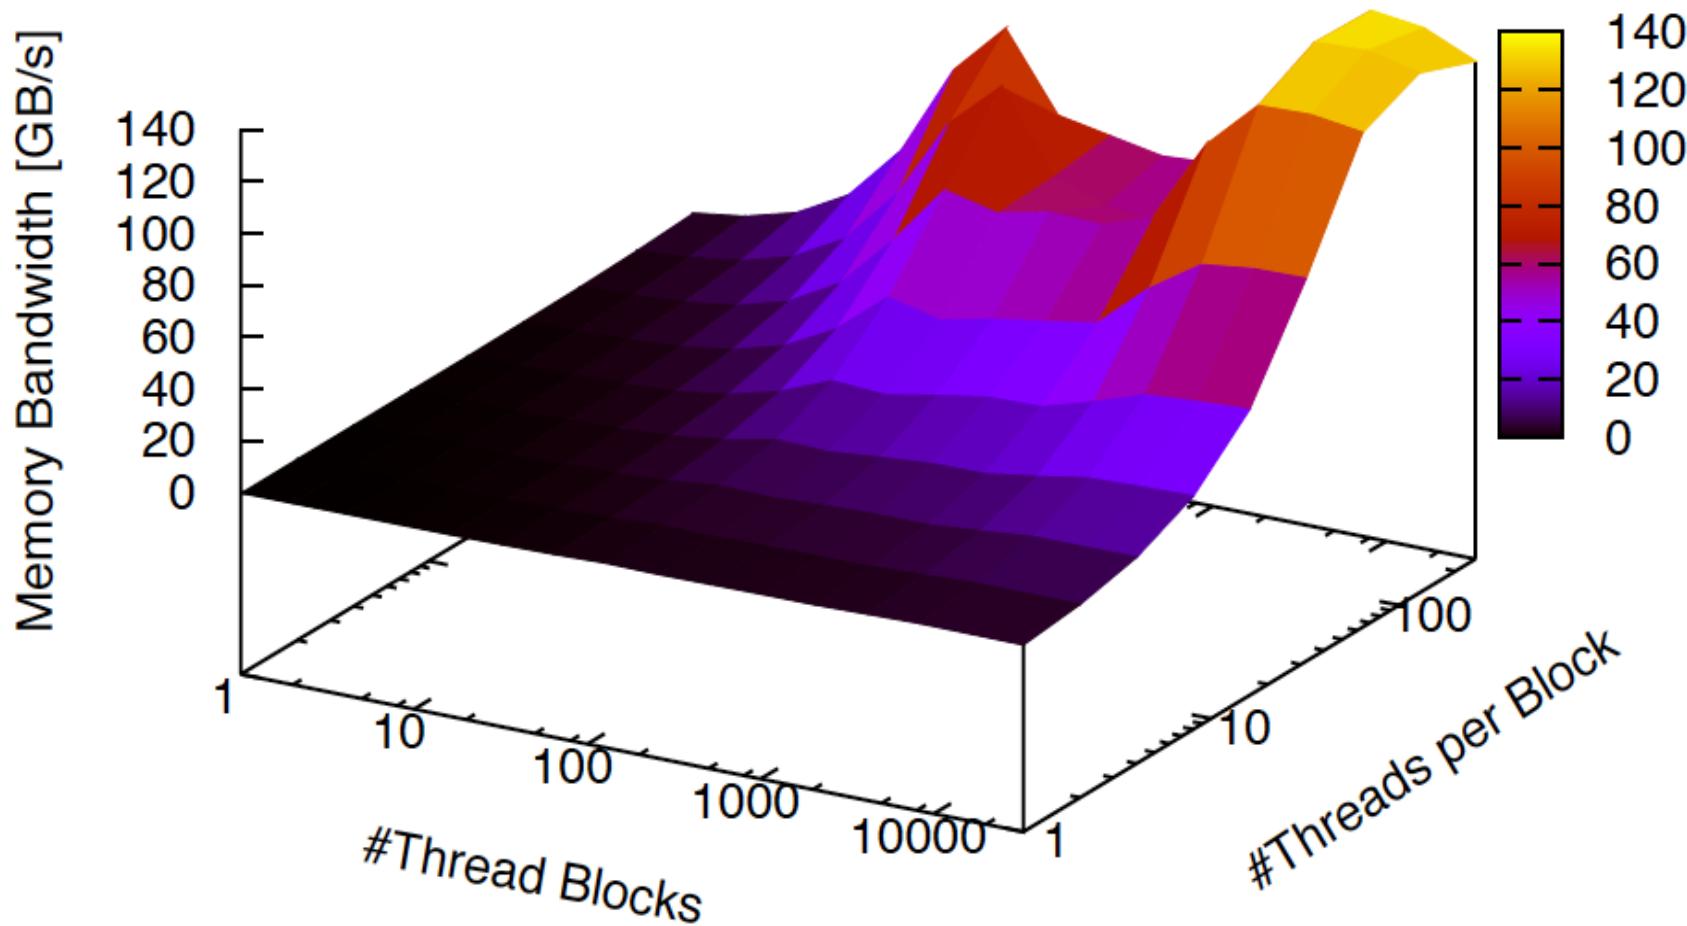
² David Yen. Opening Doors to the MultiCore Era. MultiCore Expo 2006

Parallel Processing – How does it affect memory performance?



Source: Terabyte Bandwidth Initiative, Craig Hampel - Rambus, HotChips'08

GPU's are a Poster Child of the Throughput Era



Coalesced 32-bit read access on an nVidia GTX 285 1.5GHz, GDDR3 1.2GHZ.

- Required level of concurrency to reach peak performance depends on thread configuration.

Agenda

Large Scale Data Management on the GPU

- Why GPUs for information management workloads?
 - Search as an example
- Maximizing Device memory access performance
 - Coalescing
 - Thread configuration
 - Large(r) data sets
- GPU data transfers
 - Conventional
 - CUDA Streams
 - Zero Copy Access / Universal Virtual Addressing
- Search again
 - A naïve implementation

Why search?

- Honestly, how many times a day do you visit



The Google logo, featuring the word "Google" in its signature multi-colored, rounded font. The letters are blue, red, yellow, and green. A small "TM" symbol is located at the top right of the letter "e".



The Yahoo! logo, featuring the word "YAHOO!" in a large, bold, red, sans-serif font. A small registered trademark symbol (®) is located at the top right of the letter "O".

?

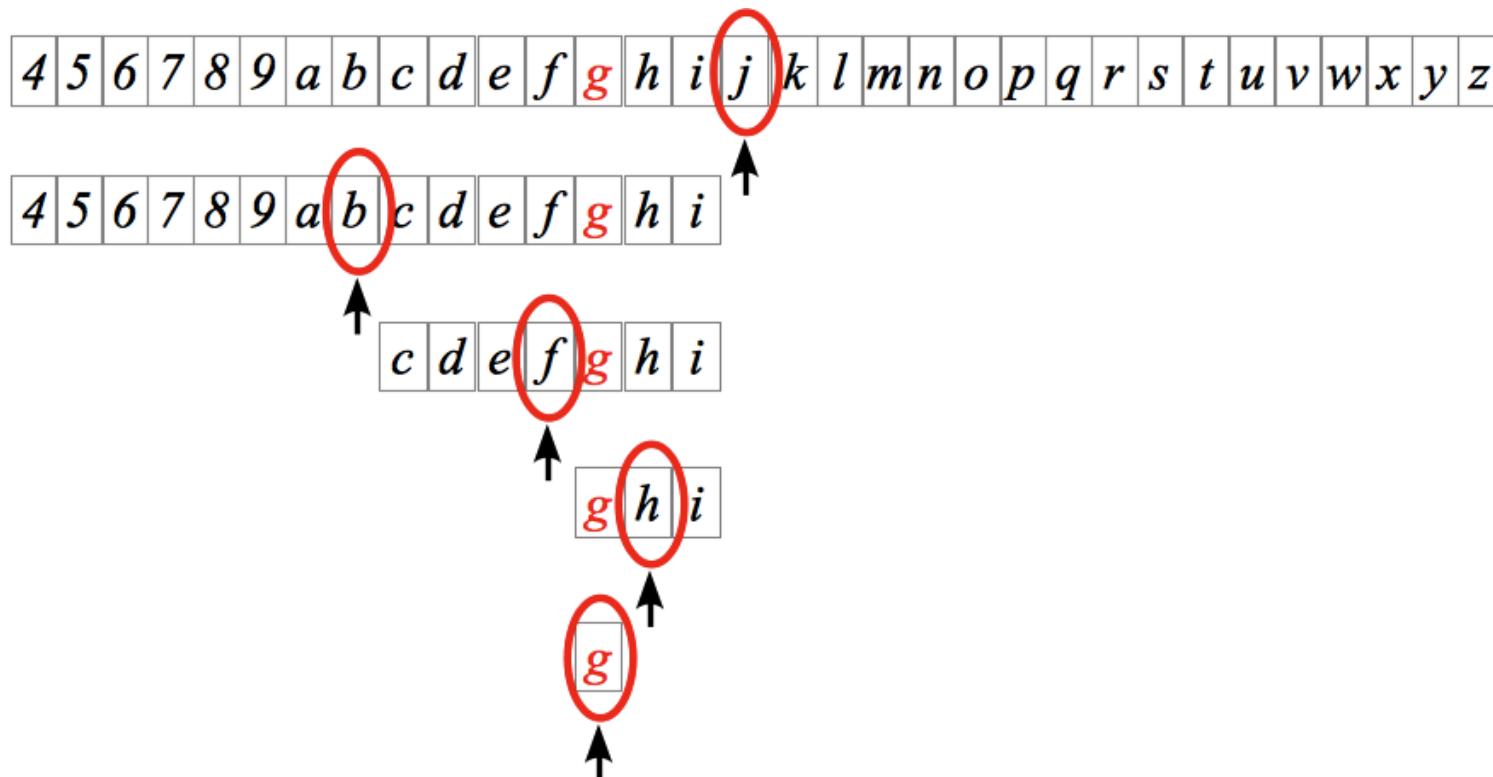
Search – Memory Access Pattern

- Binary search (in a sorted list)
 - Choose Pivot Element at position $\#elements / 2$
 - Compare if equal, larger, or smaller than search key (**g**)
 - Equal → Done
 - Smaller → proceed with lower half of the list
 - Larger → proceed with upper half



Search – Memory Access Pattern

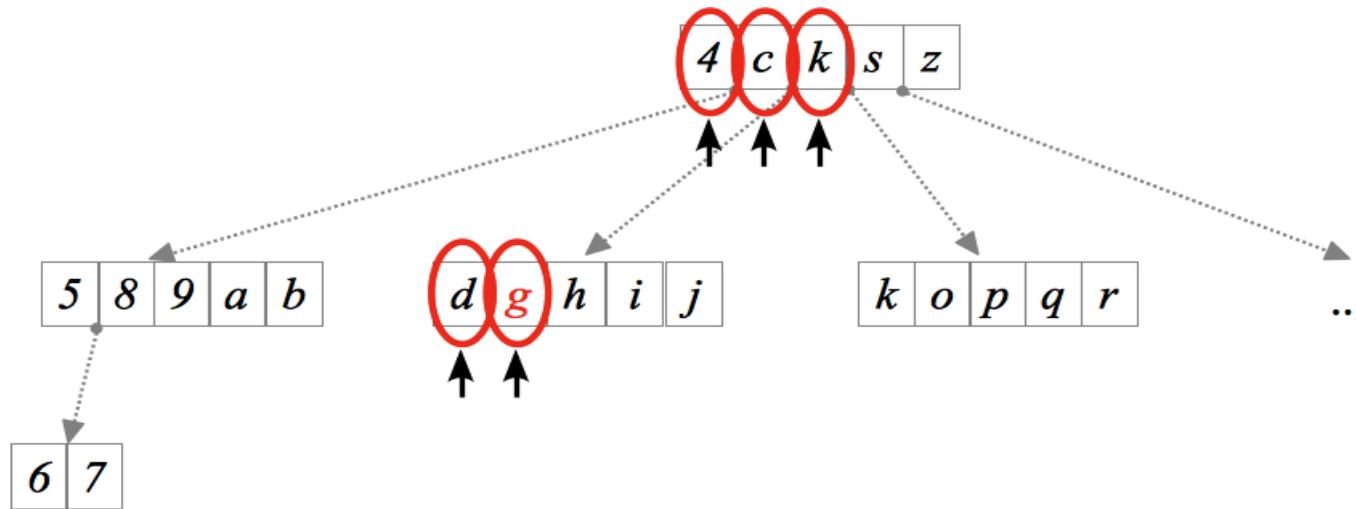
- Binary search (in a sorted list)
 - Choose Pivot Element at position $\#elements / 2$
 - Compare if equal, larger, or smaller than search key (g)
 - Equal → Done
 - Smaller → proceed with lower half of the list
 - Larger → proceed with upper half



- Data dependent, quasi-random access pattern

Search – Memory Access Pattern

- B-trees group pivot elements = making access pattern more linear:



- How do you store B-trees?
- Still have quasi random access when traversing nodes!
- Overhead of maintaining trees?
 - Requires atomics if done in parallel ...

GPU – Memory access performance



	GPU (GTX580)	CPU (i7-2600)
Peak memory bandwidth [spec] ¹⁾	179 GB/s	21 GB/s
Peak memory bandwidth [measured] ²⁾	153 GB/s	18 GB/s
Random access [measured] ²⁾	6.6 GB/s	0.8 GB/s
Compare and swap [measured] ³⁾	4.6 GB/s	0.4 GB/s

How do you
achieve this?



~ 1 order of magnitude

(1) Nvidia: 192.4×10^6 B/s \approx 179.2 GB/s

(2) 64-bit accesses over 1 GB of device memory

(3) 64-bit compare-and-swap to random locations over 1 GB device memory

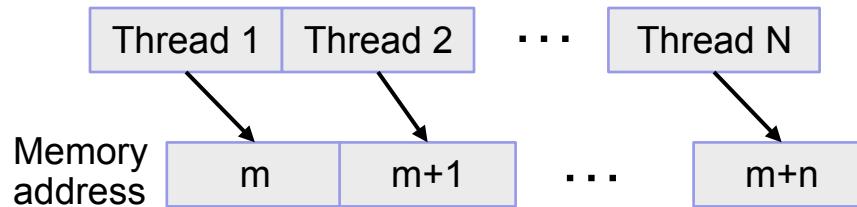
Agenda

Large Scale Data Management on the GPU

- Why GPUs for information management workloads?
 - Search as an example
- Maximizing Device memory access performance
 - Coalescing
 - Thread configuration
 - Large(r) data sets
- GPU data transfers
 - Conventional
 - CUDA Streams
 - Zero Copy Access / Universal Virtual Addressing
- Search again
 - A naïve implementation

Coalesced Memory access (reminder)

- Ideal memory access pattern is coalesced memory access
 - Threads of a block/warp access consecutive memory addresses

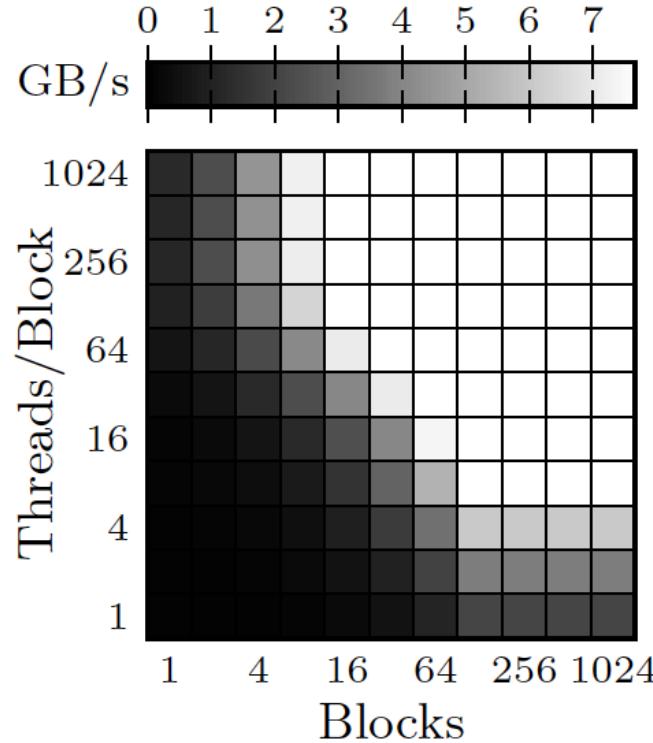


```
for (position = blockIdx.x*blockDim.x+threadIdx.x;  
     position < length;  
     position += blockDim.x*gridDim.x) {  
    x = list[position];  
    ...
```

- This works great if we just read the entire data set, e.g. image processing
- Don't caches (Fermi and newer) make this obsolete?
- Random access is often the norm for information management workloads =(

“Fast” random access

- Performance depends on thread configuration

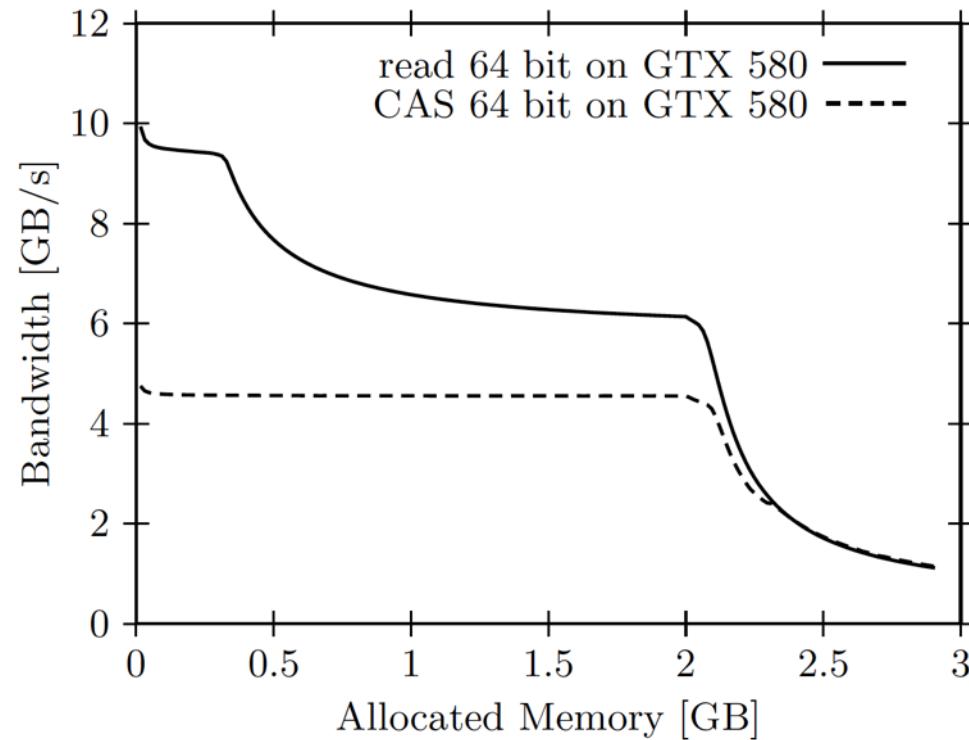


64-bit random access within 512 MB of device memory on GTX 580

- Assume that each access incurs device memory latency
- Use sufficient threads to hide memory latency
- How many?
 - Depends on specific GPU: memory interface, latency, ...

Random access in large data sets

- Performance depends on data set size



64-bit random access within 512 MB of device memory on GTX 580

- Assume sufficient threads/blocks to hide memory latency
- What causes the 3 dips?

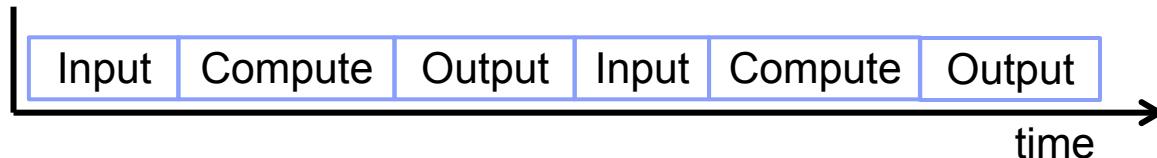
Agenda

Large Scale Data Management on the GPU

- Why GPUs for information management workloads?
 - Search as an example
- Maximizing Device memory access performance
 - Coalescing
 - Thread configuration
 - Large(r) data sets
- GPU data transfers
 - Conventional
 - CUDA Streams
 - Zero Copy Access / Universal Virtual Addressing
- Search again
 - A naïve implementation

How to get data to the GPU (and back):

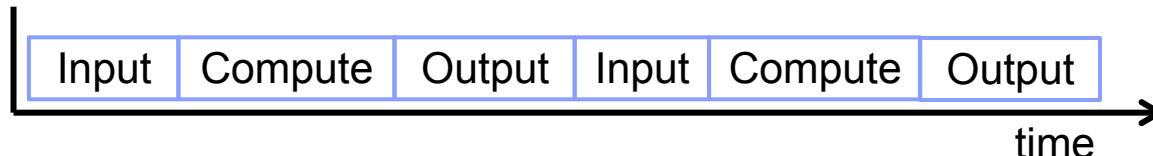
- Synchronous:



– Copy data to GPU, compute, copy results back to CPU

How to get data to the GPU (and back):

- Synchronous:



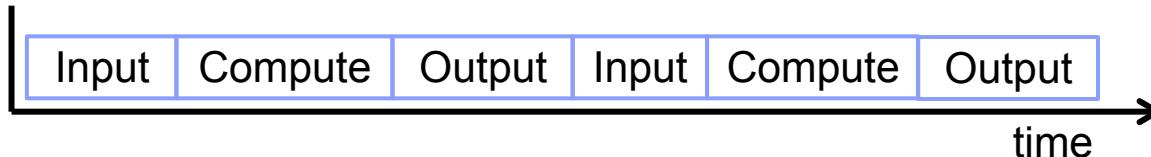
- Copy data to GPU, compute, copy results back to CPU

```
// make space for input data & results
cudaMalloc(...);
// copy input data to GPU
cudaMemcpy(Input_Dev, Input_Host, size, cudaMemcpyHostToDevice);
// call compute Kernel
computeKernel<<dimGrid, dimBlock>>(Dev* ...)

...
// copy results
cudaMemcpy(Output_Host, Ouput_dev, size, cudaMemcpyDeviceToHost);
```

How to get data to the GPU (and back):

- Synchronous:



- Copy data to GPU, compute, copy results back to CPU

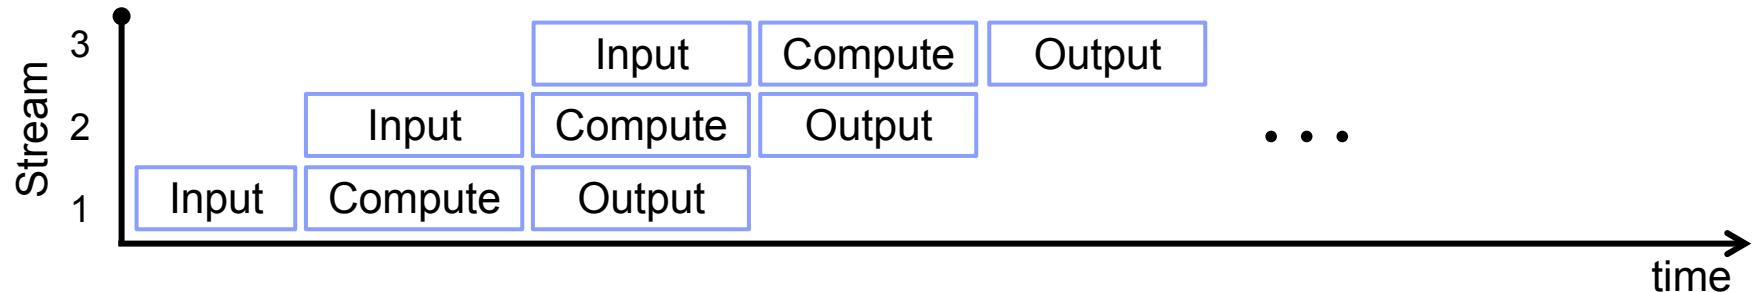
```
// make space for input data & results
cudaMalloc(...);
// copy input data to GPU
cudaMemcpy(Input_Dev, Input_Host, size, cudaMemcpyHostToDevice);
// call compute Kernel
computeKernel<<dimGrid, dimBlock>>(Dev* ...)

...
// copy results
cudaMemcpy(Output_Host, Ouput_dev, size, cudaMemcpyDeviceToHost);
```

- Common for most problems where the compute phase is long
- Caveats:
 - Data set size (input + output) limited by GPU memory size!
 - PCI-E idle during compute, cores idle during data transfer

How to get data to the GPU (and back):

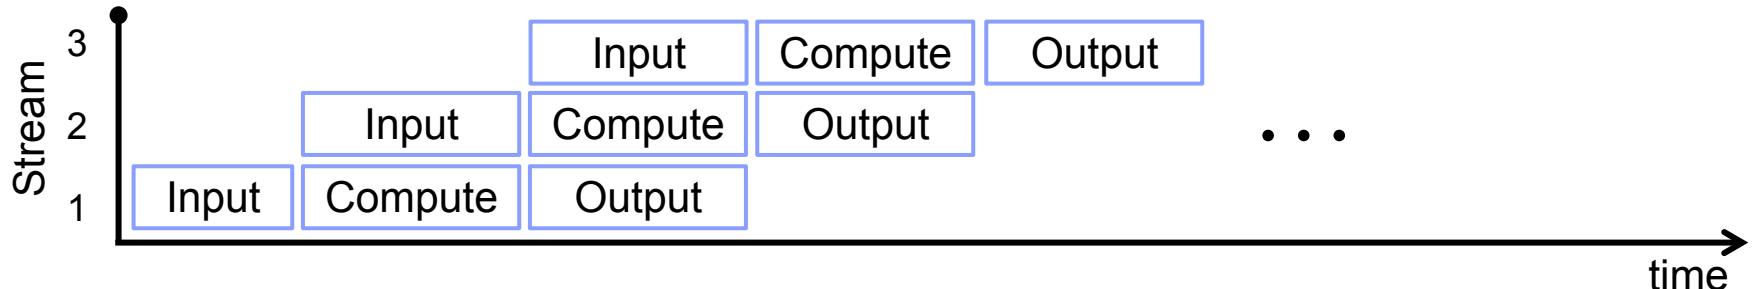
- Asynchronous data transfers, aka CUDA Streams:



- While one stream is executing a kernel, others can copy data in (and out)
 - Effectively overlap data copies with compute

How to get data to the GPU (and back):

- Asynchronous data transfers, aka CUDA Streams:

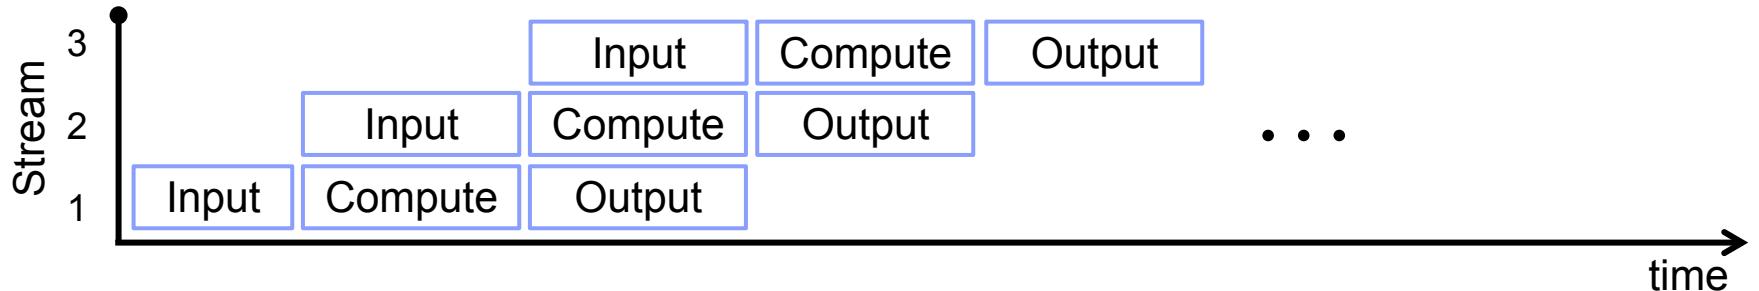


- While one stream is executing a kernel, others can copy data in (and out)
 - Effectively overlap data copies with compute

```
for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&d_ipt[offset], &h_ipt[offset], streamBytes,
                   cudaMemcpyHostToDevice, stream[i]);
    computeKernel<<..., stream[i]>>(d_ipt, d_opt, offset);
    cudaMemcpyAsync(&h_opt[offset], &d_opt[offset], streamBytes,
                   cudaMemcpyDeviceToHost, stream[i]);
}
```

How to get data to the GPU (and back):

- Asynchronous data transfers, aka CUDA Streams:



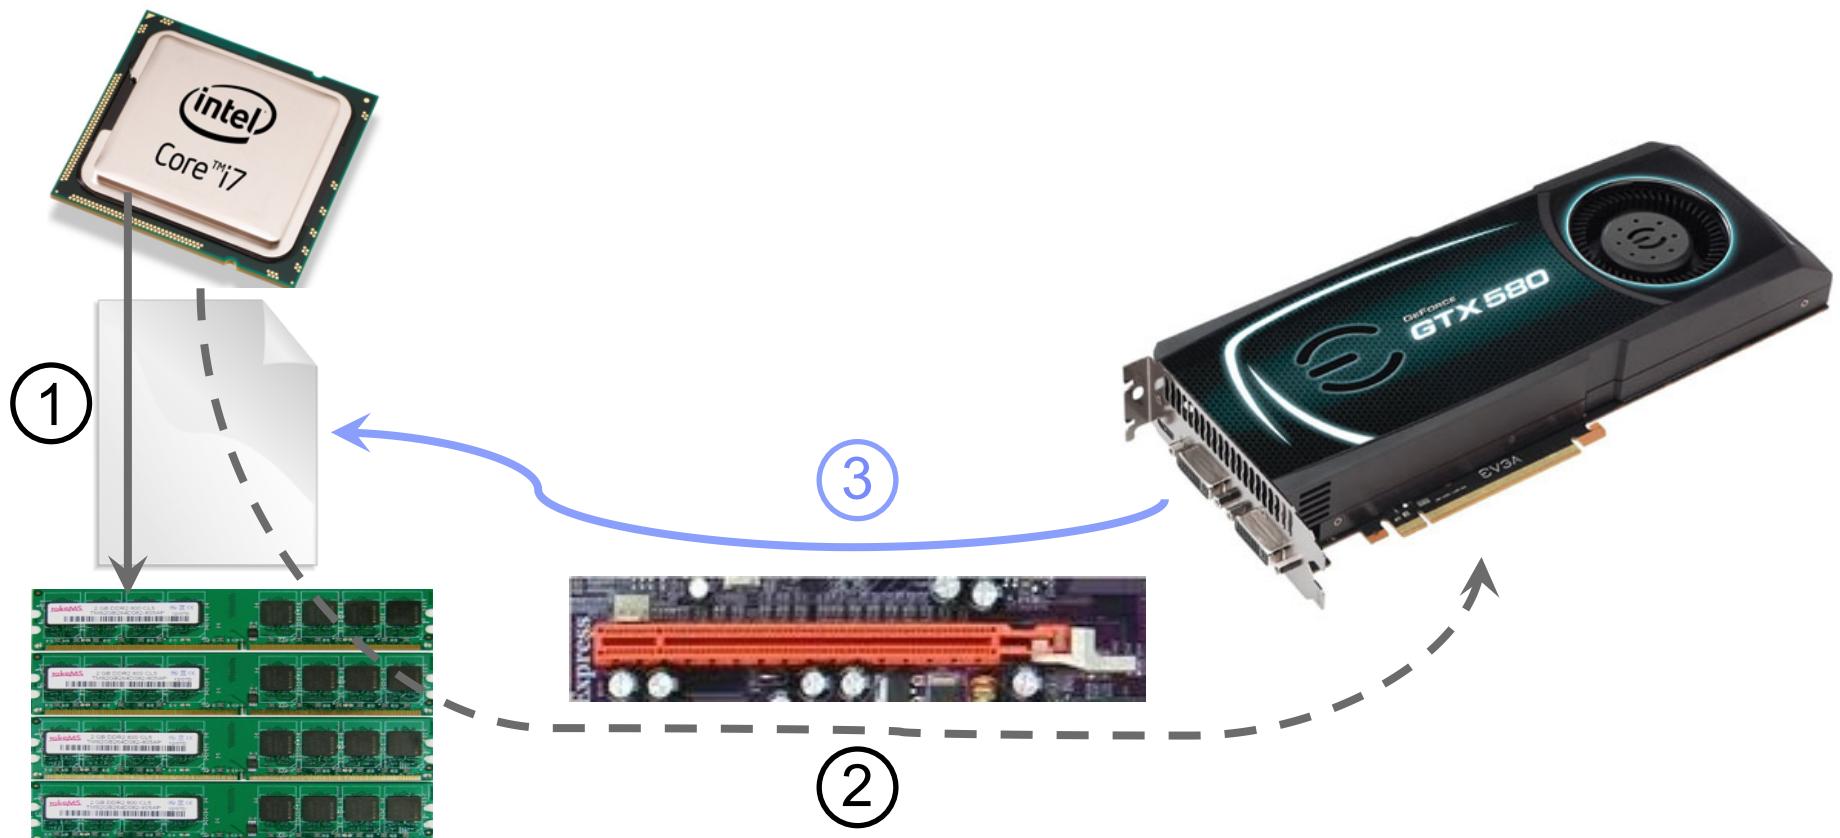
- While one stream is executing a kernel, others can copy data in (and out)
 - Effectively overlap data copies with compute

```
for (int i = 0; i < nStreams; ++i) {
    int offset = i * streamSize;
    cudaMemcpyAsync(&d_ipt[offset], &h_ipt[offset], streamBytes,
                   cudaMemcpyHostToDevice, stream[i]);
    computeKernel<<..., stream[i]>>(d_ipt, d_opt, offset);
    cudaMemcpyAsync(&h_opt[offset], &d_opt[offset], streamBytes,
                   cudaMemcpyDeviceToHost, stream[i]);
}
```

- If data transfers faster than compute, 100% time available for compute
- Caveats:
 - $2 * \text{Input set} + 2 * \text{output(results)} < \text{GPU memory size!}$
 - What if arithmetic intensity is low (like most database operations)?

How to get data to the GPU (and back):

- Zero Copy Access (ZCA)
 - Allows the GPU to access CPU memory “directly”, i.e., without placing a copy in GPU in device memory:
 1. CPU thread pins memory page & obtains phys. address
 2. Function call contains a pointer to the phys. address
 3. GPU(CUDA) threads execute **load/store** instructions



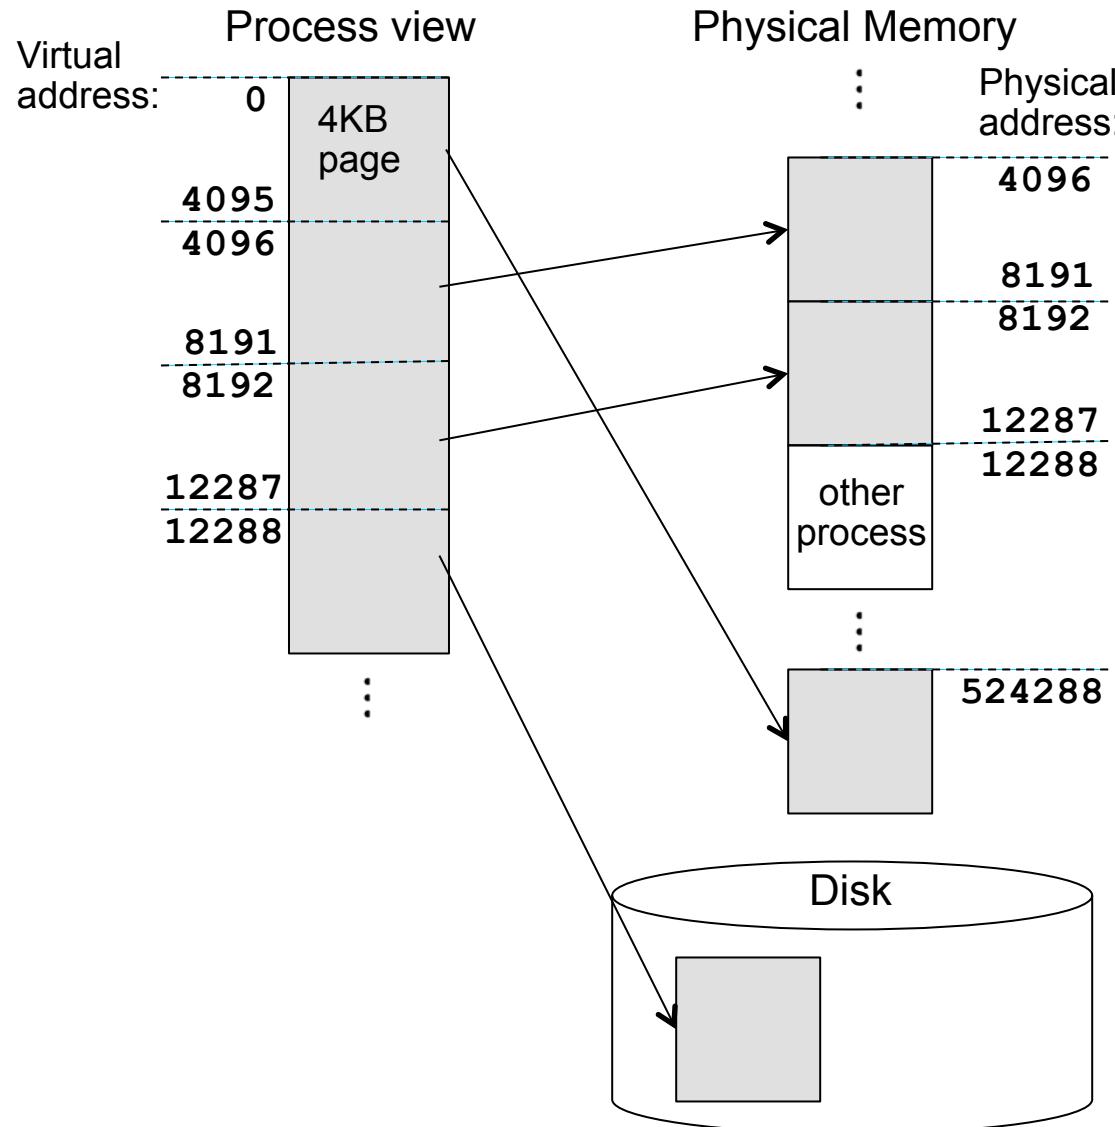
How to get data to the GPU (and back):

- Zero Copy Access (ZCA)

1. CPU thread pins memory page & creates pointer to phys. address

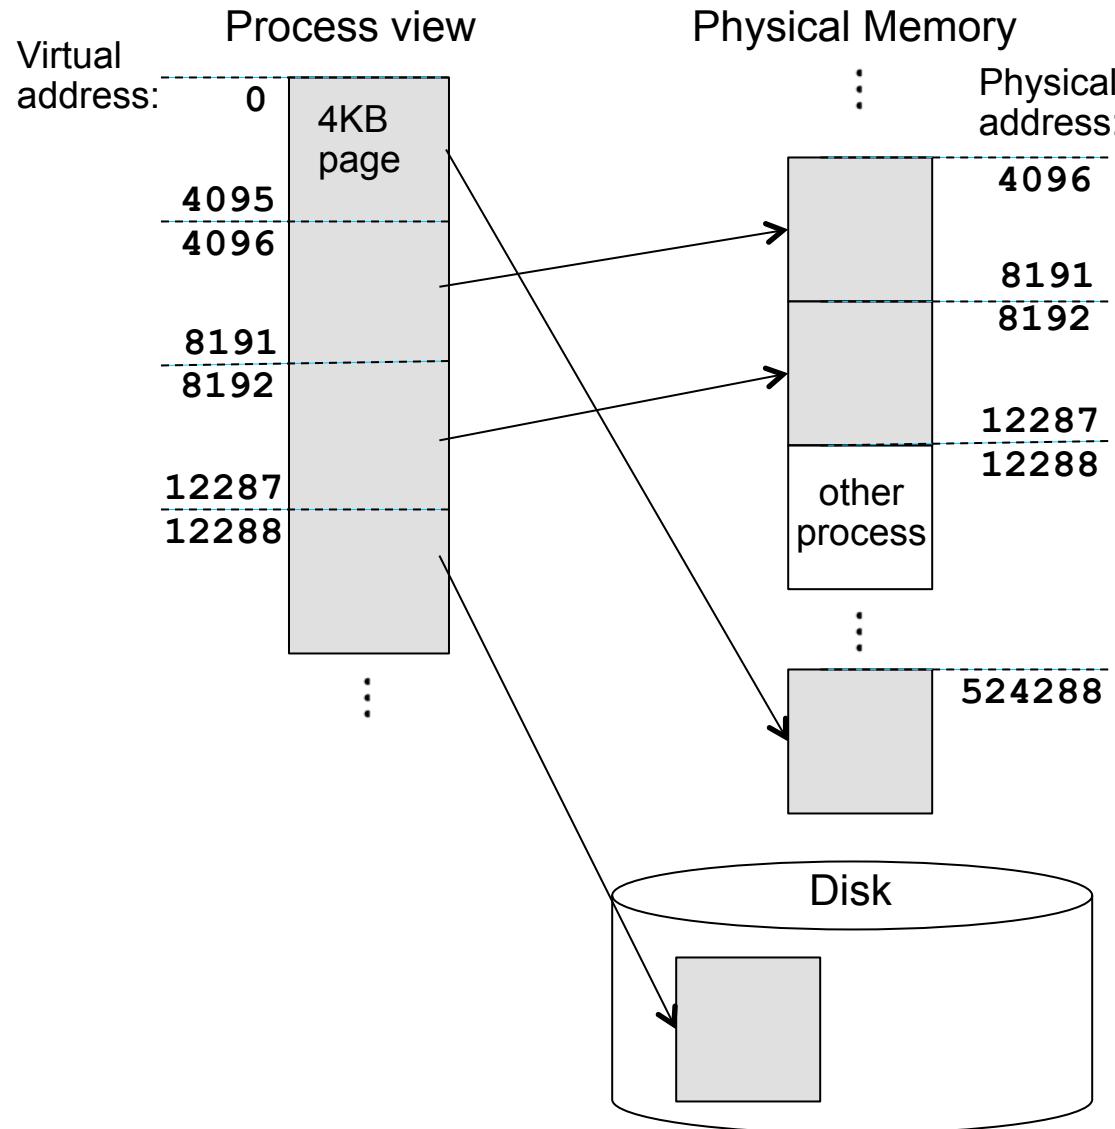
```
cudaHostRegister(h_ipt, ipt_size, cudaHostRegisterMapped) ;  
cudaHostRegister(h_opt, opt_size, cudaHostRegisterMapped) ;  
cudaHostGetDevicePointer(&ipt_ptr, h_ipt, 0) ;  
cudaHostGetDevicePointer(&opt_ptr, h_opt, 0) ;
```

Zero Copy Access & Virtual memory management



- OS provides virtual to physical address translation
- Memory pages (usually 4KB) can be physically stored on disk
- GPU kernel knows nothing about virtual memory

Zero Copy Access & Virtual memory management



- OS provides virtual to physical address translation
- Memory pages (usually 4KB) can be physically stored on disk
- GPU kernel knows nothing about virtual memory
- `cudaHostRegister()` makes a call to `mlock()` which marks page(s) as not page-able and assures pages are contiguous
- Ensures that page(s) will **ALWAYS** be present in **physical** memory
- `cudaHostGetDevicePointer()` obtains **physical** memory address

How to get data to the GPU (and back):

- Zero Copy Access (ZCA)

1. CPU thread pins memory page & creates pointer to phys. address

```
cudaHostRegister(h_ip, ip_size, cudaHostRegisterMapped);  
cudaHostRegister(h_opt, opt_size, cudaHostRegisterMapped);  
cudaHostGetDevicePointer(&ip_ptr, h_ip, 0);  
cudaHostGetDevicePointer(&opt_ptr, h_opt, 0);
```

2. Function call contains pointer(s) to host address(es)

```
computeKernel<<<dimGrid, dimBlock>>>  
    (ip_ptr, ip_size, opt_ptr);
```

How to get data to the GPU (and back):

- Zero Copy Access (ZCA)

1. CPU thread pins memory page & creates pointer to phys. address

```
cudaHostRegister(h_ip, ip_size, cudaHostRegisterMapped);  
cudaHostRegister(h_opt, opt_size, cudaHostRegisterMapped);  
cudaHostGetDevicePointer(&ipt_ptr, h_ip, 0);  
cudaHostGetDevicePointer(&opt_ptr, h_opt, 0);
```

2. Function call contains pointer(s) to host address(es)

```
computeKernel<<<dimGrid, dimBlock>>>  
    (ipt_ptr, ip_size, opt_ptr);
```

3. CUDA kernel/threads execute load/stores instructions as before

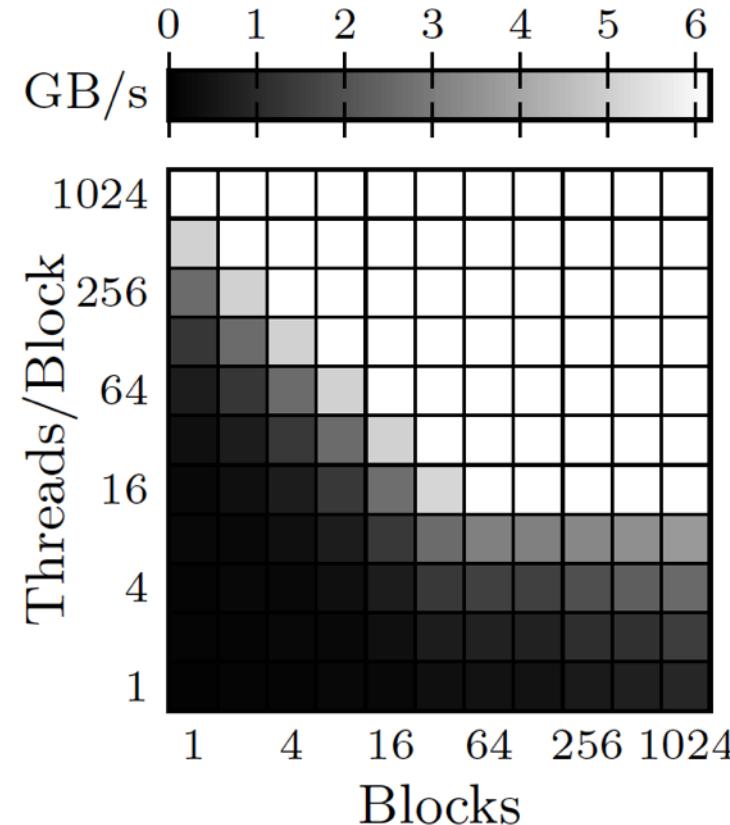
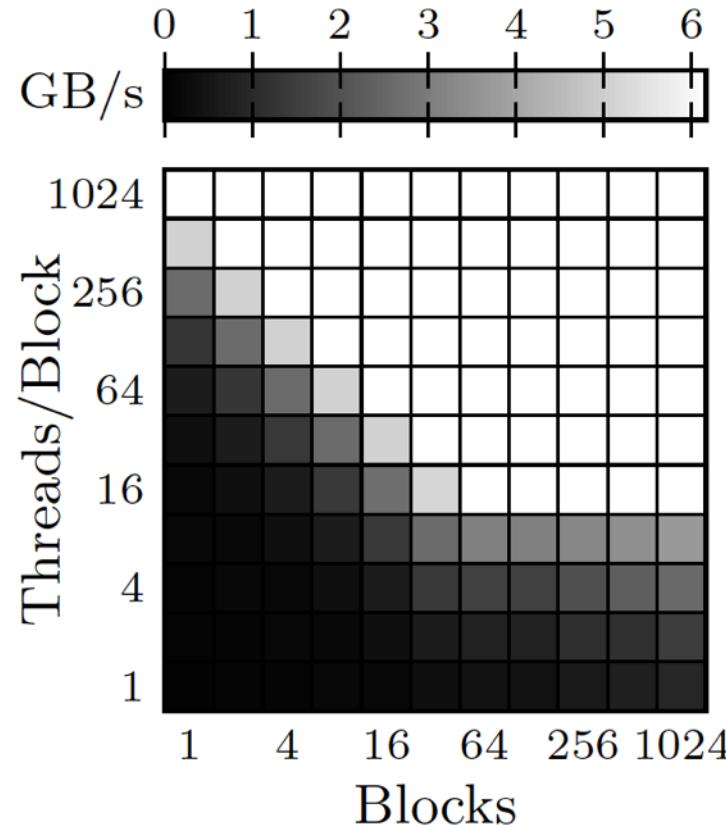
```
__global__ computeKerne(float* ipt_ptr, int ip_size,  
                        float* opt_ptr){  
    int entry=blockIdx.x*blockDim.x + threadIdx.x;  
    while(entry < ip_size){  
        opt_ptr[entry] = ipt_ptr[entry] + 42.0;  
        entry += blockDim.x*blockDim.x;  
    }  
}
```

How to get data to the GPU (and back):

- Zero Copy Access (ZCA)
- What is the maximum (theoretical) performance ?
 - PCI-E 2.0 operates at 5 GigaTransfers(bit) / second per lane
 - Most GPUs use 16 lanes
 - $5 \text{ Gbit/s} * 16 \text{ lanes} = 80 \text{ Gbit/s}$
 - 8/10 bit encoding reduces this to 64 Gbit/s payload or 8 GB/s
 - Packet protocol with 128B payload + 24B header reduces it to 6.27GB/s
- PCI-E 3.0 operates at 8 GT/s
- 16-lane GPU interface
- 128/130 encoding
- Bandwidth = ?
- Can this be achieved in practice (for PCI-E 2.0) ?

How to get data to the GPU (and back):

- Zero Copy Access (ZCA) – measured performance

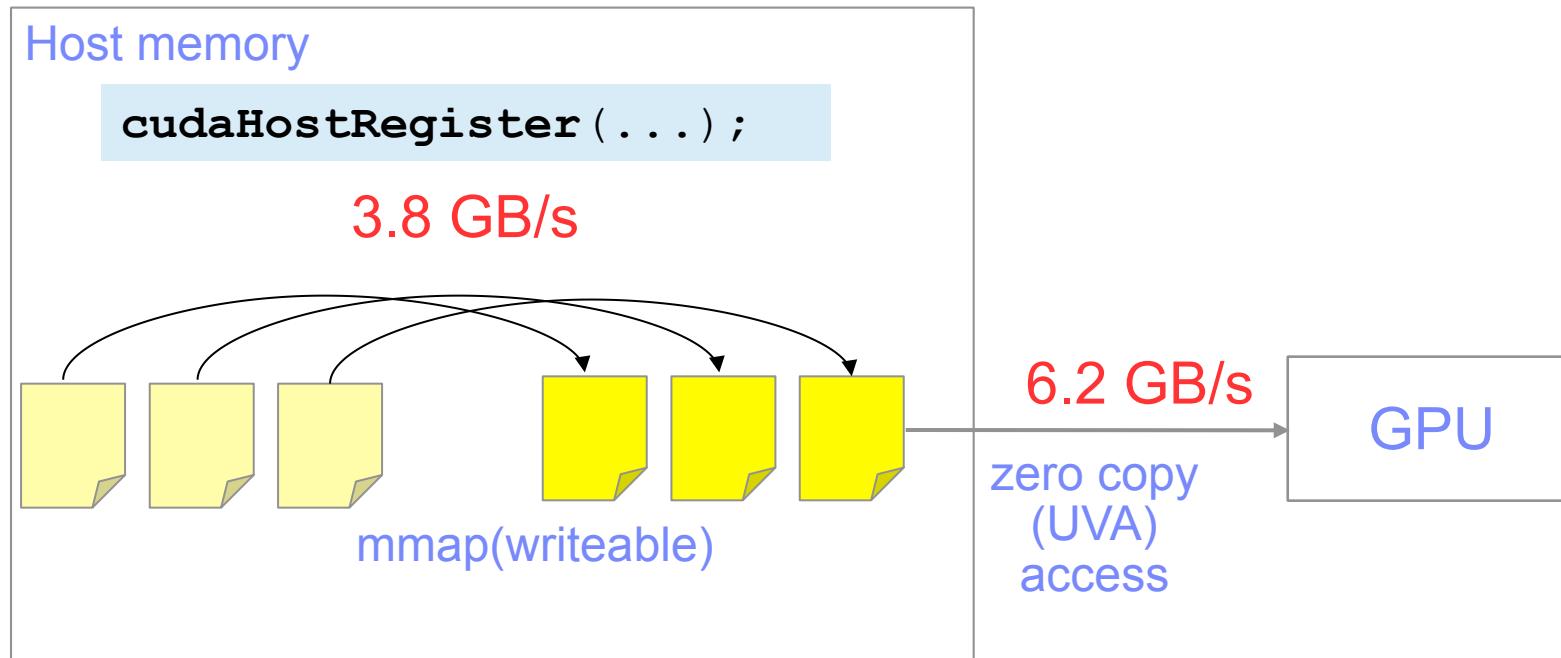


64-bit coalesced read/write access to host memory using ZCA on a GTX580

- What about read & write at the same time ?

How to get data to the GPU (and back):

- Zero Copy Access (ZCA) – pinning



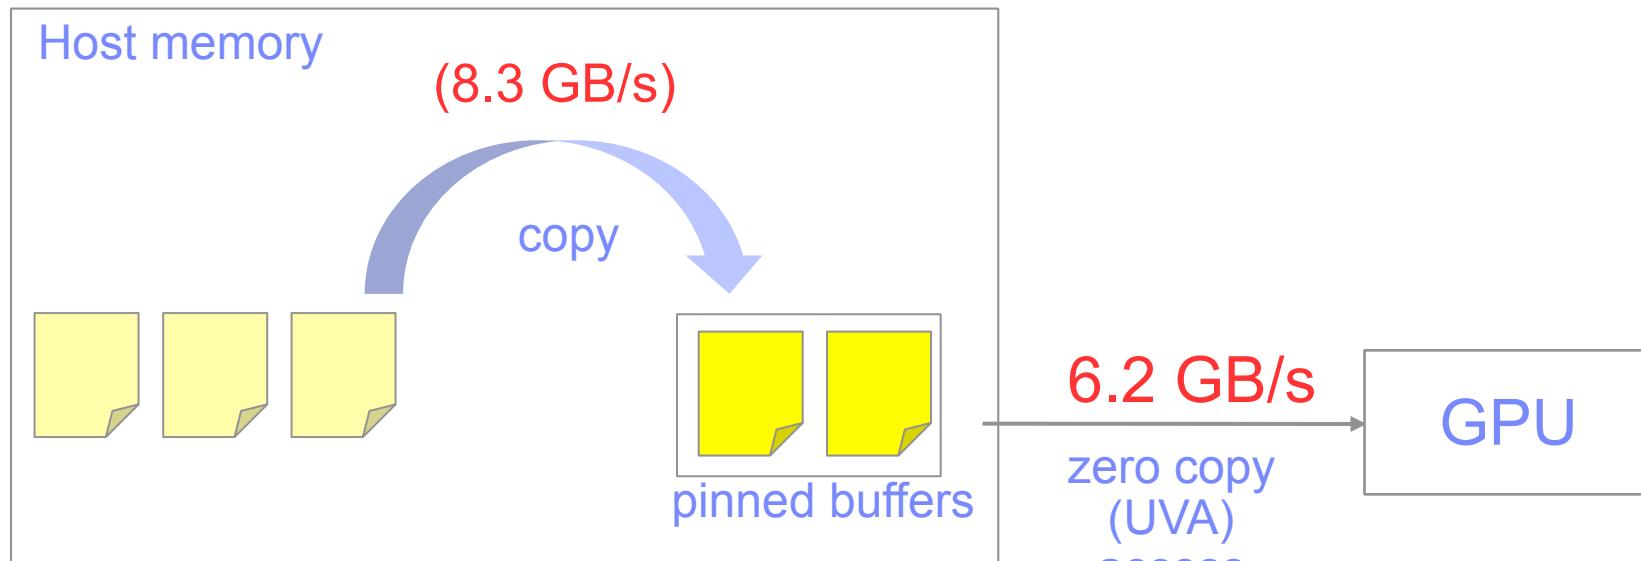
Register Pages

- GPU accessible pages must be pinned (why?)
- GPU accessible pages must be writable (why?)
 - copy page before pinning

Registering (and un-registering) is a bottleneck!

- $3.8 \text{ GB/s} < 6.2 \text{ GB/s}$ on PCIe

How to get data to the GPU (and back): Zero Copy Access (ZCA) –pinning “workaround”



Idea

- Use double buffering scheme
 - Pre-allocate pinned memory buffer
 - Copy data into one pre-allocated buffer
 - GPU kernel can operate on the second

→ Bottleneck is now PCI-E again

How about PCI-E 3.0?

Agenda

Large Scale Data Management on the GPU

- Why GPUs for information management workloads?
 - Search as an example
- Maximizing Device memory access performance
 - Coalescing
 - Thread configuration
 - Large(r) data sets
- GPU data transfers
 - Conventional
 - CUDA Streams
 - Zero Copy Access / Universal Virtual Addressing
- Search again
 - A naïve implementation

A Simple implementation of (index) search

Keyword Document ID

Adam	1, 2, 3
Bethlehem	4, 5
Character	1, 2, 3, 301, 5790
Drachenflieger	301, 317, 5790
Eva	1, 2
Flughafenbahnhof	5790
Grabdenkmal	2, 5790
Haubentaucher	300, 5790

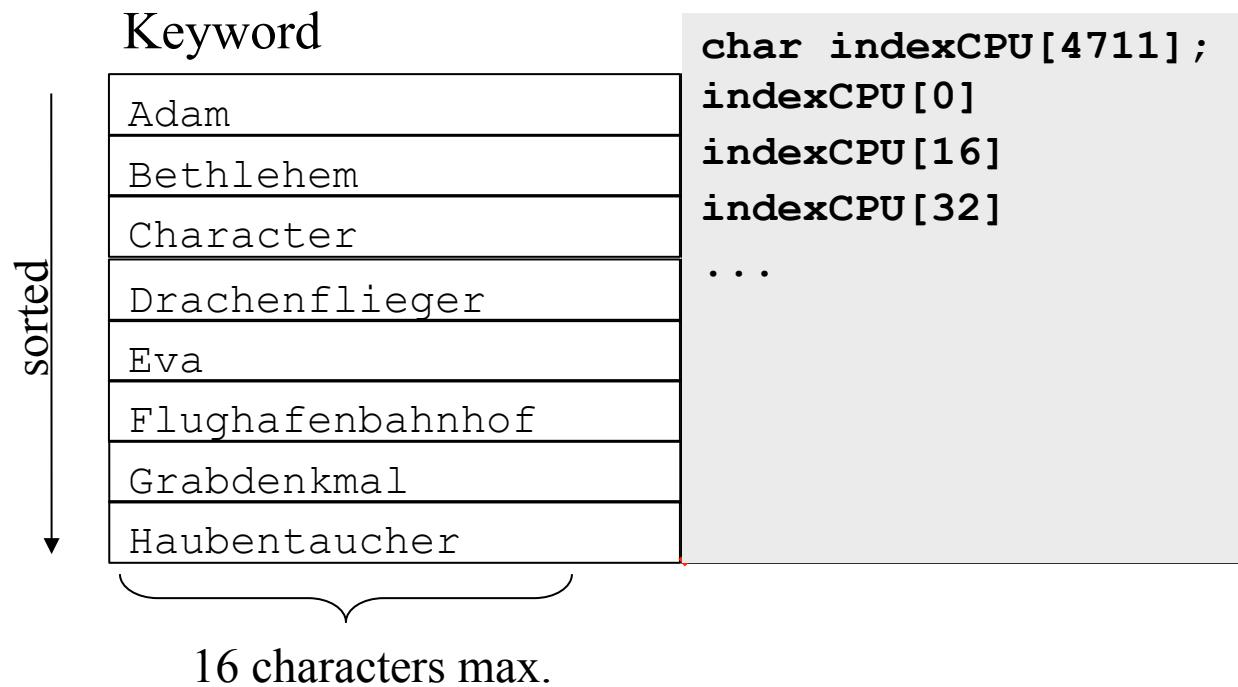
A Simple implementation of (index) search

Keyword	Document ID
Adam	1, 2, 3
Bethlehem	4, 5
Character	1, 2, 3, 301, 5790
Drachenflieger	301, 317, 5790
Eva	1, 2
Flughafenbahnhof	5790
Grabdenkmal	2, 5790
Haubentaucher	300, 5790

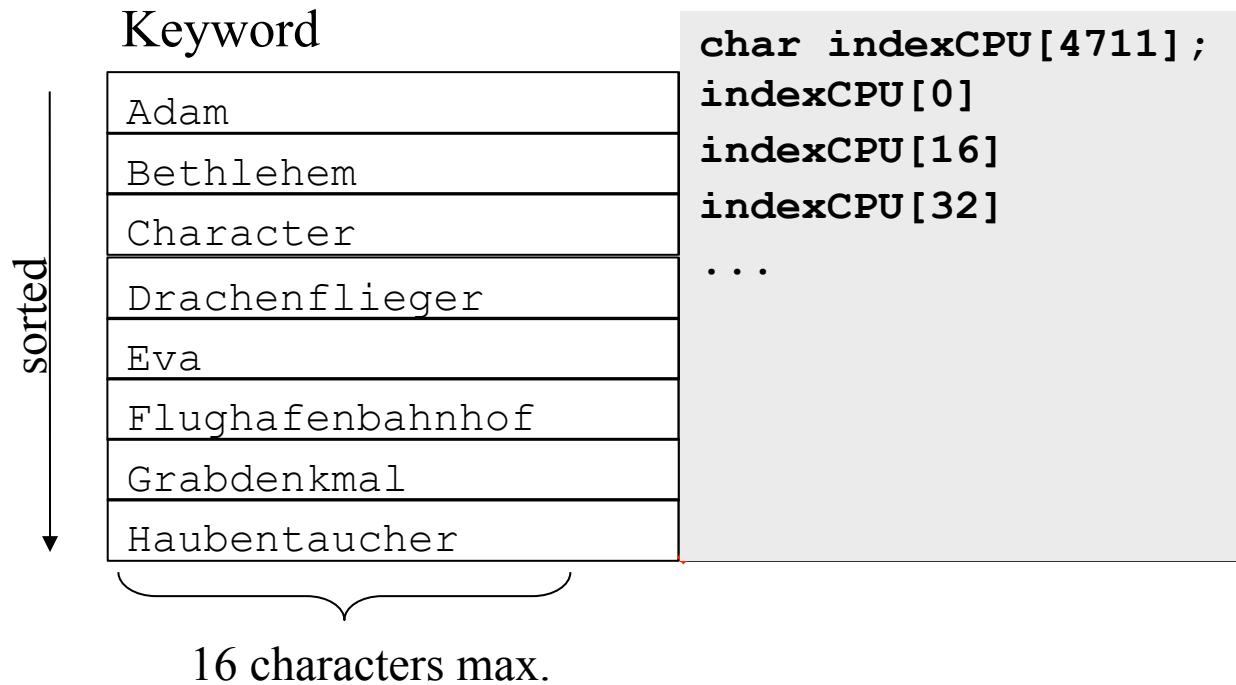
Sorted ↓

16 characters max.

A Simple implementation of (index) search



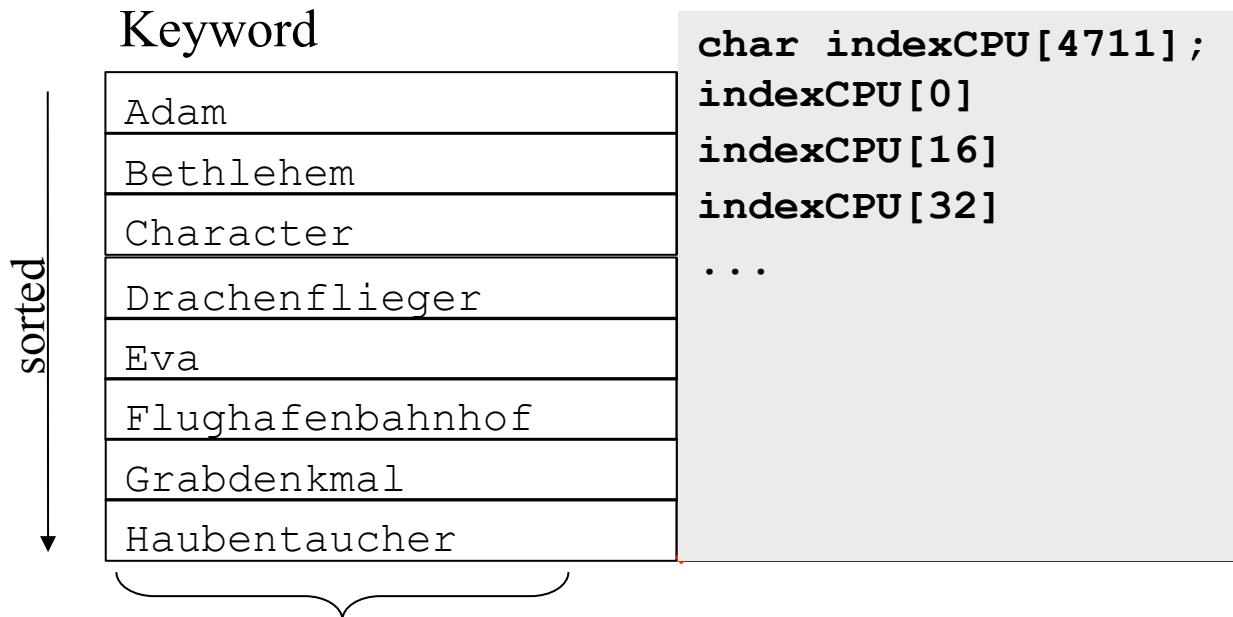
A Simple implementation of (index) search



- On the CPU we use a few library calls and we are done

```
char searchkey[16]= "Flughafenbahnhof";
result = bsearch((void*) searchkey, indexCPU,
                    numentries, sizeof(char)*16,
                    (int(*)(const void*, const void*)) strcmp);
```

A Simple implementation of (index) search



16 characters max.

- On the CPU we use a few library calls and we are done

```
char searchkey[16] = "Flughafenbahnhof";
result = bsearch((void*)searchkey, indexCPU,
                 numentries, sizeof(char)*16,
                 (int(*)(const void*, const void*)) strcmp);
```

- Can we just port a CPU implementation?

A Simple GPU implementation

- Get the data to the GPU

```
char* indexGPU;
char* searchkeysGPU;
char* resultsGPU;
// copy the data
cudaMalloc((void**)&indexGPU, sizeof(char)*wordlength*entries);
cudaMemcpy(indexGPU, indexCPU, sizeof(char)*wordlength*entries,
           CudaMemcpyHostToDevice);
// copy the searchkey(s)
cudaMalloc((void**)&searchkeysGPU, ...
cudaMemcpy(searchkeysGPU, searchkeysCPU,
           sizeof(char)*wordlength*numsearches,
           CudaMemcpyHostToDevice);
// make room for the results
cudaMalloc((void**)&resultsGPU, ...
```

A Simple GPU implementation

- Get the data to the GPU

```
char* indexGPU;
char* searchkeysGPU;
char* resultsGPU;
// copy the data
cudaMalloc((void**)&indexGPU, sizeof(char)*wordlength*entries);
cudaMemcpy(indexGPU, indexCPU, sizeof(char)*wordlength*entries,
           CudaMemcpyHostToDevice);
// copy the searchkey(s)
cudaMalloc((void**)&searchkeysGPU, ...
cudaMemcpy(searchkeysGPU, searchkeysCPU,
           sizeof(char)*wordlength*numsearches,
           CudaMemcpyHostToDevice);
// make room for the results
cudaMalloc((void**)&resultsGPU, ...
```

- Know your hardware (GTX 285, 30 SMs, 8 cores each, 240 cores)
 - Set up an execution configuration & call global function

```
dim3 Dg = dim3(30,0,0);
dim3 Db = dim3(8,0,0);
searchGPU<<<Dg,Db>>>(indexGPU, entries...;
```

A Simple GPU implementation

- The GPU kernel

```
__global__ void searchGPU(char* index, int entries, int wordlength,
                         char* search_keys, int* results) {
    char* res;
    // use block and thread numbers for indexing
    res = bsearch(&search_keys[((blockIdx.x*BLOCK_SIZE)+threadIdx.x)*wordlength],
                  index,
                  entries,
                  wordlength);
    // use block and thread numbers for indexing
    results[(blockIdx.x*BLOCK_SIZE)+threadIdx.x] = (res-data)/
                                                MAX_WORD_LENGTH;
}
```

A Simple GPU implementation

- The GPU kernel

```
__global__ void searchGPU(char* index, int entries, int wordlength,
                         char* search_keys, int* results) {
    char* res;
    // use block and thread numbers for indexing
    res = bsearch(&search_keys[((blockIdx.x*BLOCK_SIZE)+threadIdx.x)
                                *wordlength],
                  index,
                  entries,
                  wordlength);
    // use block and thread numbers for indexing
    results[(blockIdx.x*BLOCK_SIZE)+threadIdx.x] = (res-data)/
                                                MAX_WORD_LENGTH;
}
```

- There is no libc on the GPU =(
 - Just stick __device__ in front of the libc code?
 - “bsearch” is recursive, but there is no recursion on the GPU
- Write a iterative one ...

A Simple GPU binary search

```
__device__ char* bsearchGPU(char *key, char *base, int n, int size){
    char *mid_point;
    int cmp;

    while (n > 0) {
        mid_point = (char *)base + size * (n >> 1);
        if ((cmp = strcmpGPU(key, mid_point)) == 0)
            return (char *)mid_point;
        if (cmp > 0) {
            base = (char *)mid_point + size;
            n = (n - 1) >> 1;
        } // cmp < 0
        else n >>= 1;
    }
    return (char *)NULL;
}
```

- Still need strcmp

A Simple GPU binary search

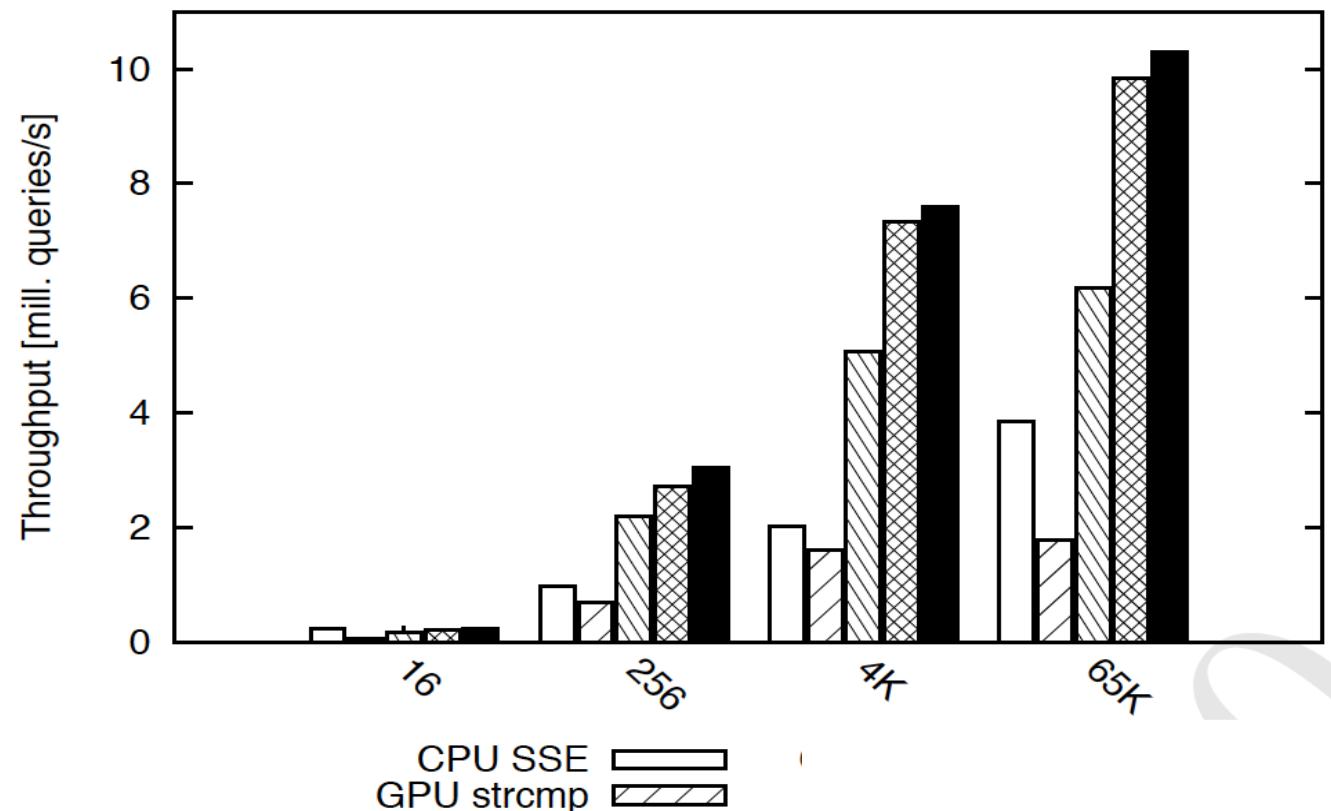
```
__device__ char* bsearchGPU(char *key, char *base, int n, int size){  
    char *mid_point;  
    int cmp;  
  
    while (n > 0) {  
        mid_point = (char *)base + size * (n >> 1);  
        if ((cmp = strcmpGPU(key, mid_point)) == 0)  
            return (char *)mid_point;  
        if (cmp > 0) {  
            base = (char *)mid_point + size;  
            n = (n - 1) >> 1;  
        } // cmp < 0  
        else n >>= 1;  
    }  
    return (char *)NULL;  
}
```

- Still need strcmp
- Again, stick `__device__` in front of the libc code

```
__device__ int strcmpGPU(char* s1, char* s2){  
    while (*s1 == *s2++)  
        if (*s1++ == 0) return 0;  
    return (*s1 - *(s2 - 1));  
}
```

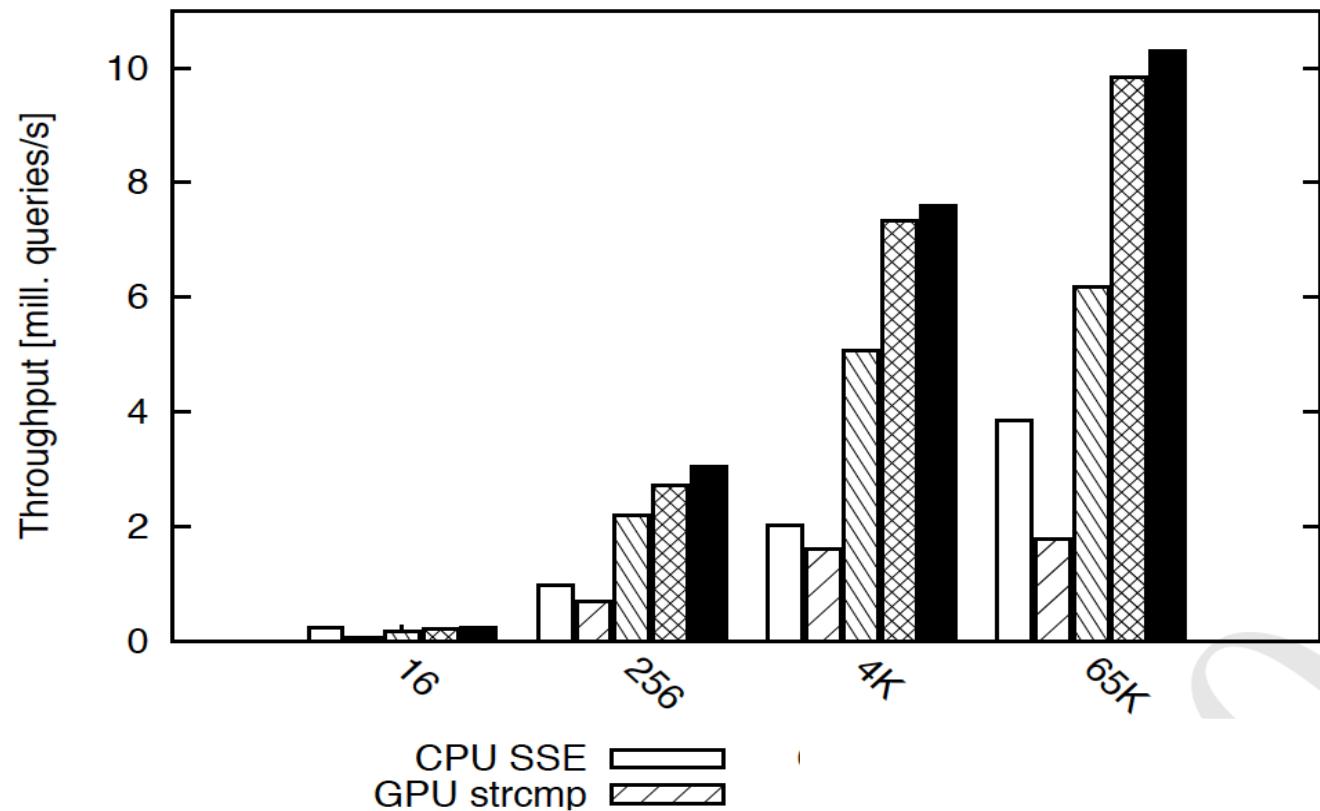
Binary Search on the GPU

- Searching a large data set (512MB) with 33 million (225) 16-character strings



Binary Search on the GPU – Why is it slow?

- Searching a large data set (512MB) with 33 million (225) 16-character strings

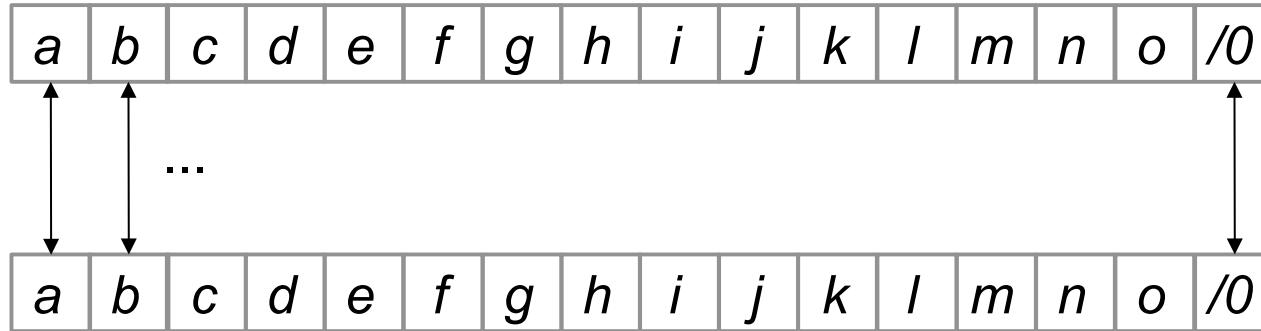


- It's slower than a CPU implementation for all data set sizes!
 - Let's try some optimizations ...

Search requires to compare

- Search naturally requires MANY comparisons
- The strcmp() library function:

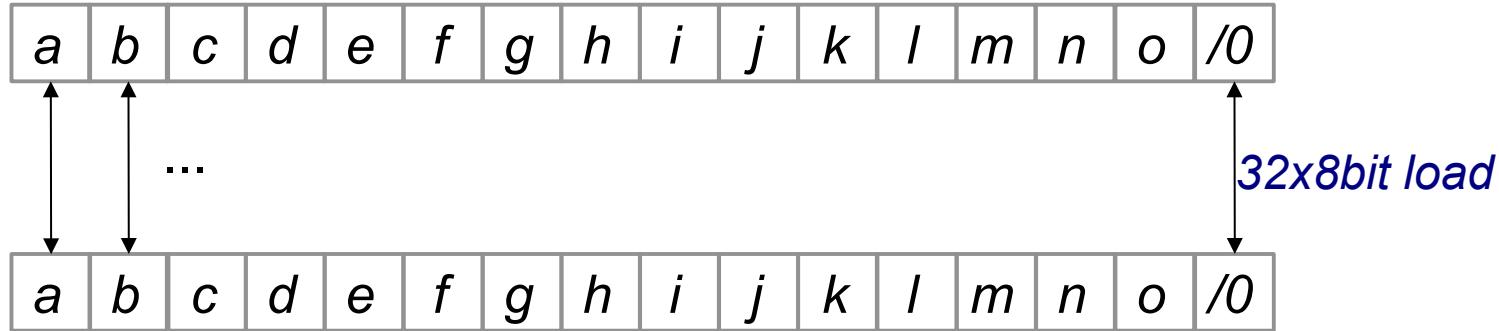
```
int strcmp(const char* s1, const char* s2) {
    while (*s1 == *s2++)
        if (*s1++ == 0) return 0;
    return (*s1 - *(s2 - 1));
}
```



Search requires to compare

- Search naturally requires MANY comparisons
- The strcmp() library function:

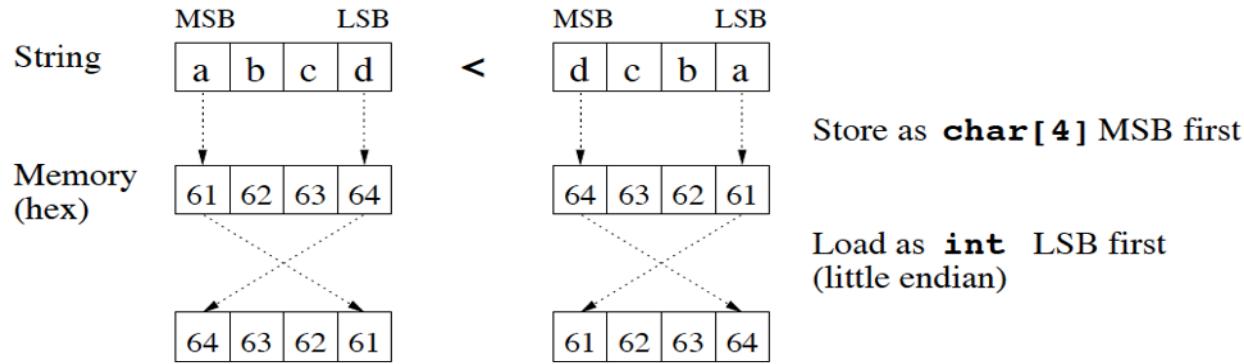
```
int strcmp(const char* s1, const char* s2){  
    while (*s1 == *s2++)  
        if (*s1++ == 0) return 0;  
    return (*s1 - *(s2 - 1));  
}
```



- Byte-wise memory access is known to be slow

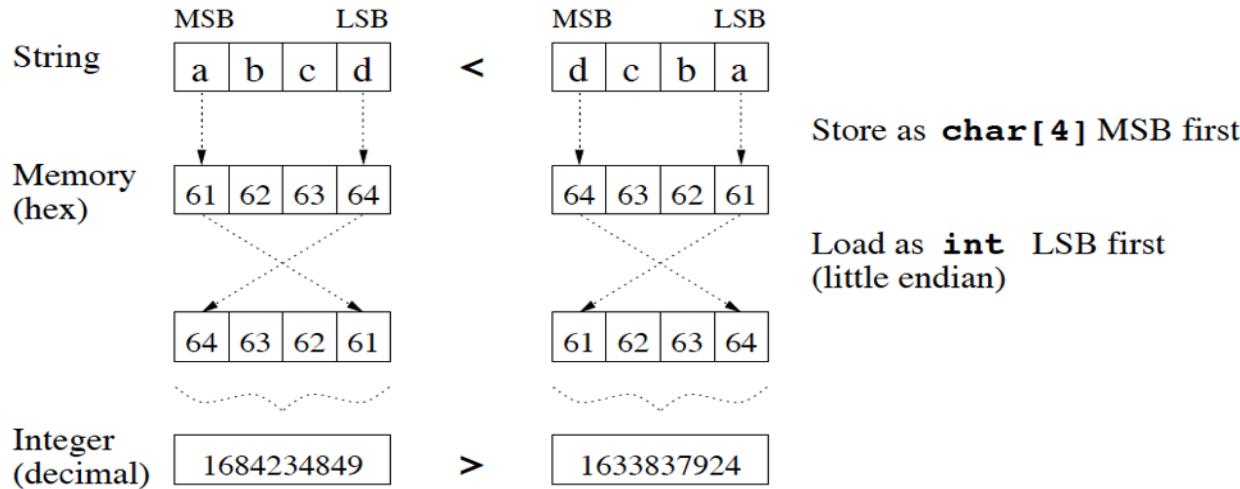
Optimizing compare operations

- How about vector string comparison, a la SSE?
- No Byte vectors on the GPU ... but Integer vectors



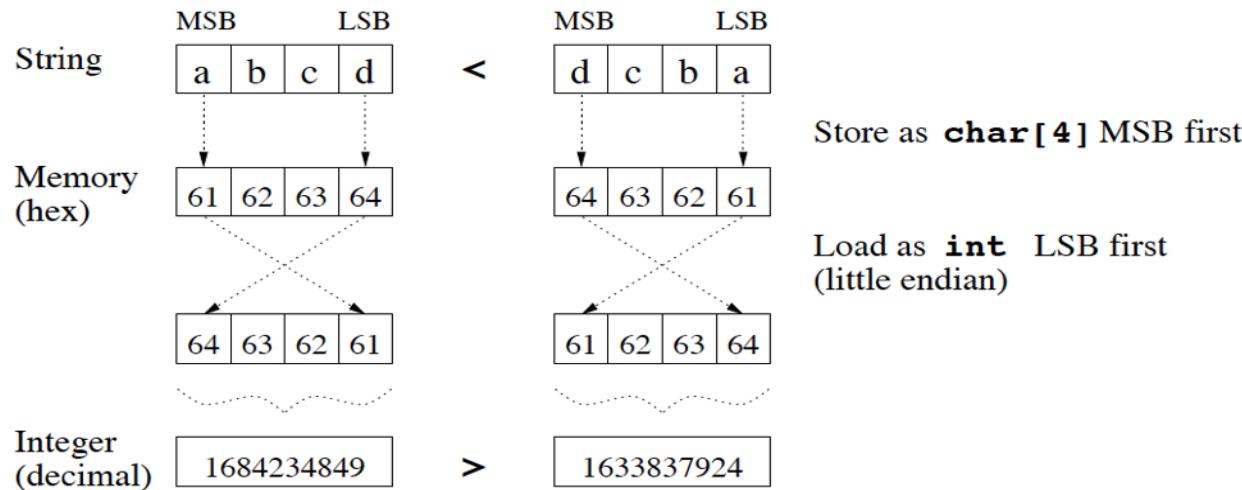
Optimizing compare operations

- How about vector string comparison, a la SSE?
- No Byte vectors on the GPU ... but Integer vectors



Optimizing compare operations

- How about vector string comparison, a la SSE?
- No Byte vectors on the GPU ... but Integer vectors



- Loading character strings as int changes endianness
- CPU has bswap, on the GPU we have to write it:

```
#define BSWP( x ) ; \
temp = ( x ) << 24 ; \
temp = temp | ( ( ( x ) << 8) & 0x00FF0000 ) ; \
temp = temp | ( ( ( unsigned ) ( x ) >> 8) & 0x0000FF00 ) ; \
x = temp | ( ( unsigned ) ( x ) >> 24 ) ;
```

Optimizing compare operations

- Comparing integer vectors (bswap for <> skipped for clarity)

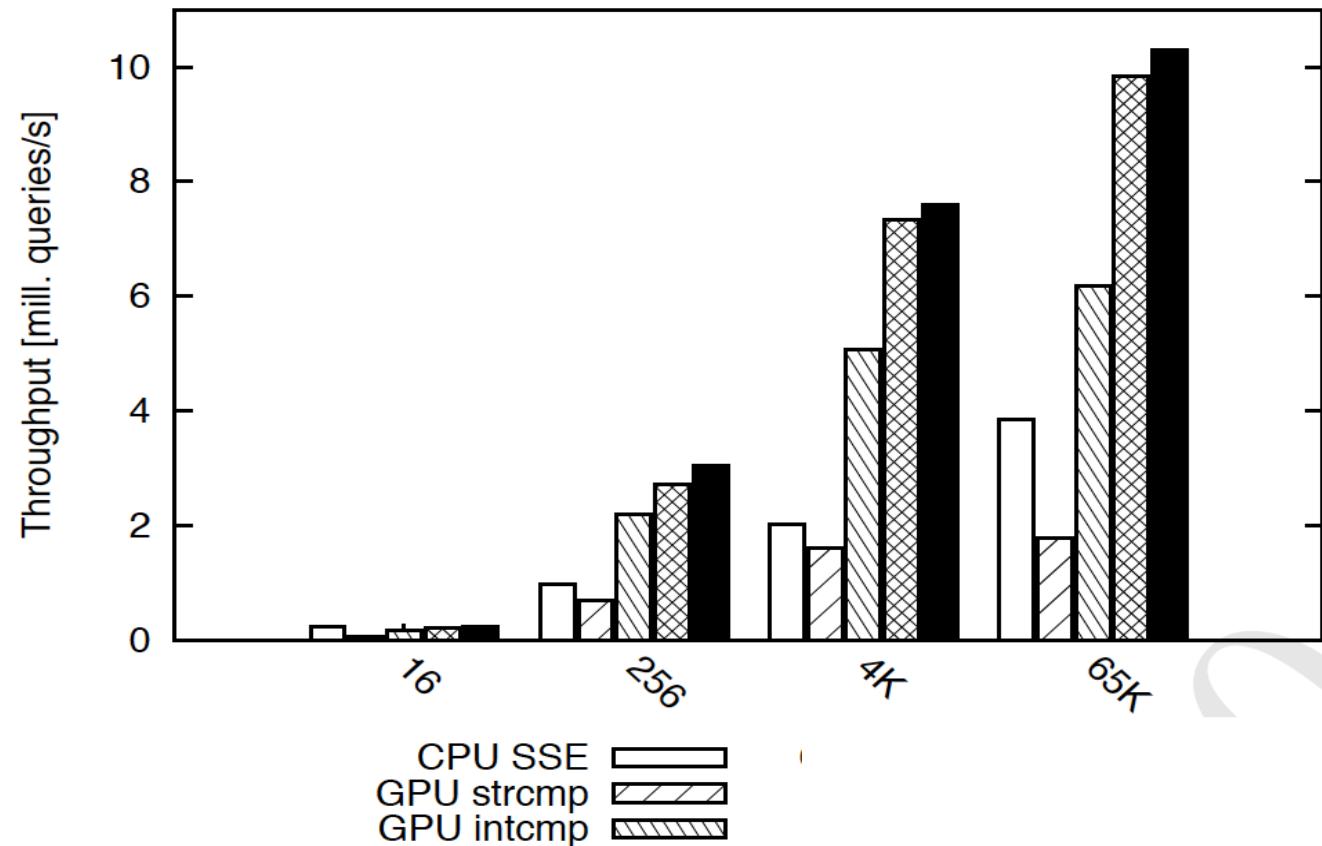
```
__device__ int intcmp(uint4* a, uint4* b) {

    int r =1;
    if ((*a).x < (*b).x)
        r=-1;
    else if ((*a).x == (*b).x) {
        if ((*a).y < (*b).y)
            r=-1;
        else if ((*a).y == (*b).y) {
            if ((*a).z < (*b).z)
                r=-1;
            else if ((*a).z == (*b).z) {
                if ((*a).w < (*b).w)
                    r=-1;
                else if ((*a).w == (*b).w)
                    r=0;
            }
        }
    }
    return r;
}
```

- Still dereferencing 16 memory pointers ...

Binary Search on the GPU – Why is it slow?

- Searching a large data set (512MB) with 33 million (225) 16-character strings



- With intcmp it's only marginally faster than a CPU implementation
- We still do pointer chasing, i.e. roundtrips to memory ...

Reducing global memory access

- Intcmp is memory latency sensitive

Processor	L1 [cyc]	L2 [cyc]	L3 [cyc]	mem [cyc]
Intel Core i7 2.6GHz	4	10	40	350
nVidia GT200b 1.5 GHz	4	n/a	n/a	500

- We can use shared memory like L1

x 16 for each comparison !!!

Reducing global memory access

- Intcmp is memory latency sensitive

Processor	L1 [cyc]	L2 [cyc]	L3 [cyc]	mem [cyc]
Intel Core i7 2.6GHz	4	10	40	350
nVidia GT200b 1.5 GHz	4	n/a	n/a	500

- We can use shared memory like L1

x 16 for each comparison !!!

```

__shared__ uint4 cache[NUM_THREADS*2];

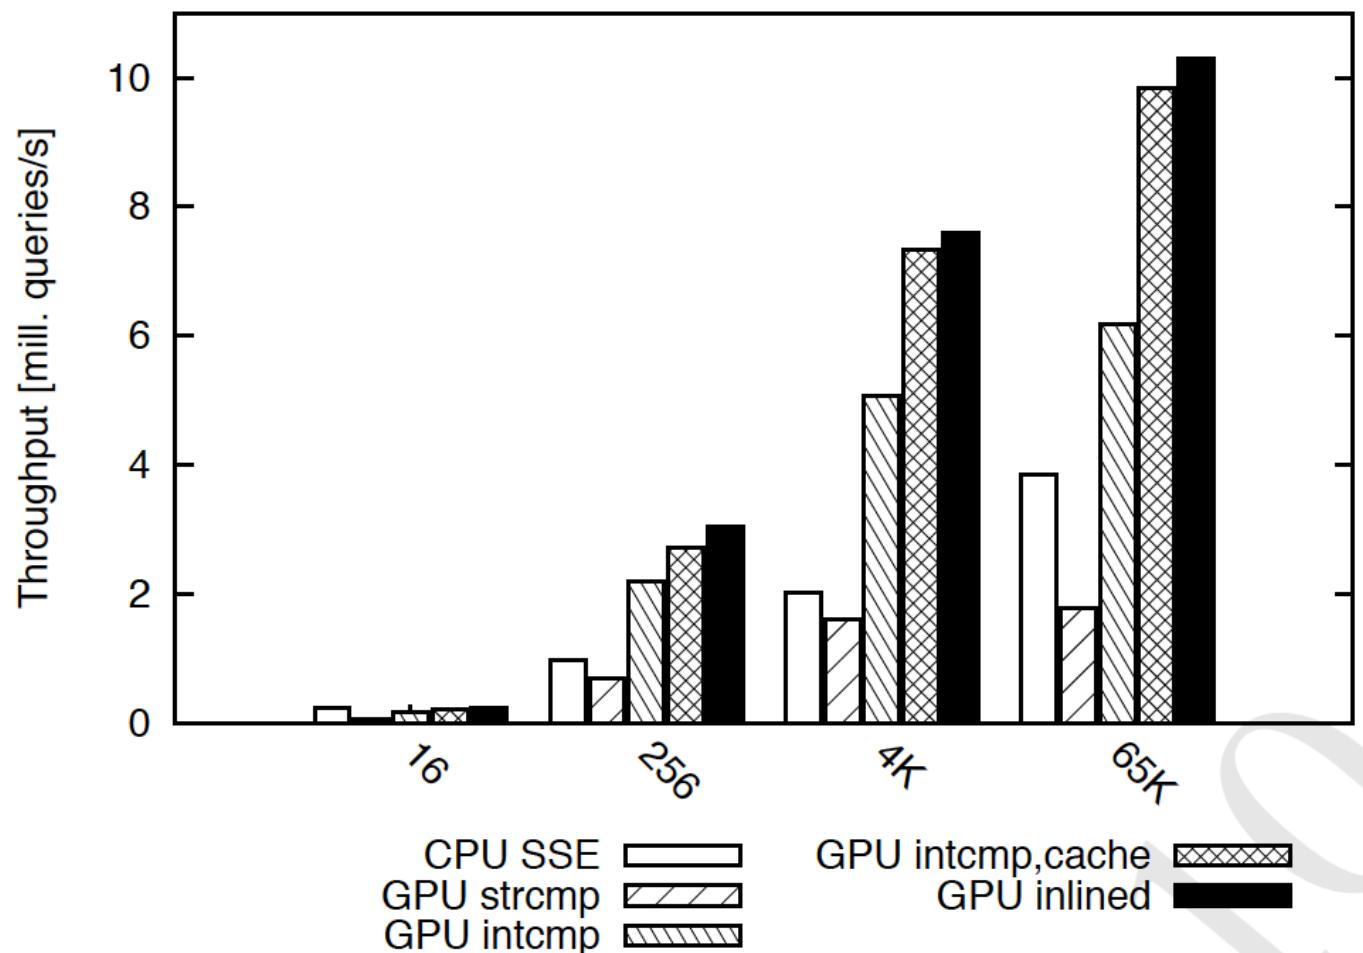
__device__ uint4* bsearchGPU( uint4 *key,  uint4 *base,
                           size_t nmemb,   size_t size)
{
    uint4 *mid_point;
    int cmp;
    cache[threadIdx.x*2] = *key;

    while (nmemb > 0) {
        mid_point = (uint4 *)base + size * (nmemb >> 1);
        cache[threadIdx.x*2+1] = *mid_point;
        if ((cmp = intcmp(&cache[threadIdx.x*2],
                          &cache[threadIdx.x*2+1])) == 0)
            return (uint4 *)mid_point;
    }
}

```

Binary Search on the GPU – optimized

- Searching a large data set (512MB) with 33 million (225) 16-character strings



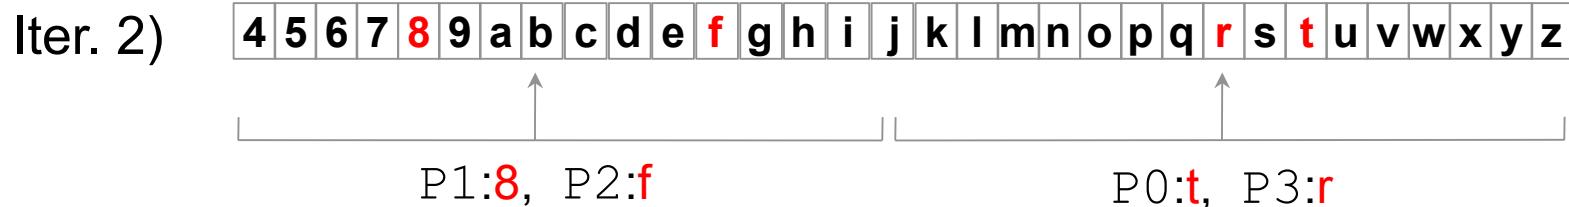
Multi-threaded Binary Search – Example

- Index: a sorted char array 32 entries
- 4 queries: t , 8 , f , r
- 4 processor cores: P1 – P4
- 1 processor core – 1 search: P0:t , P1:8 , P2:f , P3:r
- Theoretical worst-case execution time: $\log_2(32)=5$

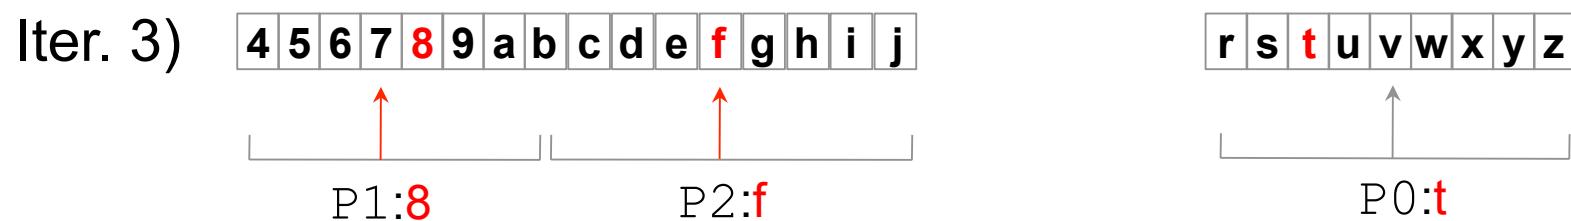
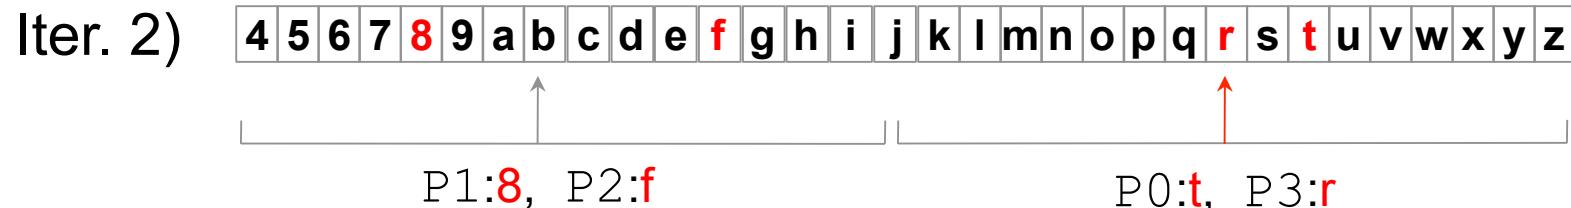
4	5	6	7	8	9	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Multi-threaded Binary Search – Example

- Index: a sorted char array 32 entries
- 4 queries: t , 8 , f , r
- 4 processor cores: P1–P4
- 1 processor core – 1 search: P0:t , P1:8 , P2:f , P3:r
- Theoretical worst-case execution time: $\log_2(32)=5$



Multi-threaded Binary Search – Example



Conventional multi-threading – Analysis

- 100% utilization requires $\# \text{cores}$ concurrent queries
- Queries finishing early
→ utilization < 100%
- Memory access collisions
→ serialized memory access
- $\#\text{memory accesses } \log_2(n)$
- More threads
→ more results
→ response time likely to be worst case: $\log_2(n)$



Can we improve the worst case?

Questions?

Appendix: CPU caches

Caches – the good

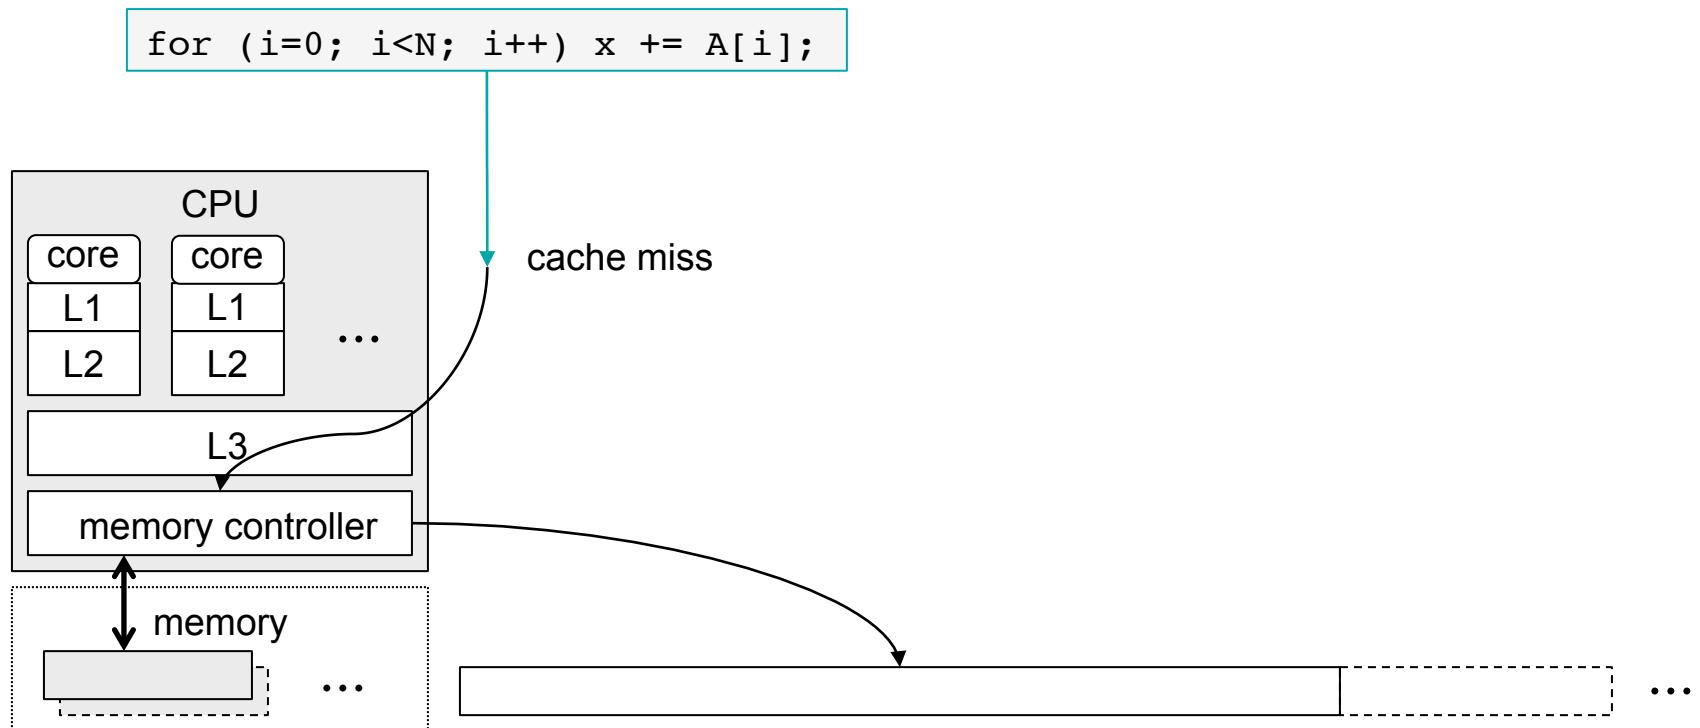
- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Sequential memory access:

```
for (i=0; i<N; i++) x += A[i];
```



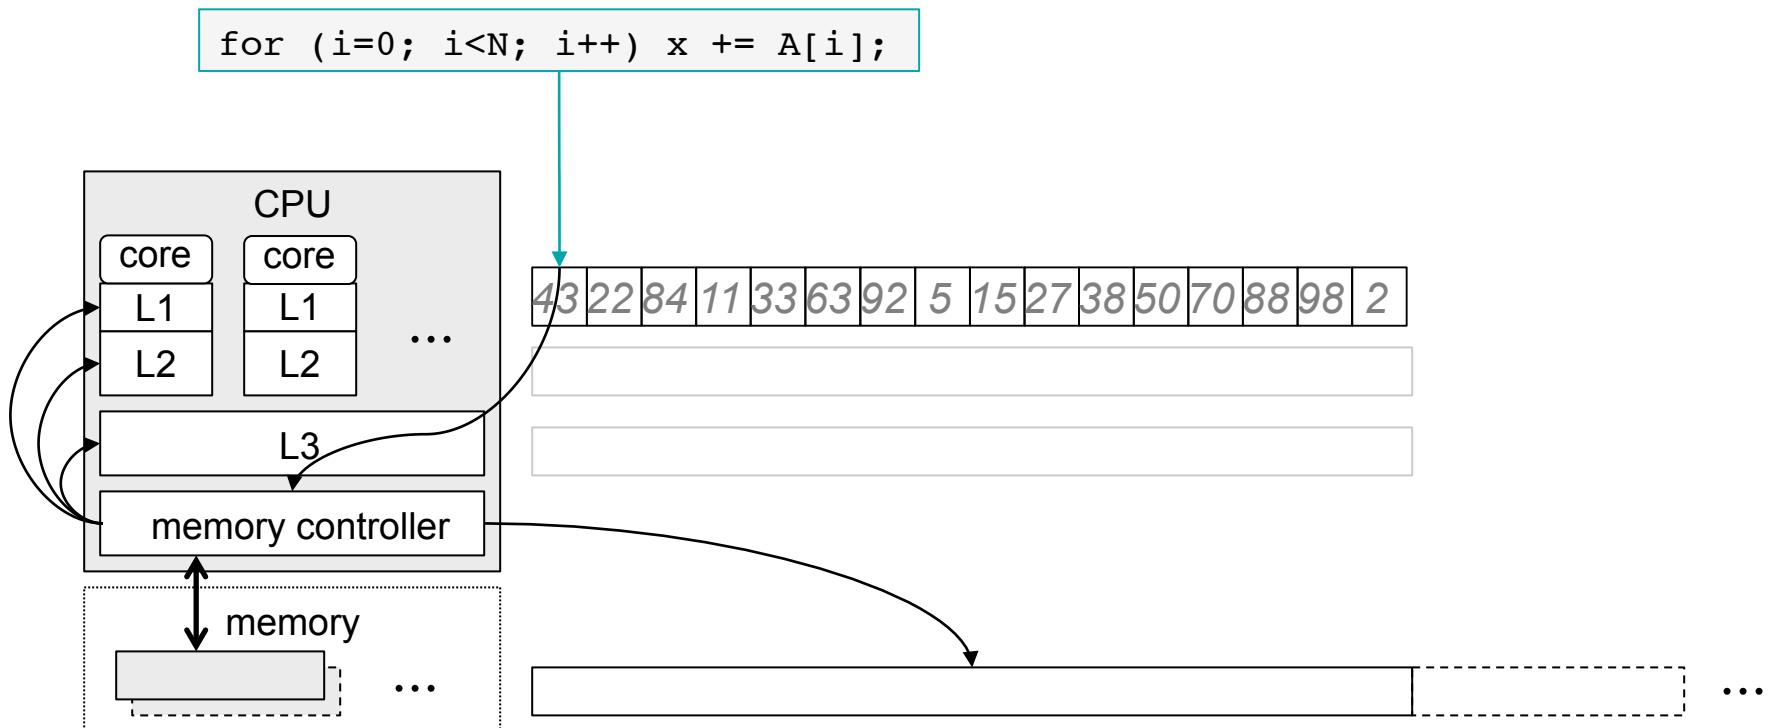
Caches – the good

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Sequential memory access:



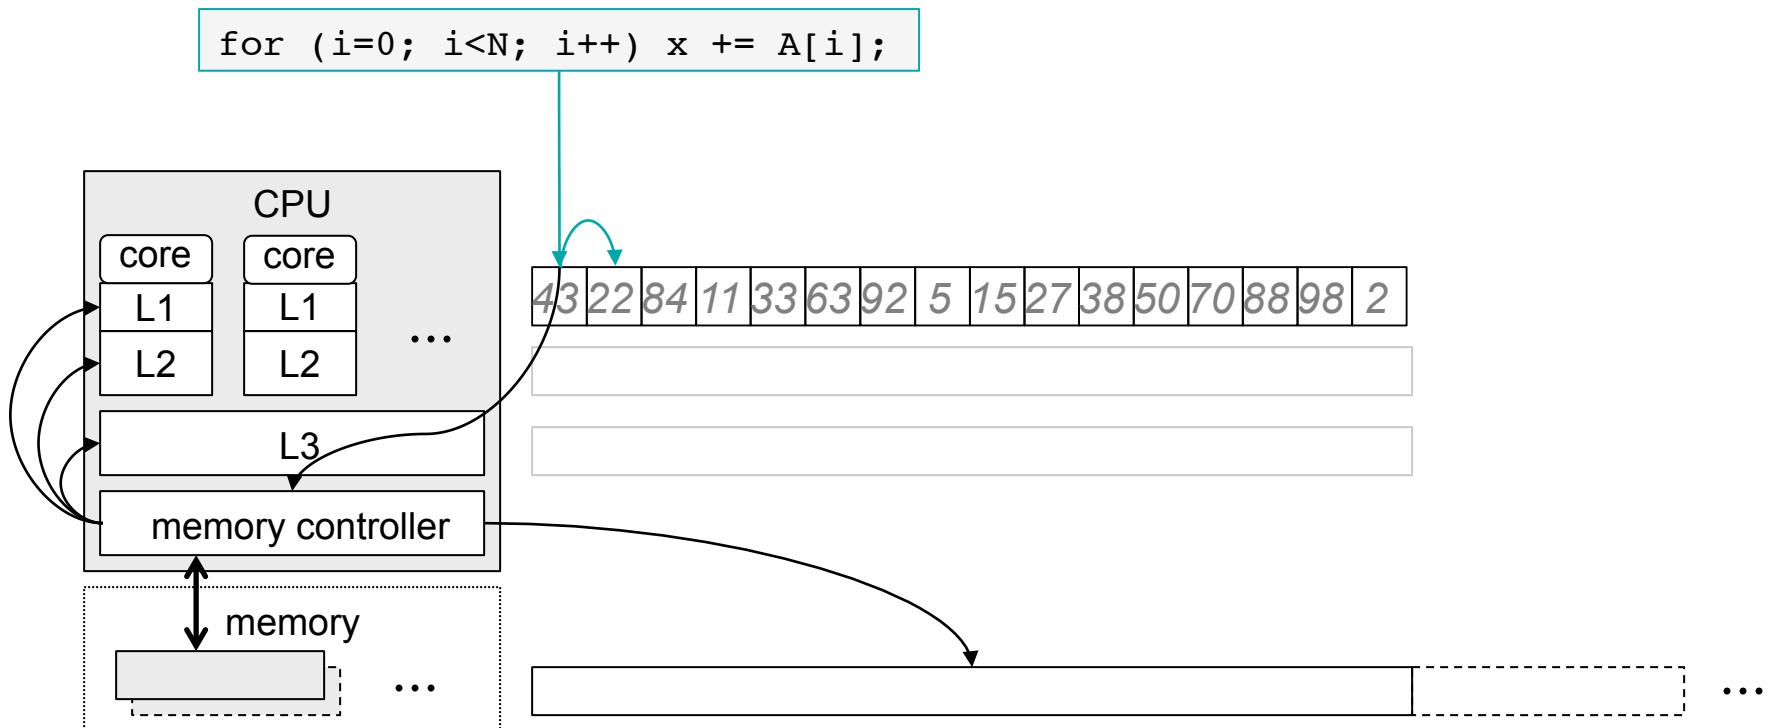
Caches – the good

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Sequential memory access:



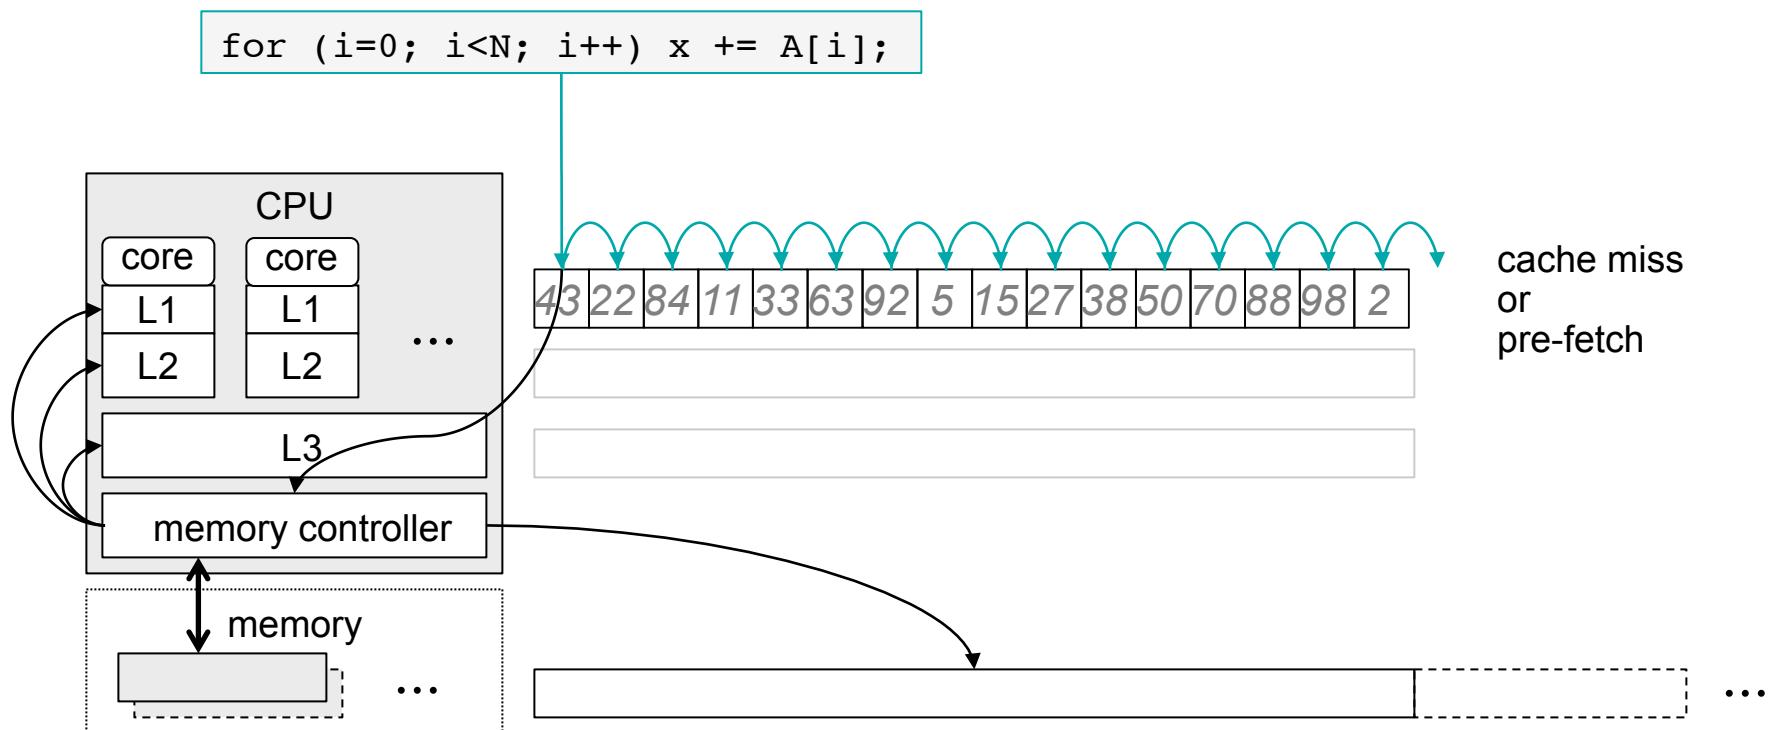
Caches – the good

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Sequential memory access:



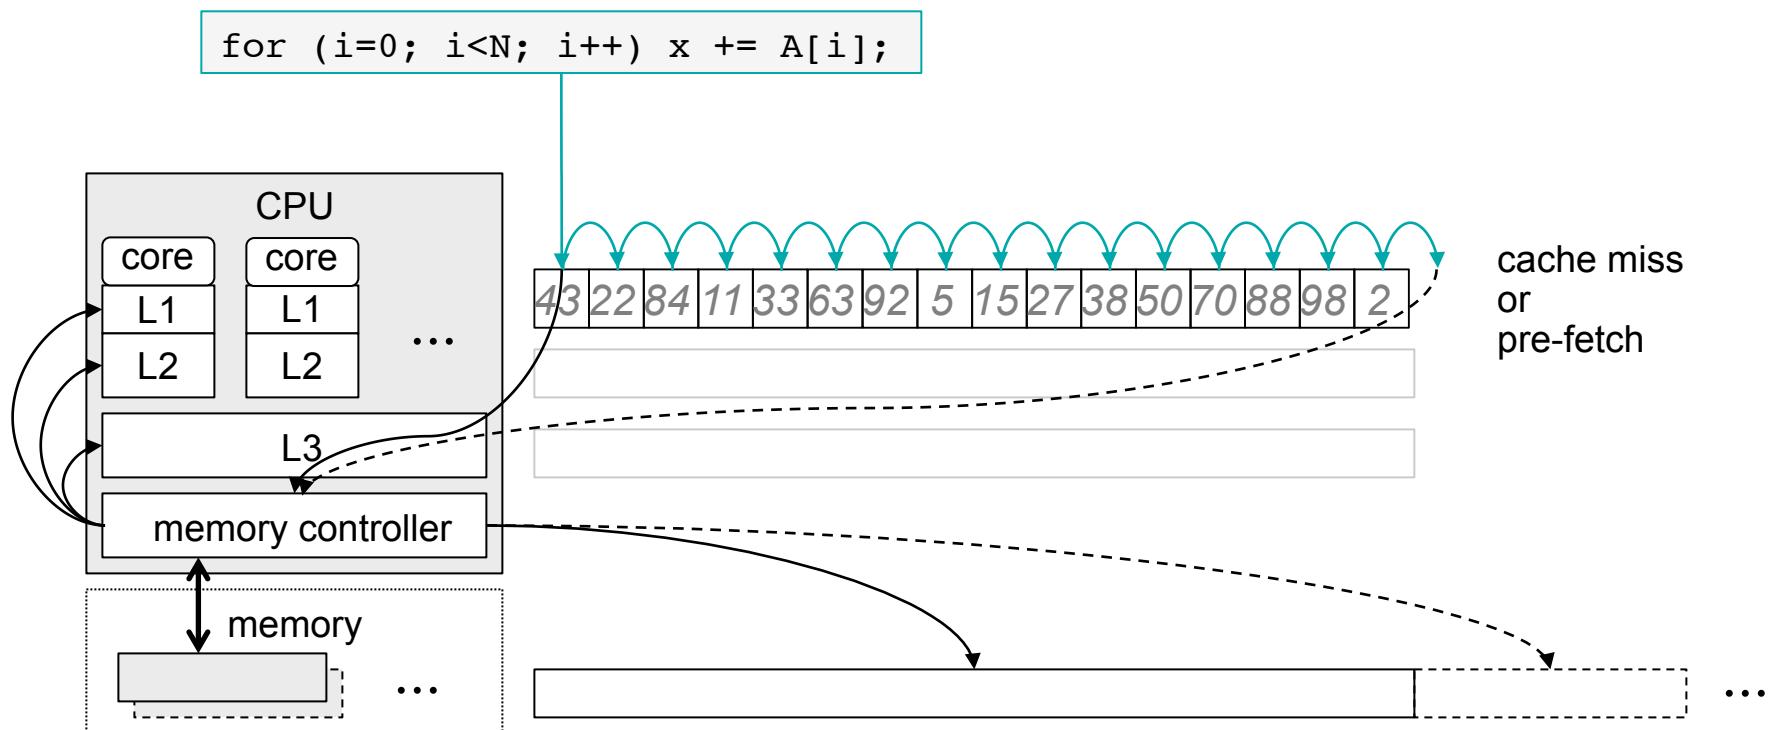
Caches – the good

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Sequential memory access:



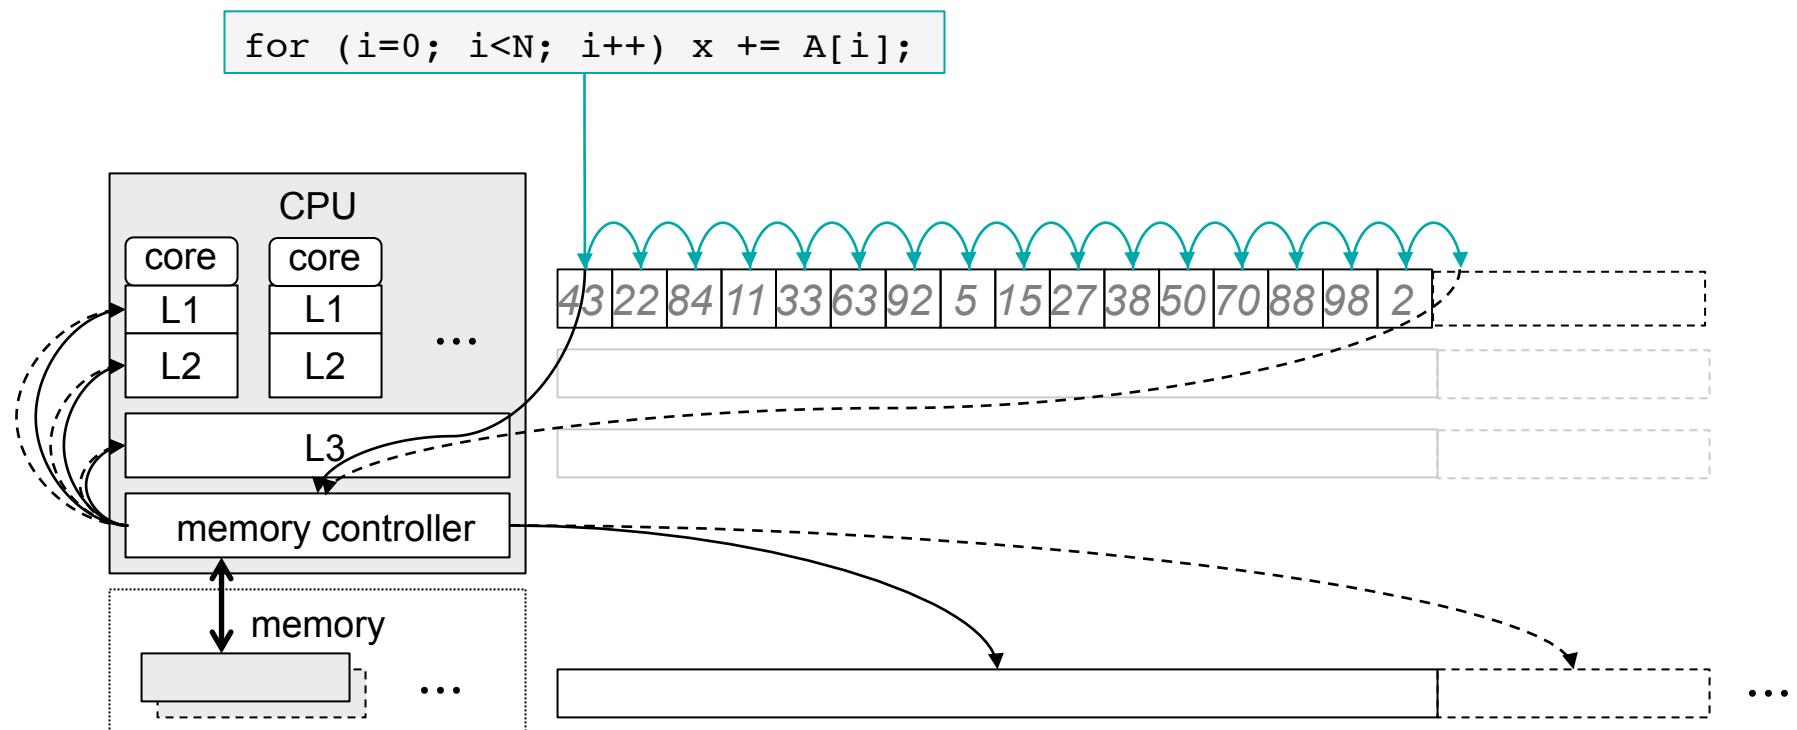
Caches – the good

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Sequential memory access:



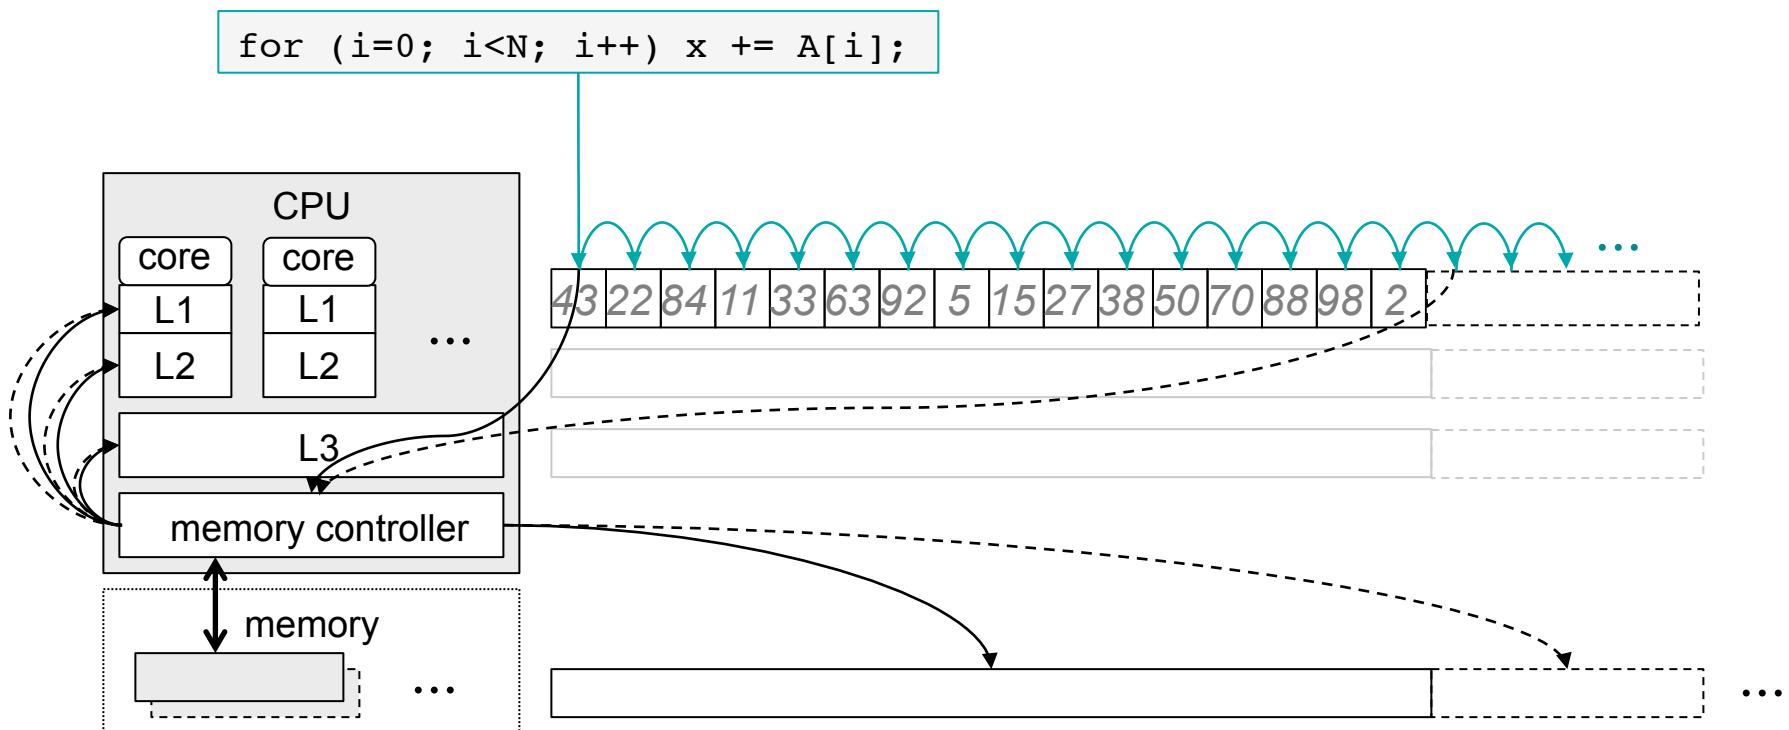
Caches – the good

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Sequential memory access:



Caches – the good

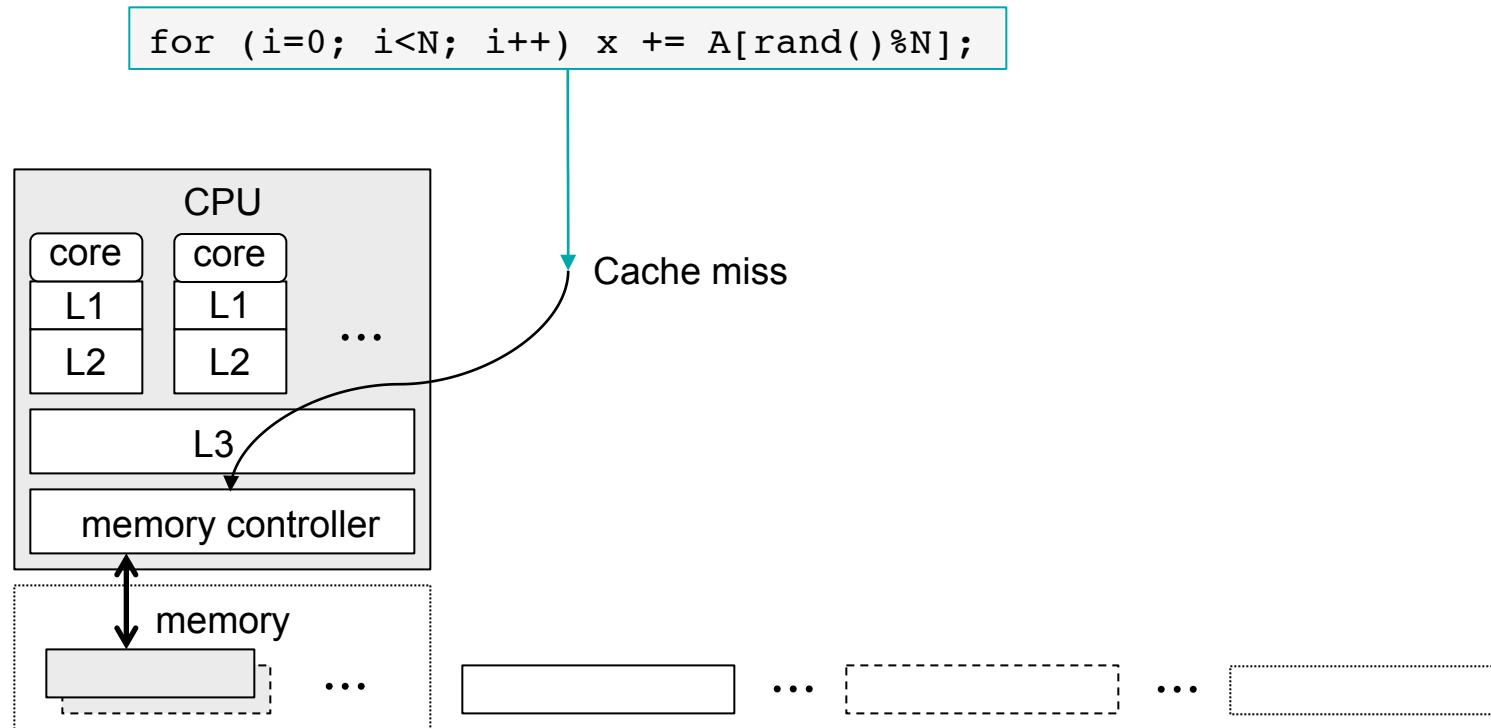
- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Sequential memory access:



→ linear memory access maximizes cache & bandwidth utilization

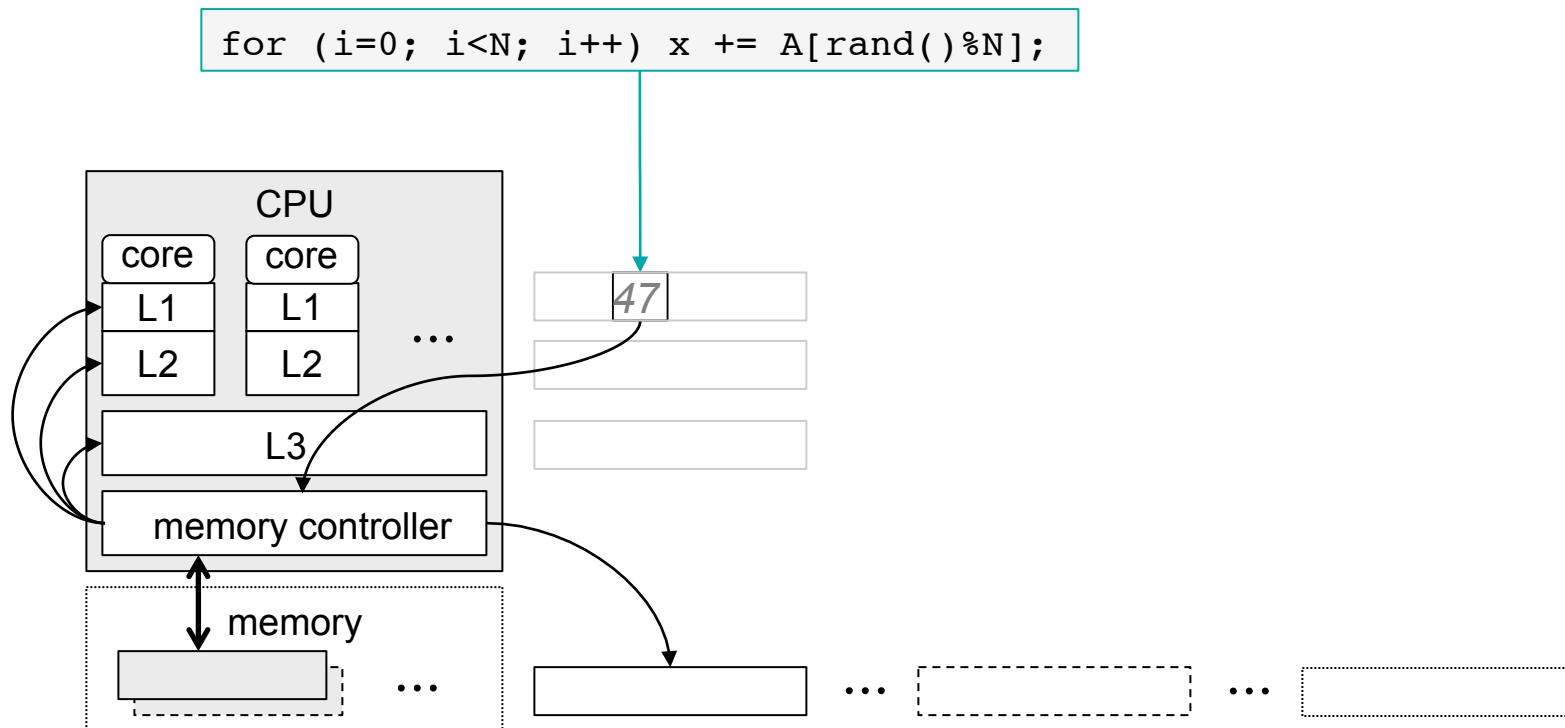
Caches – the bad

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Random memory access:



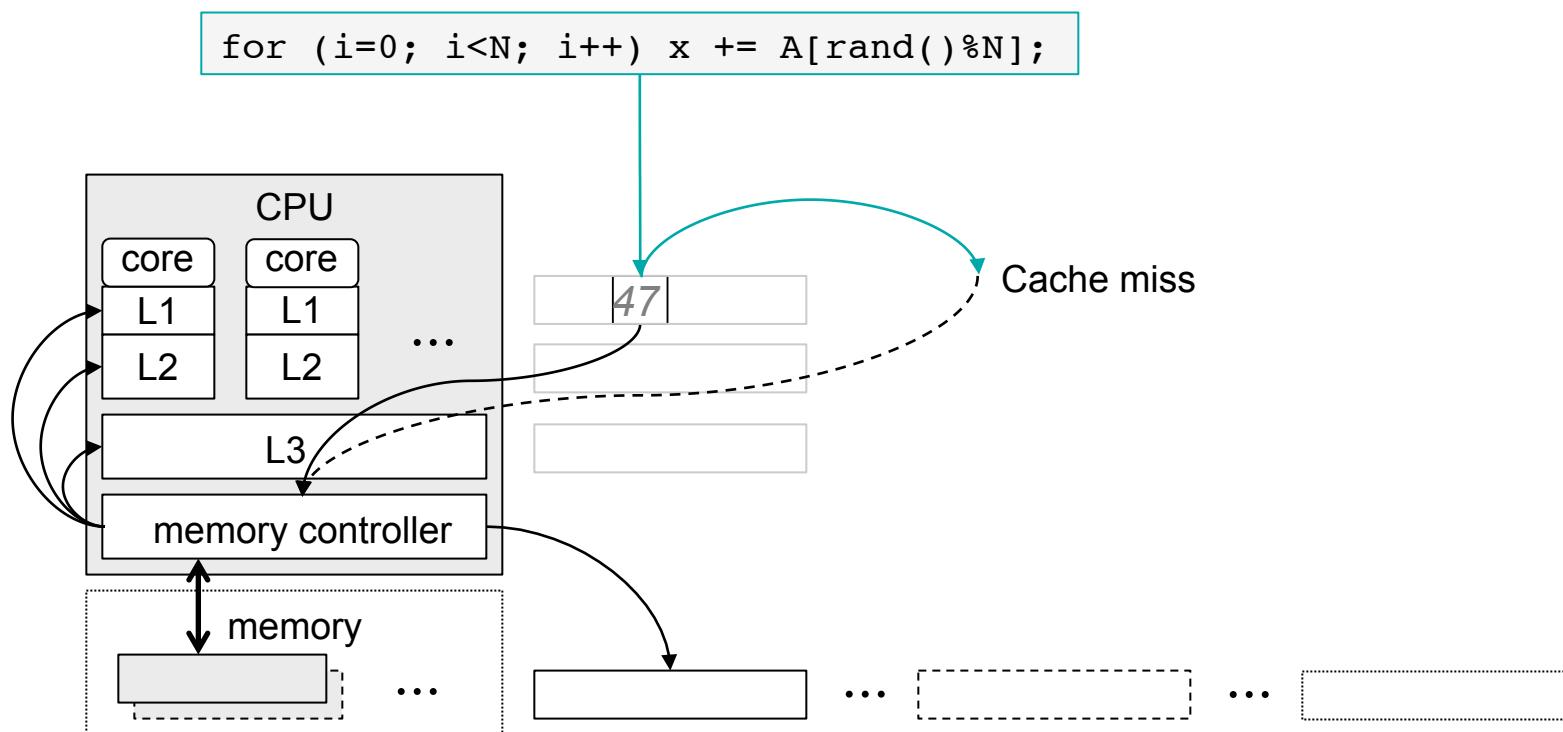
Caches – the bad

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Random memory access:



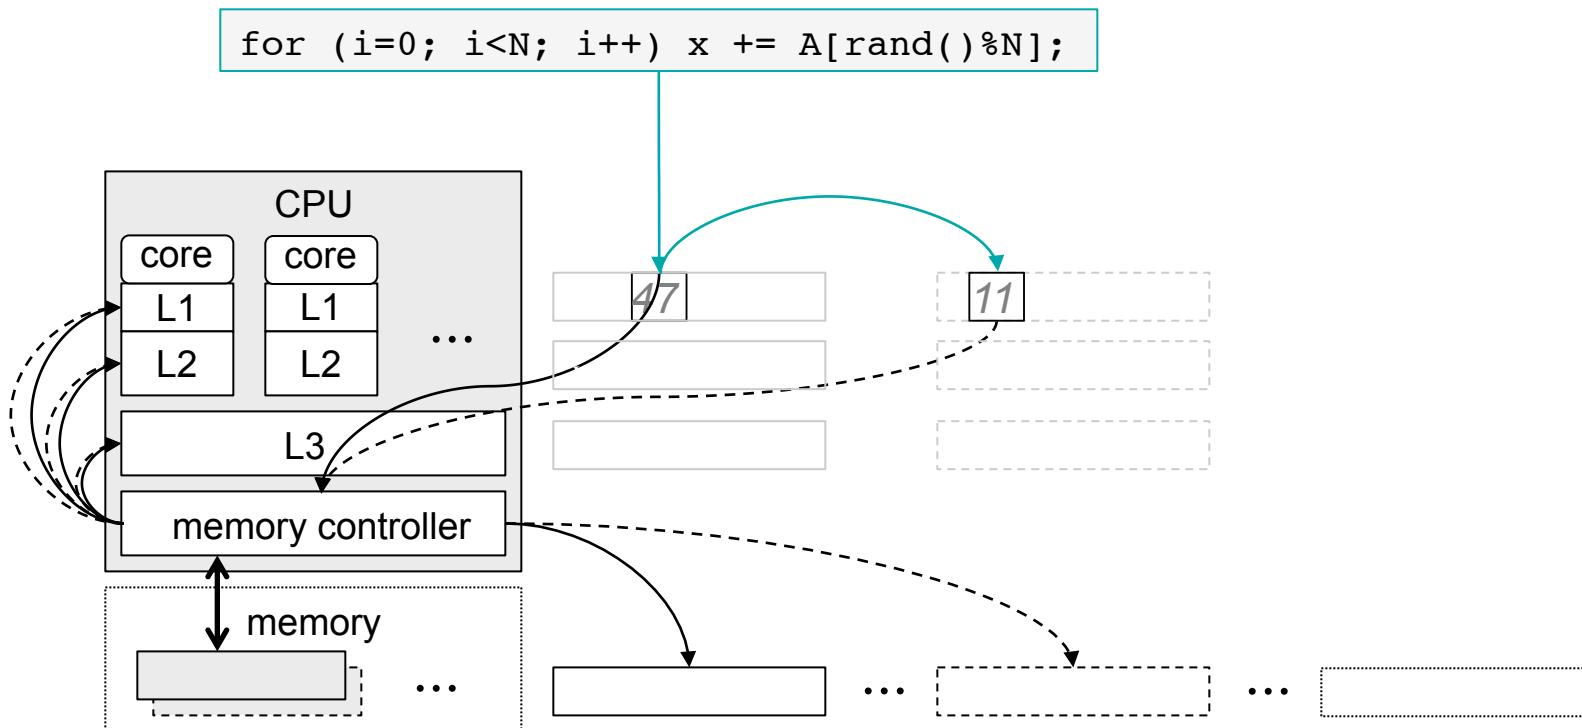
Caches – the bad

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Random memory access:



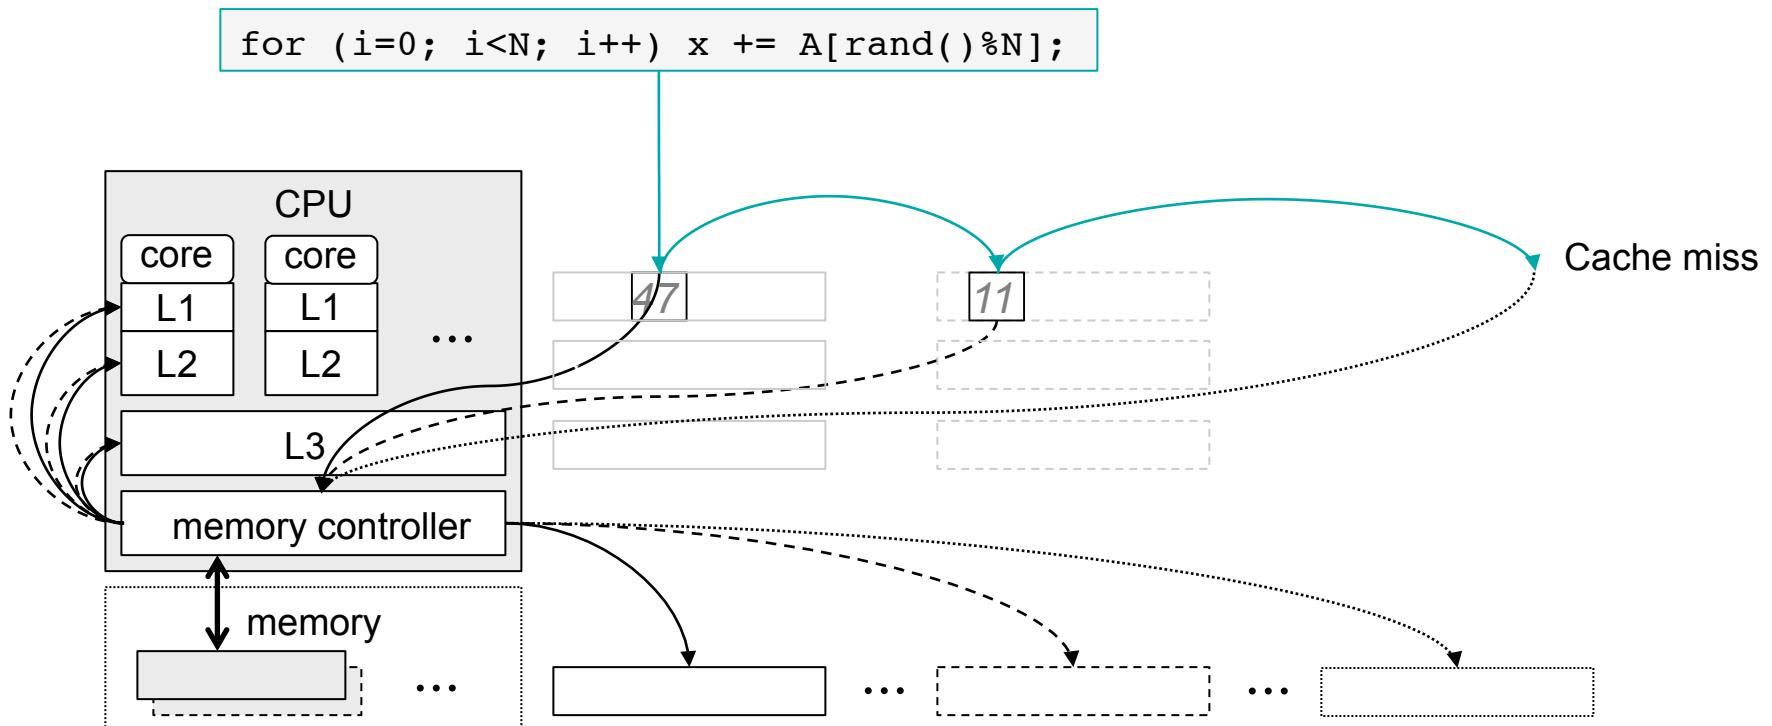
Caches – the bad

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Random memory access:



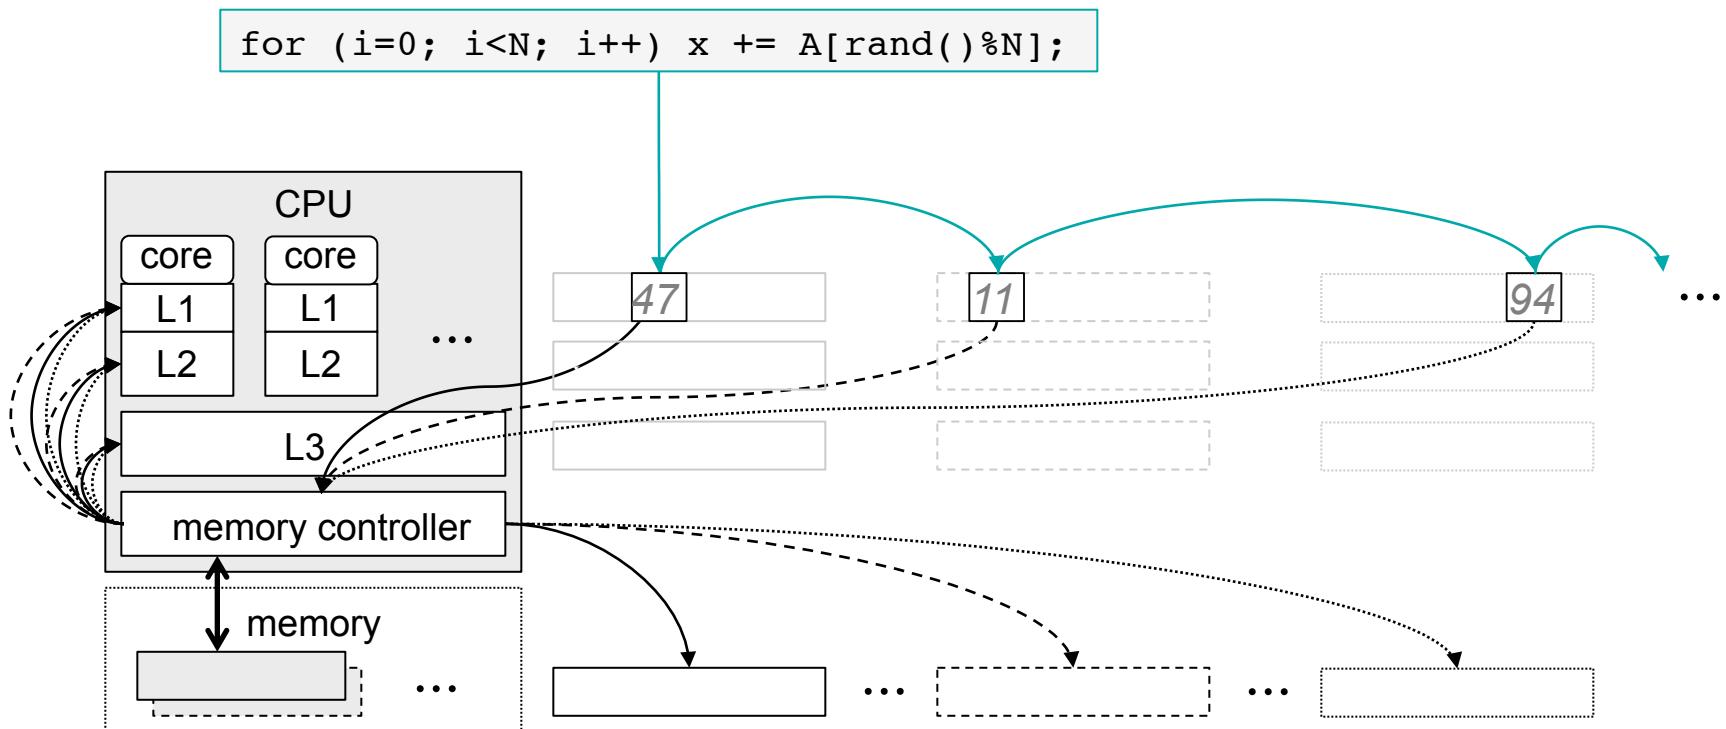
Caches – the bad

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Random memory access:



Caches – the bad

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Random memory access:



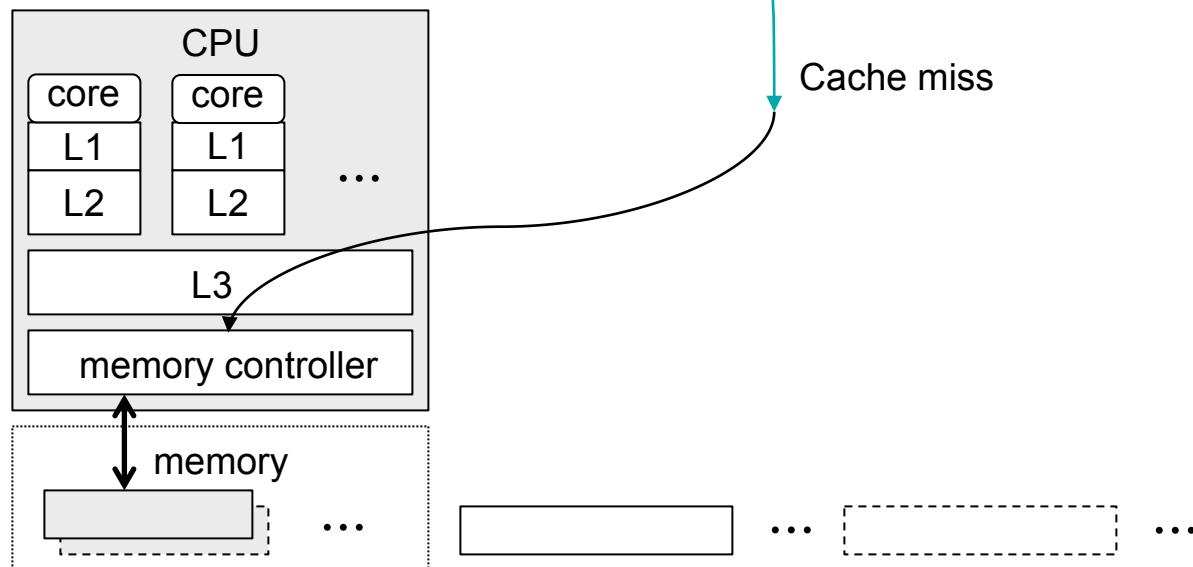
→ Random memory access wastes up to 98.5%* of bandwidth

* assuming we only need 1 Byte of a 64 Byte cache line

Caches – the ugly

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Writes:

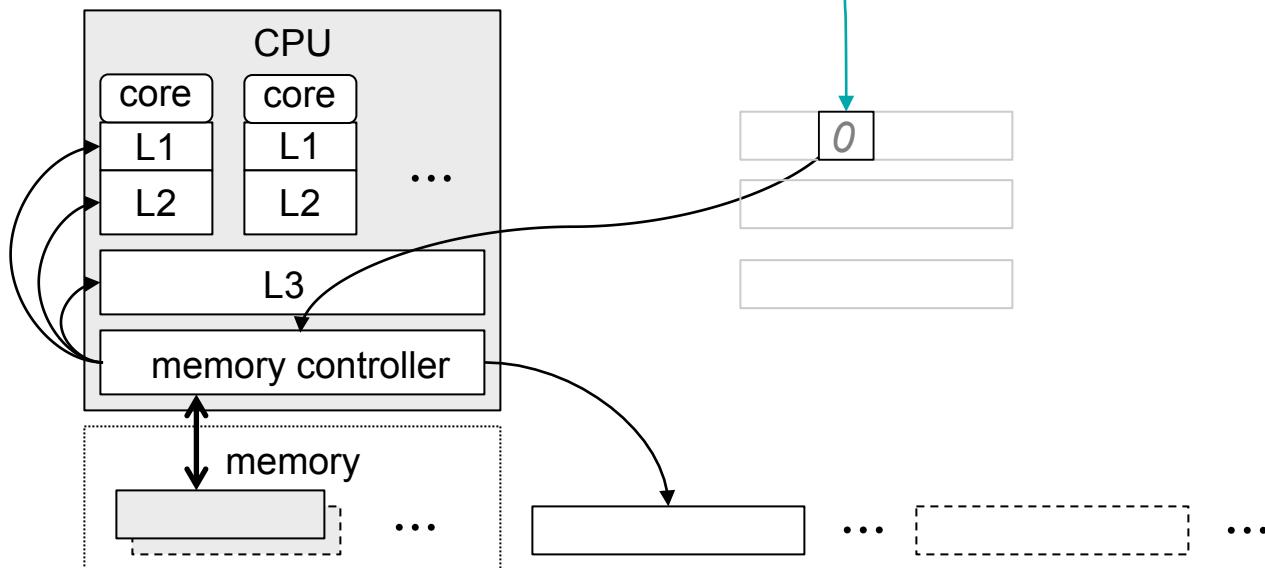
```
for (i=0; i<N; i++) A[rand()%N] = i;
```



Caches – the ugly

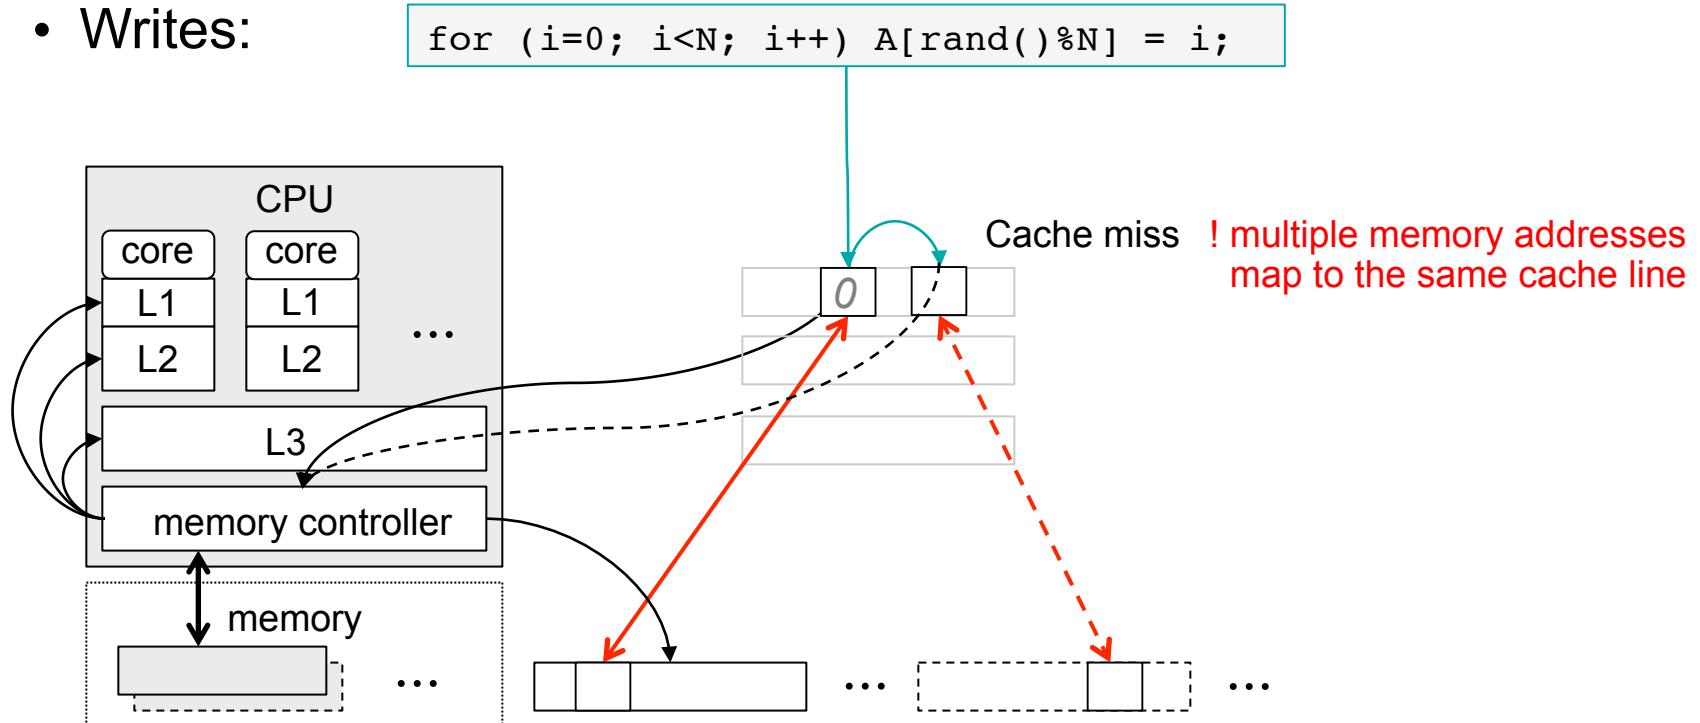
- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Writes:

```
for (i=0; i<N; i++) A[rand()%N] = i;
```



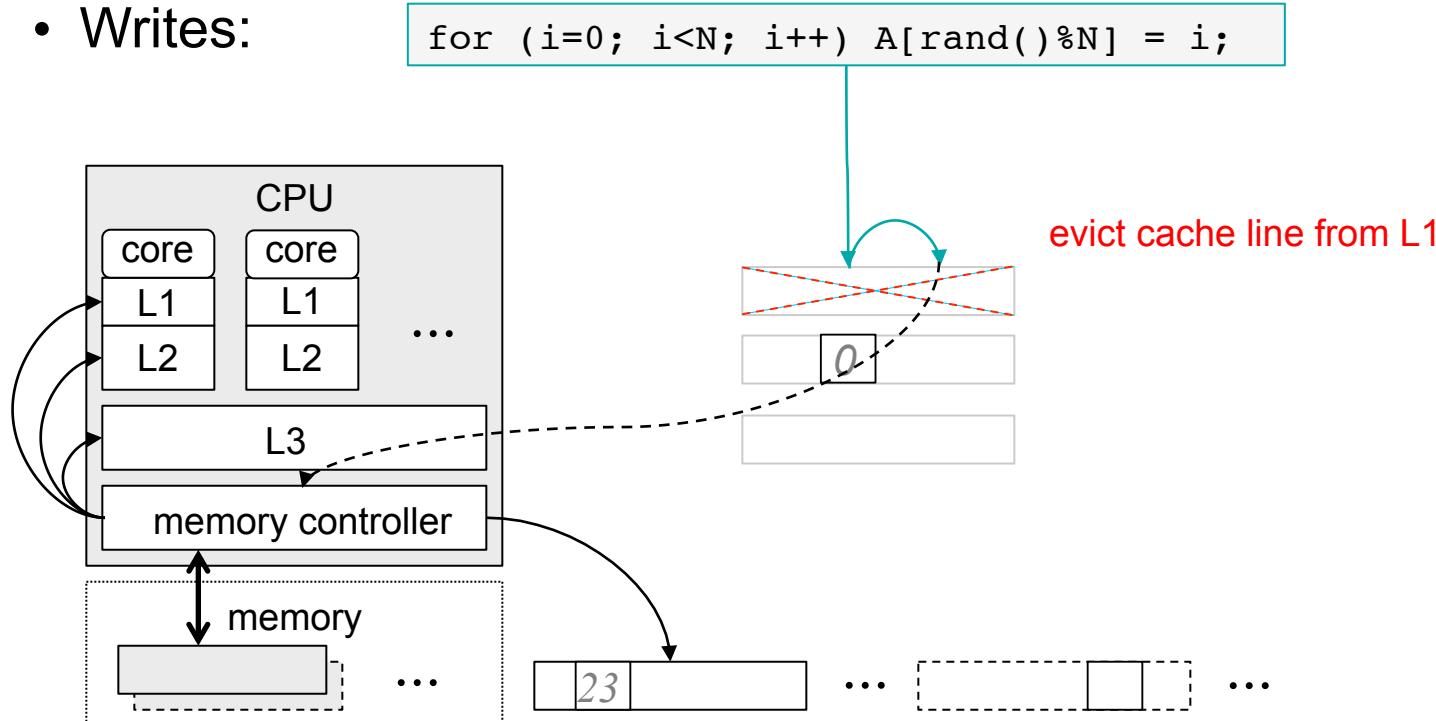
Caches – the ugly

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Writes:



Caches – the ugly

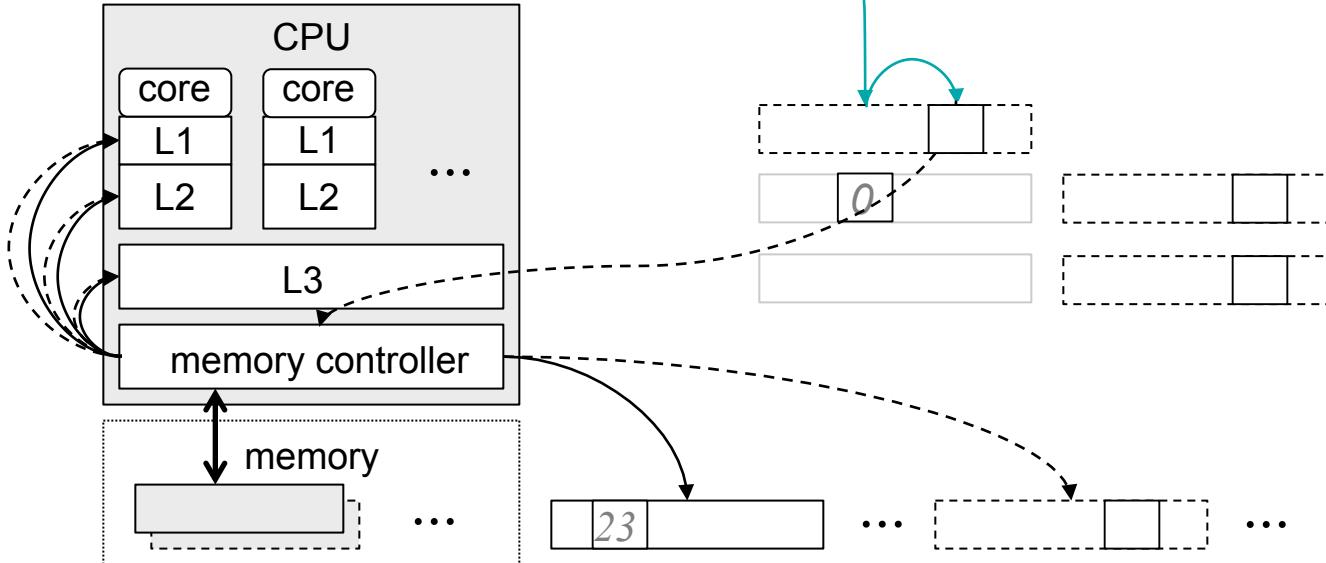
- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Writes:



Caches – the ugly

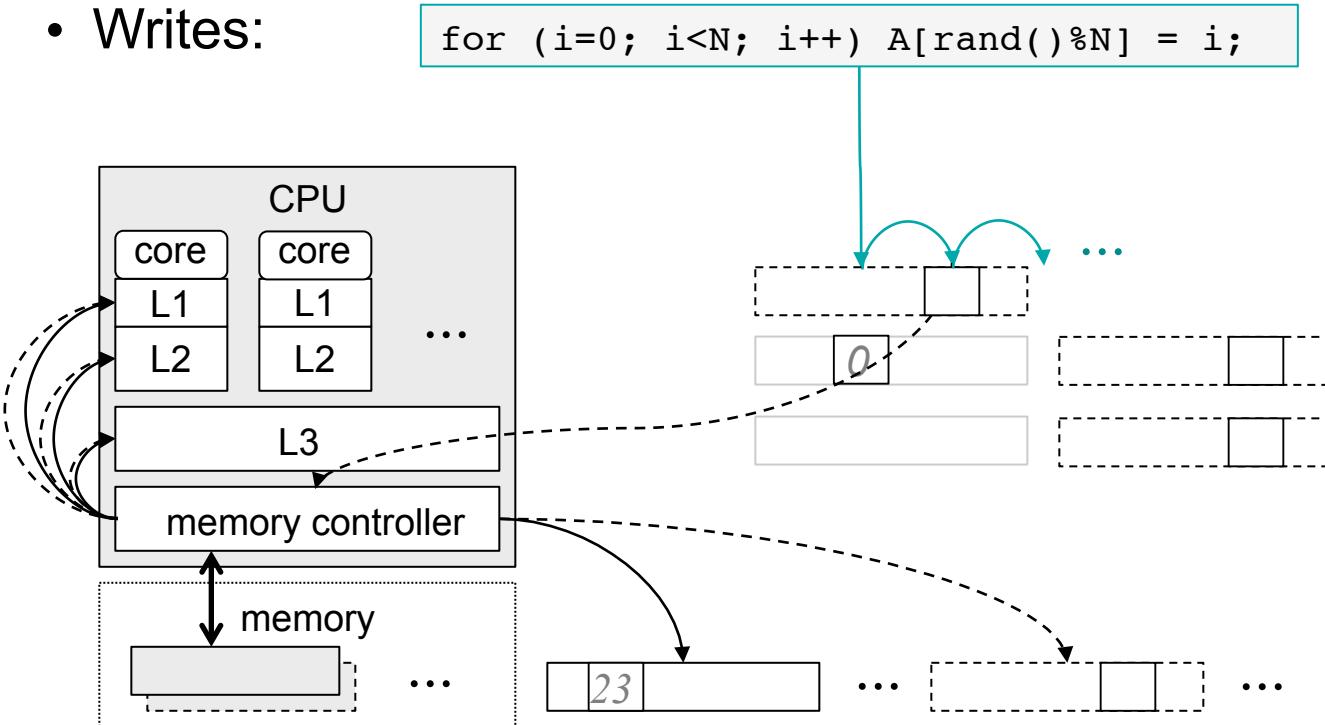
- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Writes:

```
for (i=0; i<N; i++) A[rand()%N] = i;
```



Caches – the ugly

- Memory is physically accessed at cache line granularity, e.g. 64Byte
- Writes:



- Many memory addresses map into the same cache line(s)
 - “Dirty” cache line needs to be evicted before new one loads
- Writes effectively turn into read-modify-write, c.f. disks