

# Programming GPUs for database operations



**Tim Kaldewey**

Research Staff Member

IBM TJ Watson Research Center

[tkaldew@us.ibm.com](mailto:tkaldew@us.ibm.com)

## Disclaimer

The author's views expressed in this presentation do not necessarily reflect the views of IBM.

## Acknowledgements

I would like to thank all my co-authors from IBM and my prior positions at Oracle and UCSC whose work I am also showing in this presentation.

I would also like to thank Patrick Cozzi for inviting me to teach in this class multiple years in a row.

## Agenda

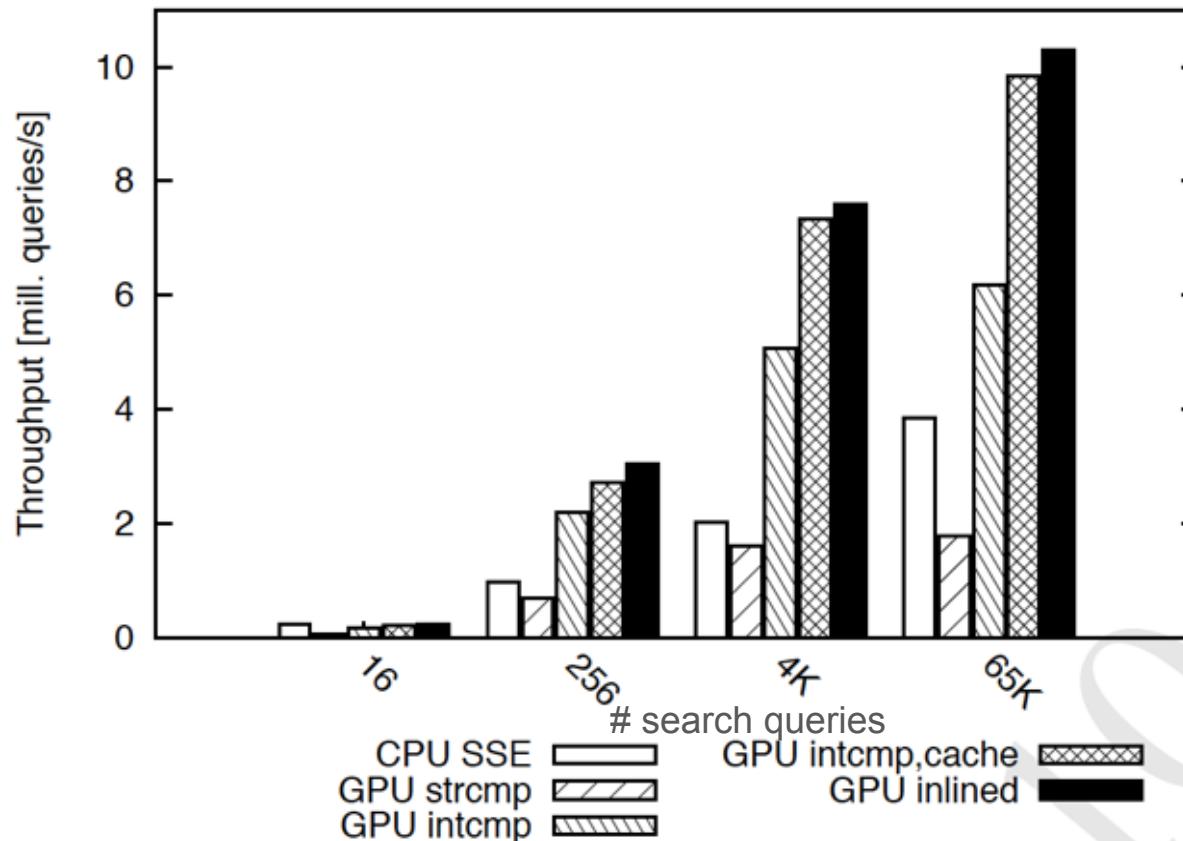
- GPU search
  - Reminder: Porting CPU search
  - Back to the drawing board:
    - P-ary search
    - Experimental evaluation
    - Why it works
- Building a GPU based data warehouse solution
  - From a query to operators
  - What to accelerate
  - What are the bottlenecks/limitations
- Maximizing data path efficiency
  - Extremely fast storage solution
  - Storage to host to device
- Putting it all together
  - Prototype demo

## Binary Search on the GPU – optimized

- Replace byte-wise strcmp with larger word size (uint4)
  - What happens if we load character strings as integers ?
- Prefetch (cache) intermediate values in shared memory
  - Don't newer GPUs have caches ?
- Inline the function calls

## Binary Search on the GPU – optimized

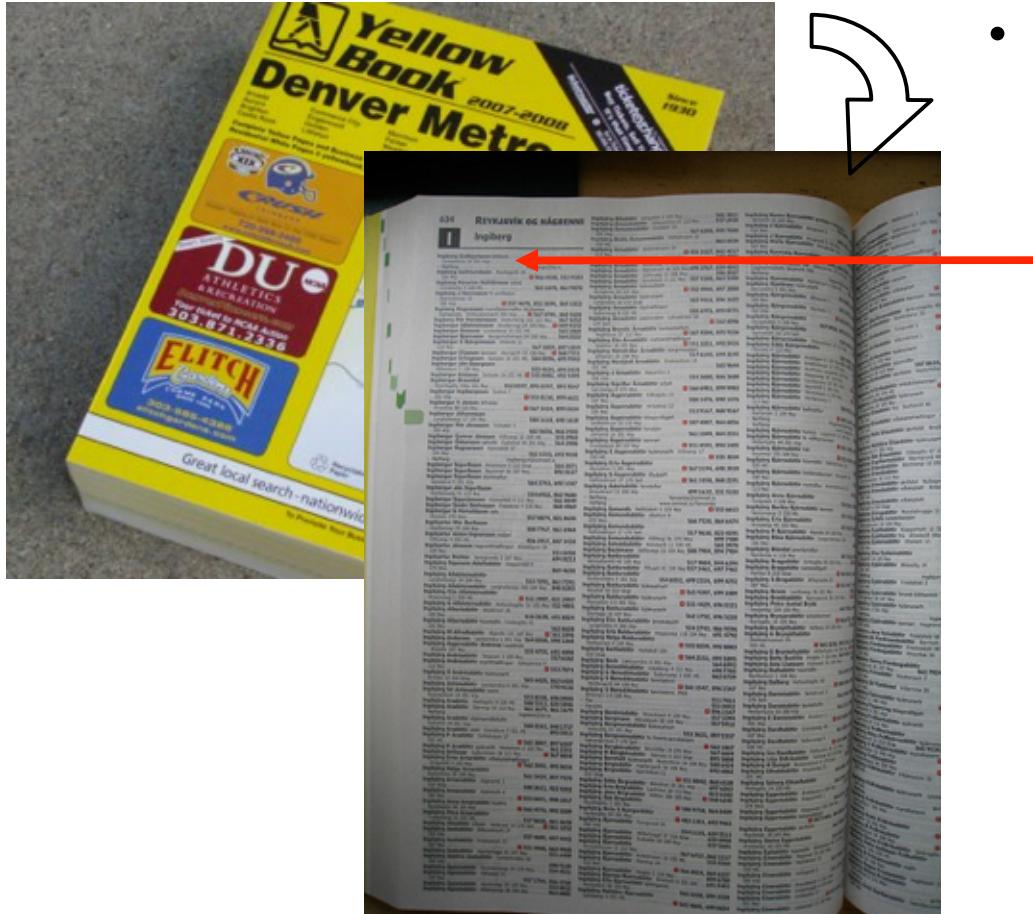
- Replace byte-wise strcmp with larger word size (uint4)
  - What happens if we load character strings as integers ?
- Prefetch (cache) intermediate values in shared memory
  - Don't newer GPUs have caches ?
- Inline the function calls



Searching a large data set (512MB) with 33 million ( $2^{25}$ ) 16-character strings

# Binary Search

- How Do you (efficiently) search an index?

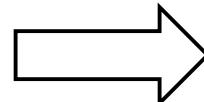


- Open phone book ~middle

- 1st name = whom you are looking for?
- < , > ?
- Iterate
  - Each iteration:  $\# \text{entries}/2$  ( $n/2$ )
  - Total time:  
→  $\log_2(n)$

## Parallel (Binary) Search

- What if you have some friends (3) to help you ?

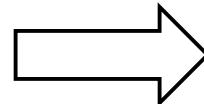


- Divide et impera !
  - Each is using binary search takes  $\log_2(n/4)$
  - All can work in parallel → faster:  $\log_2(n/4) < \log_2(n)$
- Give each of them  $\frac{1}{4}$  \*

\* You probably want to tear it a little more intelligent than that, e.g. at the binding ;-)

## Parallel (Binary) Search

- What if you have some friends (3) to help you ?

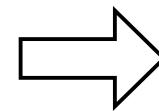
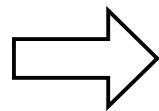


- Divide et impera !
  - Each is using binary search takes  $\log_2(n/4)$
  - All can work in parallel → faster:  $\log_2(n/4) < \log_2(n)$
  - 3 of you are **wasting time** !
- Give each of them  $\frac{1}{4}$  \*

\* You probably want to tear it a little more intelligent than that, e.g. at the binding ;-)

## P-ary Search

- Divide et impera !!



...

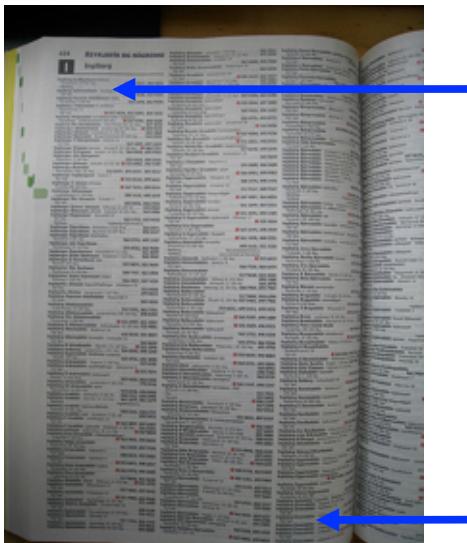
- How do we know who has the right piece ?

## P-ary Search

- Divide et impera !!



- How do we know who has the right piece ?



- It's a sorted list:
  - Look at first and last entry of a subset
  - If **first entry** < searched name < **last entry**
    - Redistribute
    - Otherwise ... throw it away
  - Iterate

## P-ary Search

- What do we get?



+

- Each iteration:  $n/4$   
→  $\log_4(n)$
- Assuming redistribution time is negligible:  
 $\log_4(n) < \log_2(n/4) < \log_2(n)$
- But each does 2 lookups !
- How time consuming are lookup and redistribution ?

## P-ary Search

- What do we get?



+

- Each iteration:  $n/4$   
→  $\log_4(n)$
- Assuming redistribution time is negligible:  
 $\log_4(n) < \log_2(n/4) < \log_2(n)$
- But each does 2 lookups !
- How time consuming are lookup and redistribution ?

||

||

memory      synchronization  
access

## P-ary Search

- What do we get?



+

- Each iteration:  $n/4$   
→  $\log_4(n)$
- Assuming redistribution time is negligible:  
 $\log_4(n) < \log_2(n/4) < \log_2(n)$
- But each does 2 lookups !
- How time consuming are lookup and redistribution ?

II

memory access

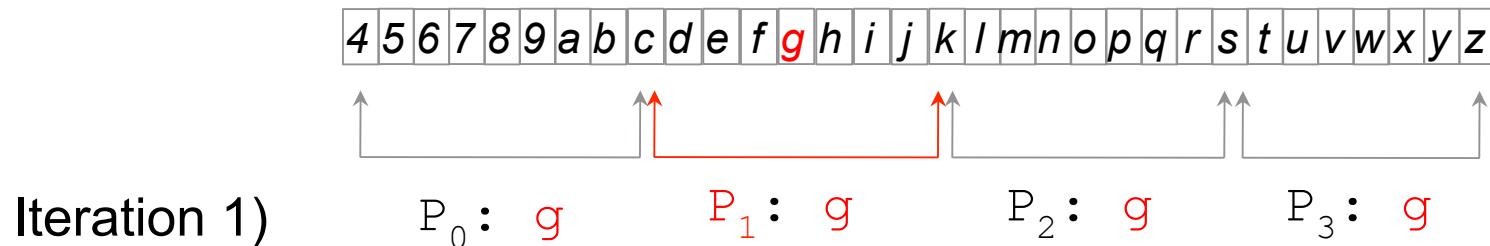
II

synchronization

- Searching a database index can be implemented the same way
  - Friends = Processor cores (threads)
  - Without destroying anything ;-)

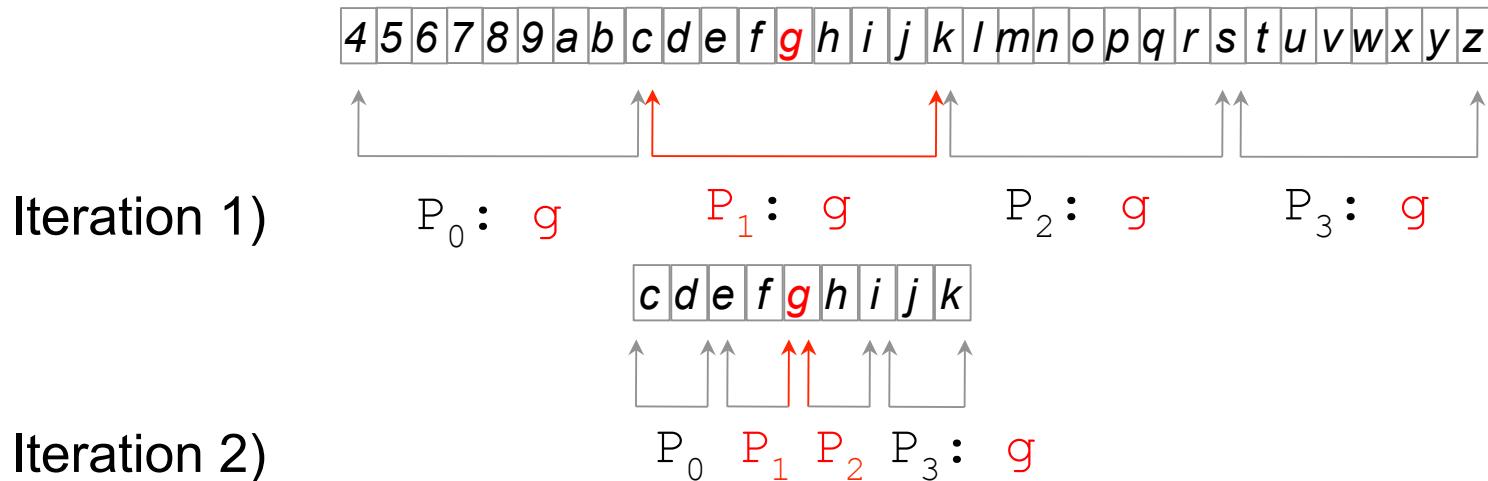
## P-ary Search - Implementation

- Strongly relies on fast synchronization
  - friends = threads / vector elements



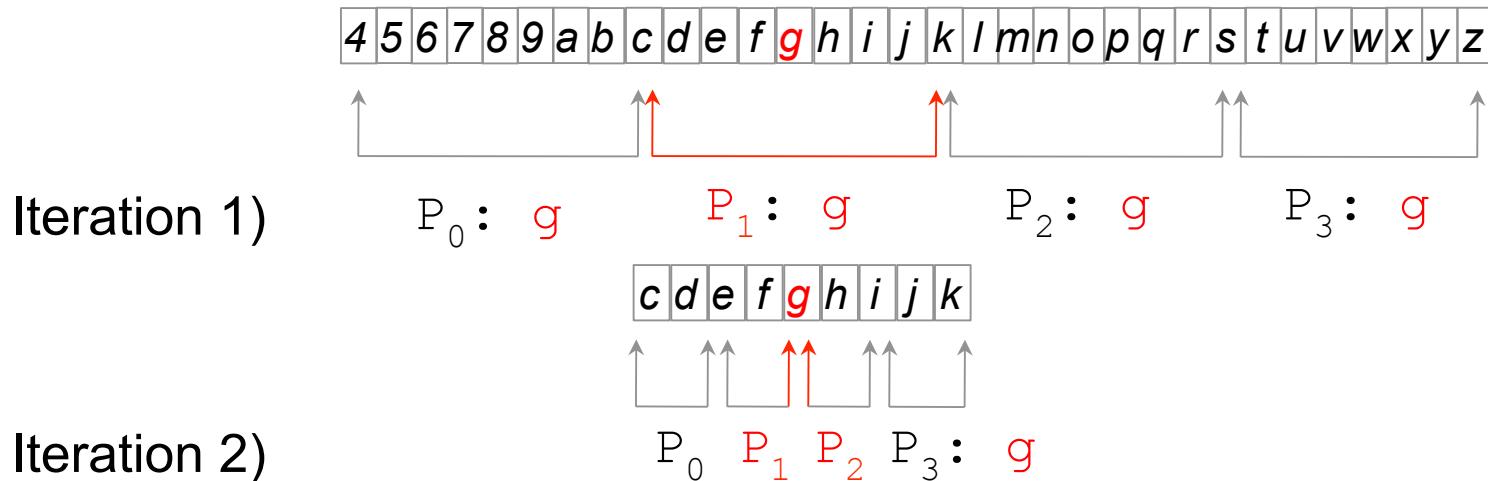
## P-ary Search - Implementation

- Strongly relies on fast synchronization
  - friends = threads / vector elements



## P-ary Search - Implementation

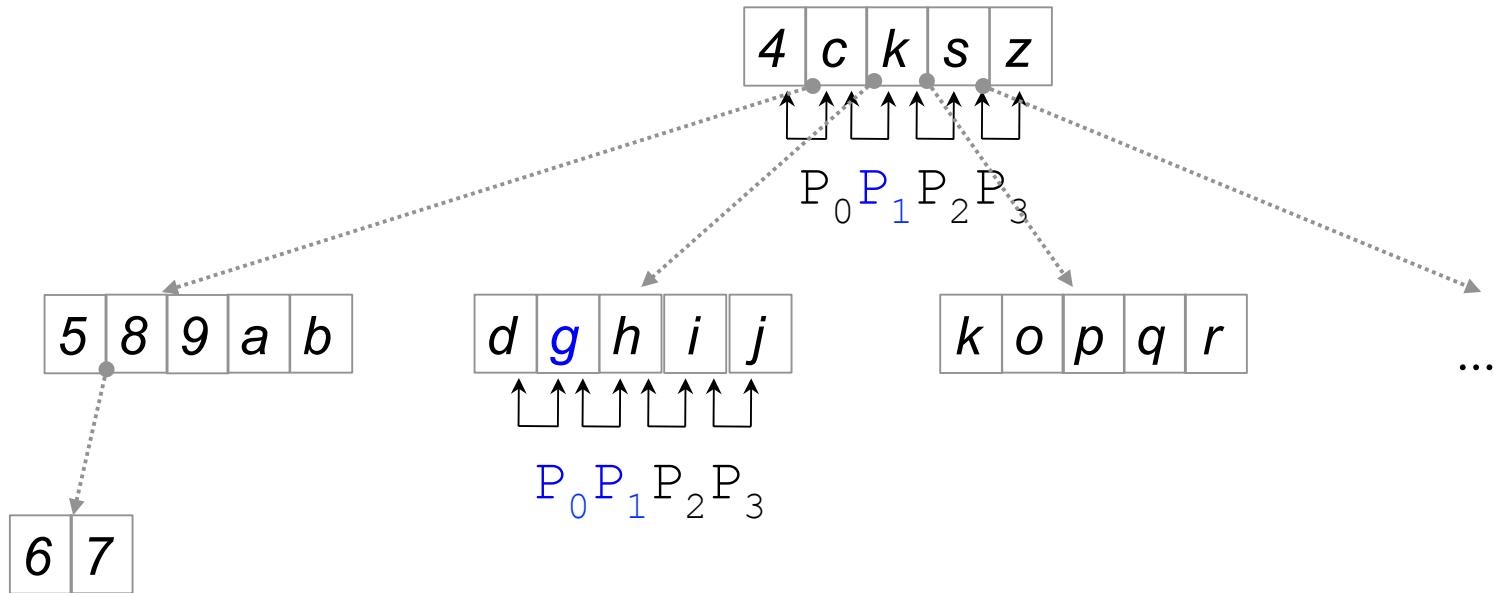
- Strongly relies on fast synchronization
  - friends = threads / vector elements



- Synchronization ~ repartition cost
- pthreads (\$\$), **cmpxchng(\$)**
- SIMD SSE-vector, GPU threads via shared memory (~0)
- Implementation using a B-tree is similar and (obviously) faster

## P-ary Search - Implementation

- B-trees group pivot elements into nodes



- Access to pivot elements is coalesced instead of a gather
- Nodes can also be mapped to
  - Cache Lines (CSB+ trees)
  - Vectors (SSE)
  - #Threads per block

## P-ary Search on a sorted integer list – Implementation (1)

```
__shared__ int offset;
__shared__ int cache[BLOCKSIZE+2]

__global__ void parySearchGPU(int* data, int length,
                           int* list_of_search_keys, int* results)

int start, sk;
int old_length = length;
// initialize search range starting with the whole data set
if (threadIdx.x == 0) {
    offset = 0;
    // cache search key and upper bound in shared memory
    cache[BLOCKSIZE] = 0x7FFFFFFF;
    cache[BLOCKSIZE+1] = list_of_search_keys[blockIdx.x];
    results[blockIdx.x] = -1;
}
__syncthreads();
//
sk = cache[BLOCKSIZE+1];
```

## P-ary Search on a sorted integer list – Implementation (1)

```
__shared__ int offset;
__shared__ int cache[BLOCKSIZE+2]

__global__ void parySearchGPU(int* data, int length,
                           int* list_of_search_keys, int* results)

int start, sk;
int old_length = length;
// initialize search range starting with the whole data set
if (threadIdx.x == 0) {
    offset = 0;
    // cache search key and upper bound in shared memory
    cache[BLOCKSIZE] = 0xFFFFFFFF;
    cache[BLOCKSIZE+1] = list_of_search_keys[blockIdx.x];
    results[blockIdx.x] = -1;
}
syncthreads();
//
sk = cache[BLOCKSIZE+1];
```

Why?

## P-ary Search on a sorted list – Implementation (2)

```
// repeat until the #keys in the search range < #threads
while (length > BLOCKSIZE) {
    // calculate search range for this thread
    length = length/BLOCKSIZE;
    if (length * BLOCKSIZE < old_length) length += 1;
    old_length = length;
    // why don't we just use floating point?
    start = offset + threadIdx.x * length;
    // cache the boundary keys
    cache[threadIdx.x] = data[start];
    __syncthreads();
    // if the searched key is within this thread's subset,
    // make it the one for the next iteration
    if (sk >= cache[threadIdx.x] && sk < cache[threadIdx.x+1]) {
        offset = start;
    }
    __syncthreads();
    // all threads start next iteration with the new subset
}
```

## P-ary Search on a sorted list – Implementation (2)

```
// repeat until the #keys in the search range < #threads
while (length > BLOCKSIZE) {
    // calculate search range for this thread
    length = length/BLOCKSIZE;
    if (length * BLOCKSIZE < old_length) length += 1;
    old_length = length;
    // why don't we just use floating point?
    start = offset + threadIdx.x * length;
    // cache the boundary keys
    cache[threadIdx.x] = data[start];
    syncthreads();
    // if the searched key is within this thread's subset,
    // make it the one for the next iteration
    if (sk >= cache[threadIdx.x] && sk < cache[threadIdx.x+1]) {
        offset = start;
    }
    syncthreads();
    // all threads start next iteration with the new subset
}
```

Why?

## P-ary Search on a sorted list – Implementation (3)

```
// last iteration
start = offset + threadIdx.x;
if (sk == data[start])
    results[blockIdx.x] = start;
}
```

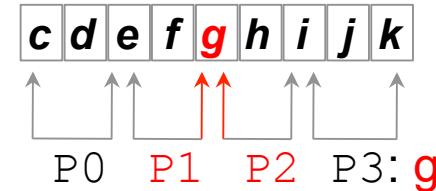
## P-ary Search on a sorted list – Implementation (3)

```
// last iteration  
start = offset + threadIdx.x;  
if (sk == data[start])  
    results[blockIdx.x] = start;  
}
```

Why don't cache?

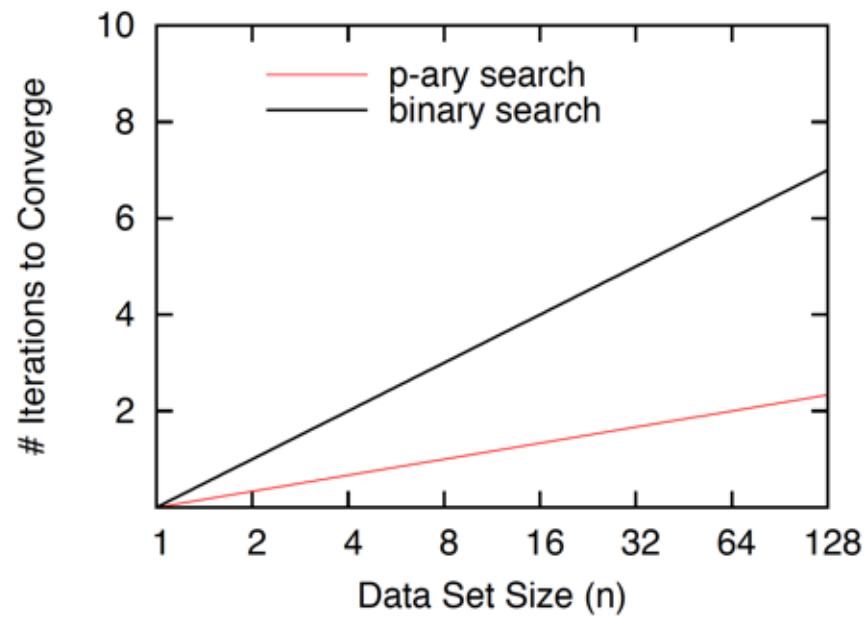
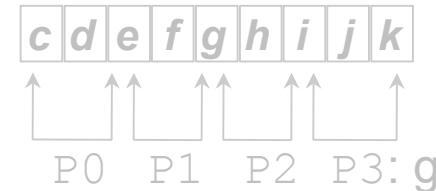
## P-ary Search – Analysis

- 100% processor utilization for each query
- Multiple threads can find a result
  - How does this impact correctness?



## P-ary Search – Analysis

- 100% processor utilization for each query
- Multiple threads can find a result
  - How does this impact correctness?
- Convergence depends on #threads
- GTX285: 1 SM, 8 cores(threads) →  $p=8$
- Better Response time
  - $\log_p(n)$  vs  $\log_2(n)$

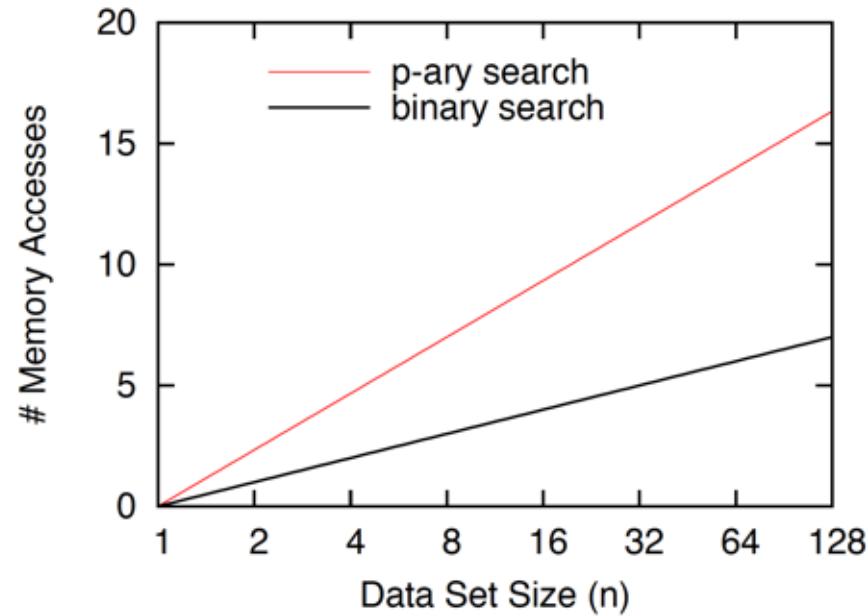
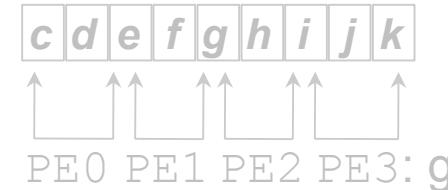


## P-ary Search – Analysis

- 100% processor utilization for each query
- Multiple threads can find a result
  - Does not change correctness
- Convergence depends on #threads

GTX285: 1 SM, 8 cores(threads)  $\rightarrow p=8$

- Better Response time
  - $\log_p(n)$  vs  $\log_2(n)$
- More memory access
  - $(p^2 \text{ per iteration}) * \log_p(n)$
  - Caching
  - $(p-1) * \log_p(n)$  vs.  $\log_2(n)$

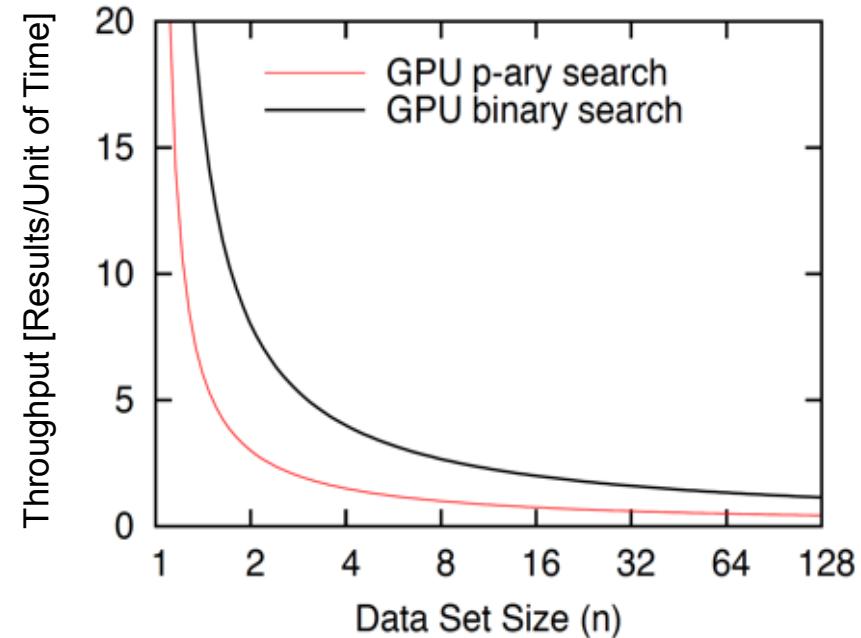
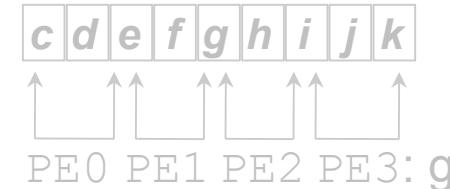


## P-ary Search – Analysis

- 100% processor utilization for each query
- Multiple threads can find a result
  - Does not change correctness
- Convergence depends on #threads

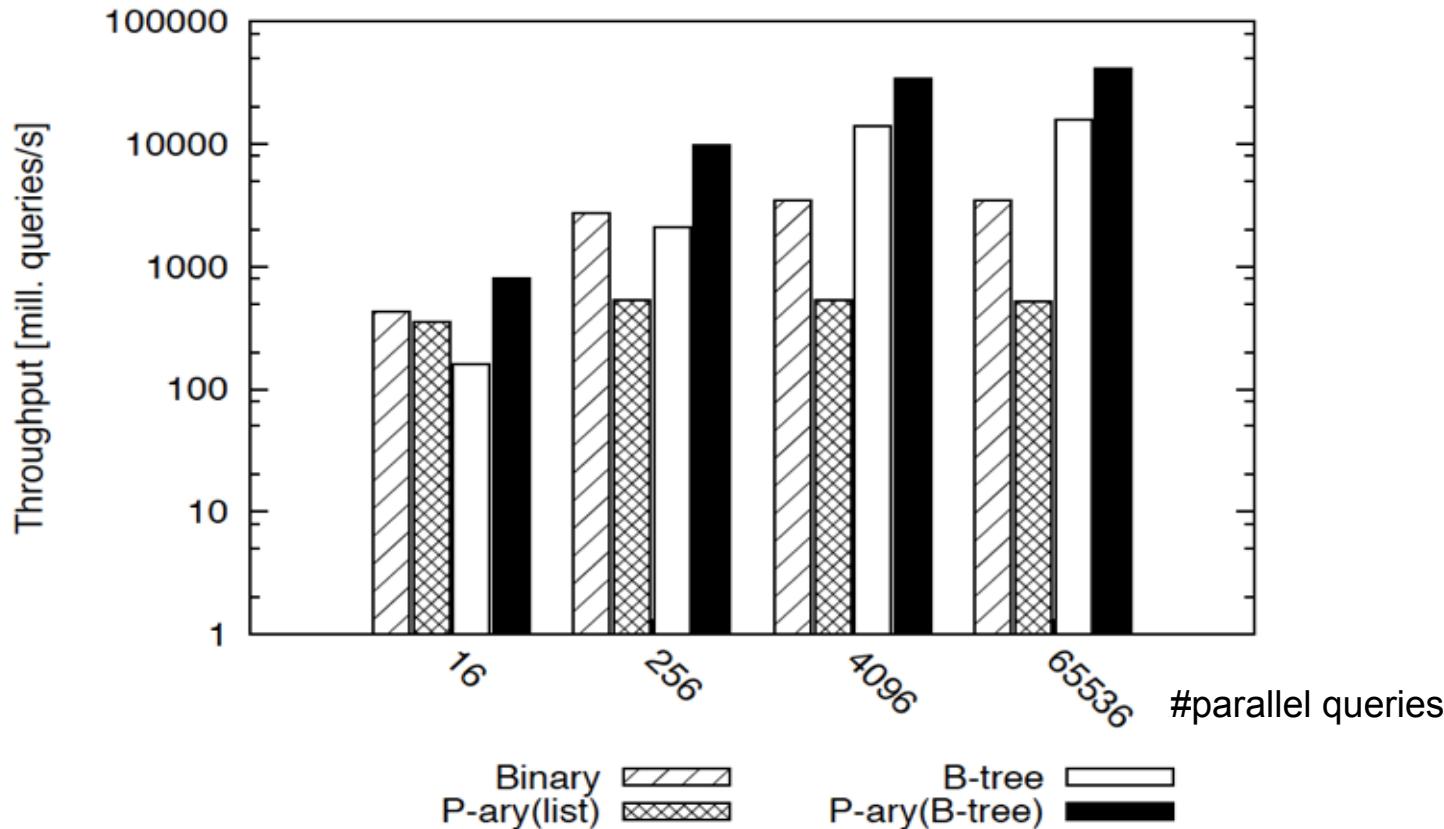
GTX285: 1 SM, 8 cores(threads)  $\rightarrow p=8$

- Better Response time
  - $\log_p(n)$  vs  $\log_2(n)$
- More memory access
  - $p^2$  per iteration \*  $\log_p(n)$
  - Caching
  - $(p-1) \cdot \log_p(n)$  vs.  $\log_2(n)$
- Lower Throughput
  - $1/\log_p(n)$  vs  $p/\log_2(n)$



## P-ary Search (GPU) – Throughput

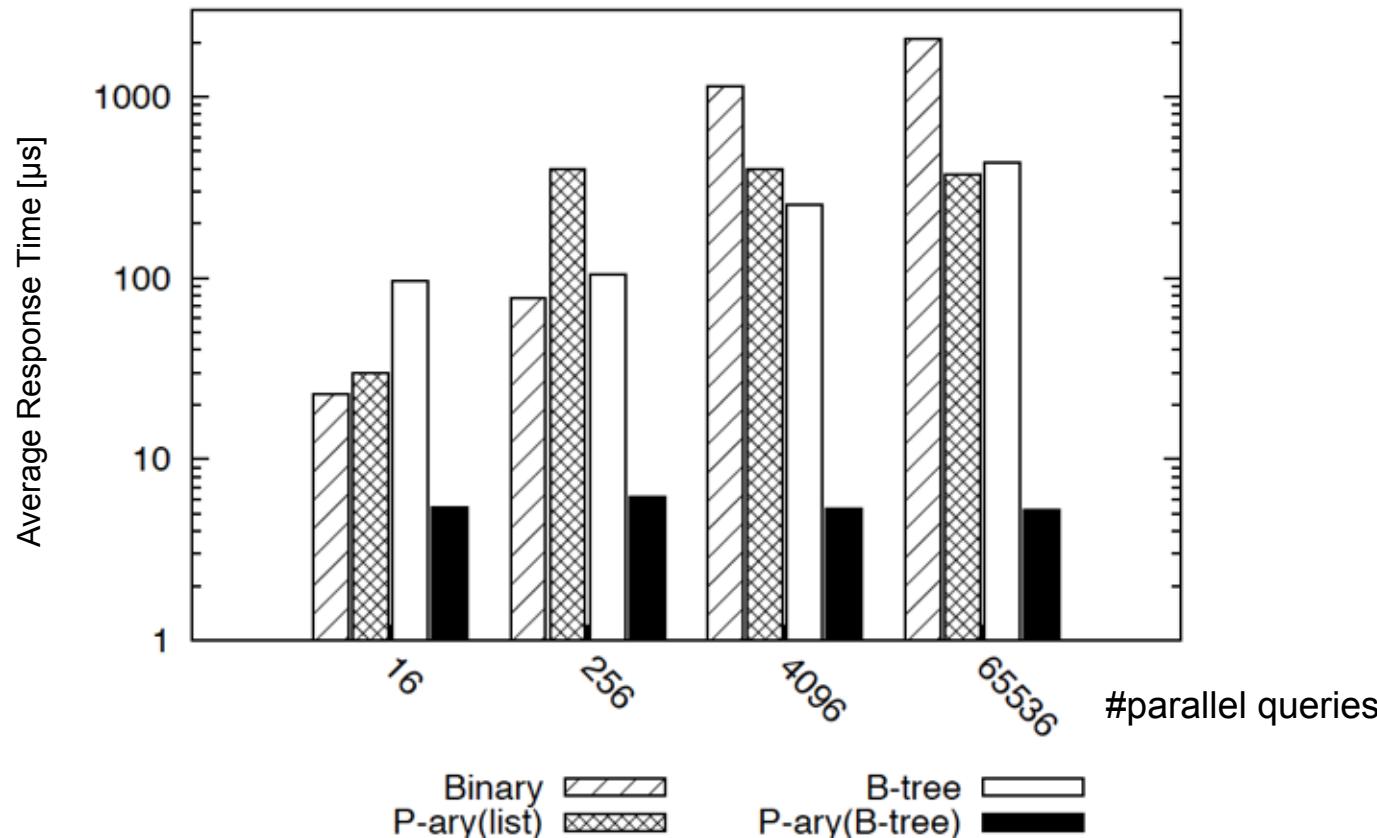
- Superior throughput compared to conventional algorithms



Searching a 512MB data set with 134mill. 4-byte integer entries,  
Results for a nVidia GT200b, 1.5GHz, GDDR3 1.2GHz.

## P-ary Search (GPU) – Response Time

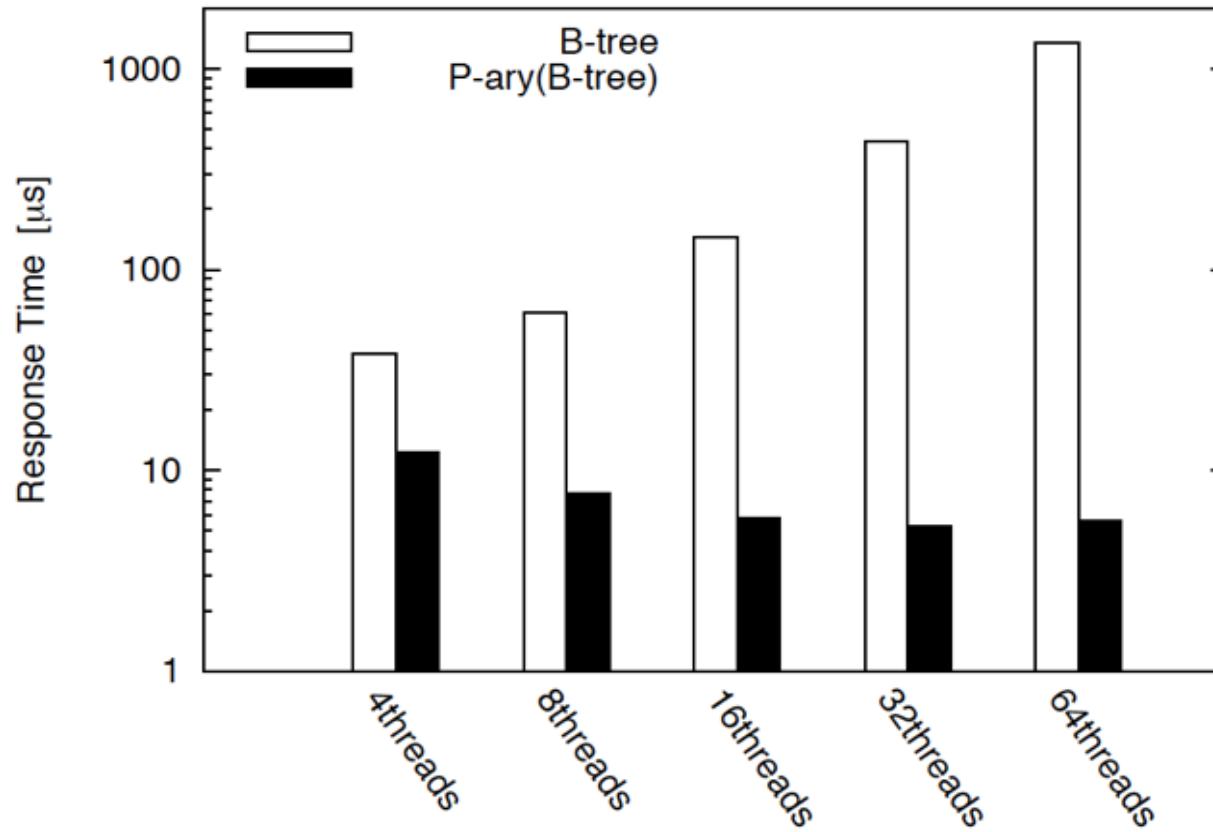
- Response time is workload independent for B-tree implementation



Searching a 512MB data set with 134mill. 4-byte integer entries,  
Results for a nVidia GT200b, 1.5GHz, GDDR3 1.2GHz.

## P-ary Search (GPU) – Scalability

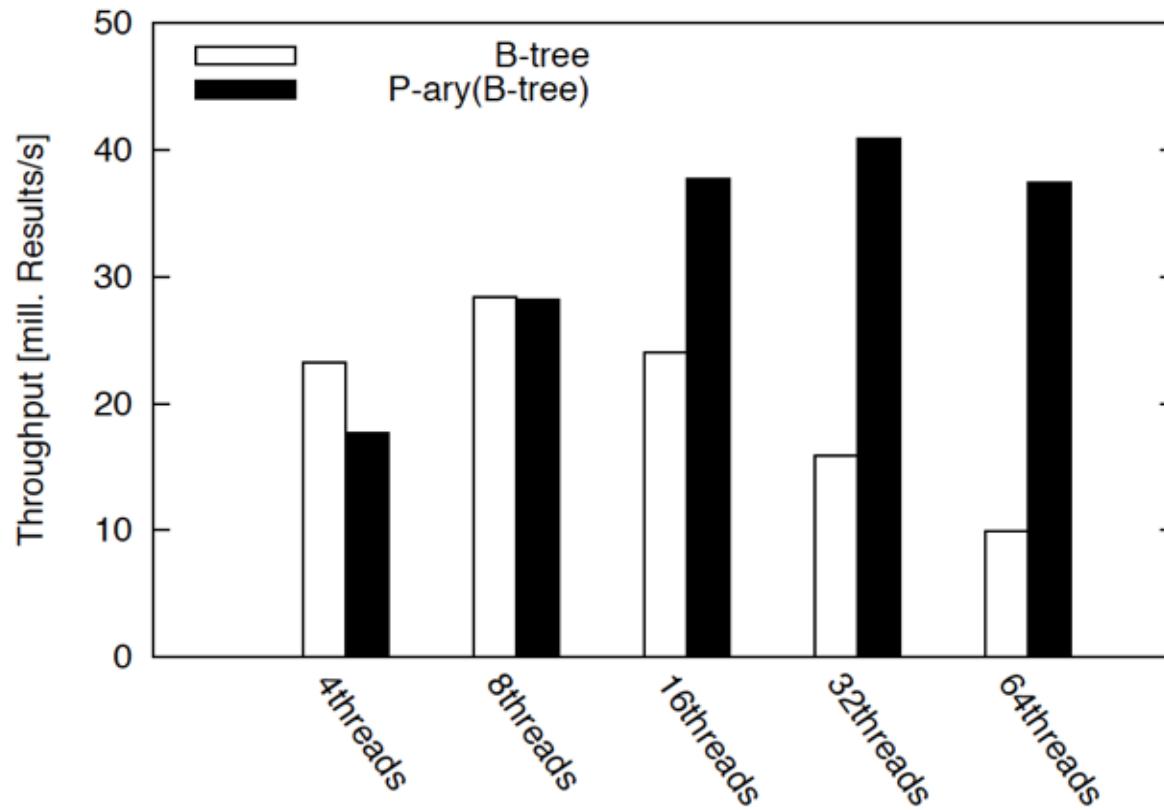
- GPU Implementation using SIMT (SIMD threads)
- Scalability with increasing #threads (P)



64K search queries against a 512MB data set with 134mill. 4-byte integer entries,  
Results for a nVidia GT200b, 1.5GHz, GDDR3 1.2GHz.

## P-ary Search (GPU) – Scalability

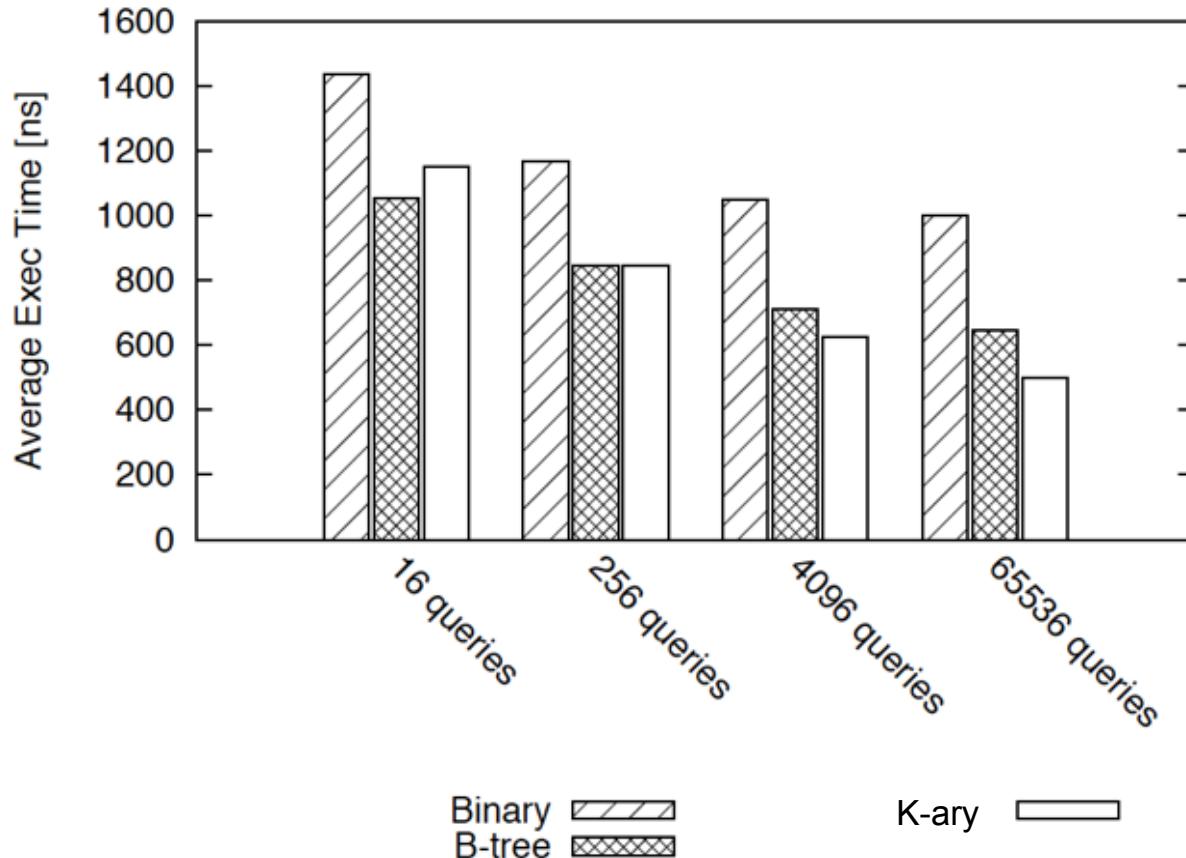
- GPU Implementation using SIMT (SIMD threads)
- Scalability with increasing #threads (P)



64K search queries against a 512MB data set with 134mill. 4-byte integer entries,  
Results for a nVidia GT200b, 1.5GHz, GDDR3 1.2GHz.

## P-ary Search(CPU) = K-ary Search<sup>1</sup>

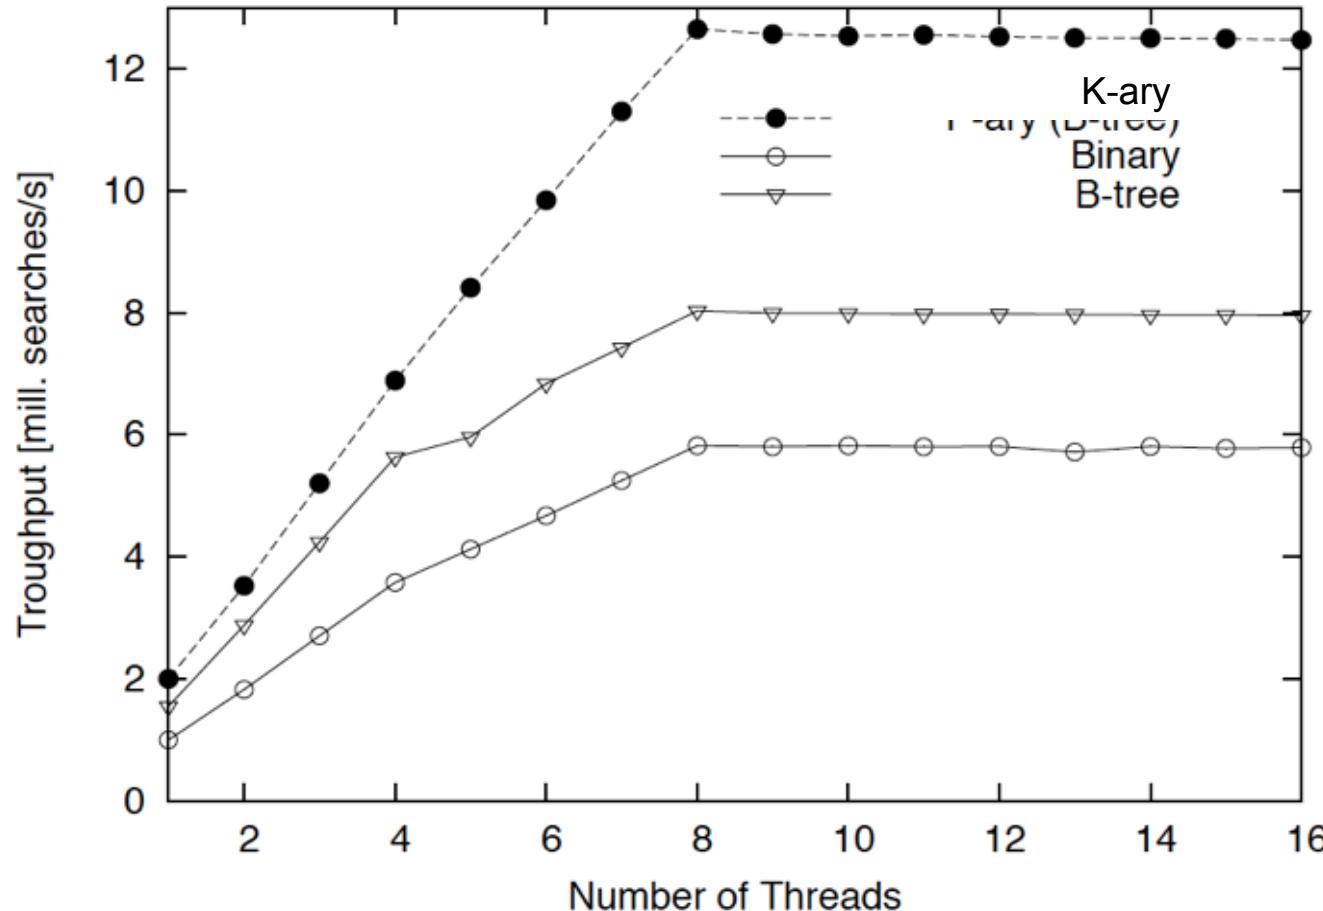
- K-ary search is the same algorithm ported to the CPU using SSE vectors (int4) → convergence rate  $\log_4(n)$



Searching a 512MB data set with 134mill. 4-byte integer entries,  
Core i7 2.66GHz, DDR3 1666.

## P-ary Search(CPU) = K-ary Search<sup>1</sup>

- Throughput scales proportional to #threads



64K search queries against a 512MB data set with 134mill. 4-byte integer entries,  
Core i7 2.66GHz, DDR3 1666.

## Agenda

- GPU search
  - Reminder: Porting CPU search
  - Back to the drawing board:
    - P-ary search
    - Experimental evaluation
    - Why it works
- Building a GPU based data warehouse solution
  - From a query to operators
  - What to accelerate
  - What are the bottlenecks/limitations
- Maximizing data path efficiency
  - Extremely fast storage solution
  - Storage to host to device
- Putting it all together
  - Prototype demo

## A data warehousing query in multiple languages

- **English:** Show me the annual development of revenue from US sales of US products for the last 5 years by city

## A data warehousing query in multiple languages

- **English:** Show me the **annual** development of **revenue** from **US sales** of **US products** for the last **5 years** by **city**

- **SQL:**

```
SELECT c.city, s.city, d.year, SUM(lo.revenue)
  FROM lineorder lo, customer c, supplier s, date d
 WHERE lo.custkey = c.custkey
   AND lo.suppkey = s.suppkey
   AND lo.orderdate = d.datekey
   AND c.nation = 'UNITED STATES'
   AND s.nation = 'UNITED STATES'
   AND d.year >= 1998 AND d.year <= 2012
 GROUP BY c.city, s.city, d.year
 ORDER BY d.year asc, revenue desc;
```

## A data warehousing query in multiple languages

- **English:** Show me the annual development of revenue from US sales of US products for the last 5 years by city

- **SQL:**

```
SELECT c.city, s.city, d.year, SUM(lo.revenue)
  FROM lineorder lo, customer c, supplier s, date d
 WHERE lo.custkey = c.custkey
   AND lo.suppkey = s.suppkey
   AND lo.orderdate = d.datekey ?
   AND c.nation = 'UNITED STATES'
   AND s.nation = 'UNITED STATES'
   AND d.year >= 1998 AND d.year <= 2012
 GROUP BY c.city, s.city, d.year
 ORDER BY d.year asc, revenue desc;
```

## Star Schema – typical for DWH

*Customer*

CUSTKEY
NAME
ADDRESS
CITY
...

*Supplier*

SUPPKEY
NAME
ADDRESS
CITY
...

*Lineorder*

ORDERKEY
LINENUMBER
CUSTKEY
PARTKEY
SUPPKEY
ORDERDATE
ORDPRIORITY
...
...
COMMITDATE
SHIPMODE

*Part*

PARTKEY
NAME
MFGR
CATEGORY
BRAND
...

*Date*

DATEKEY
DATE
DAYOFWEEK
MONTH
YEAR
...

Query:

```

SELECT c.city, s.city, d.year, SUM(lo.revenue) FROM lineorder lo, customer c, supplier s, date d
WHERE lo.custkey = c.custkey AND lo.suppkey = s.suppkey AND lo.orderdate = d.datekey AND
c.nation = 'UNITED STATES' AND s.nation = 'UNITED STATES' AND d.year >= 1998 AND d.year <= 2012
GROUP BY c.city, s.city, d.year ORDER BY d.year asc, revenue desc;
    
```

## A data warehousing query in multiple languages

- **English:** Show me the annual development of revenue from US sales of US products for the last 5 years by city

- **SQL:**

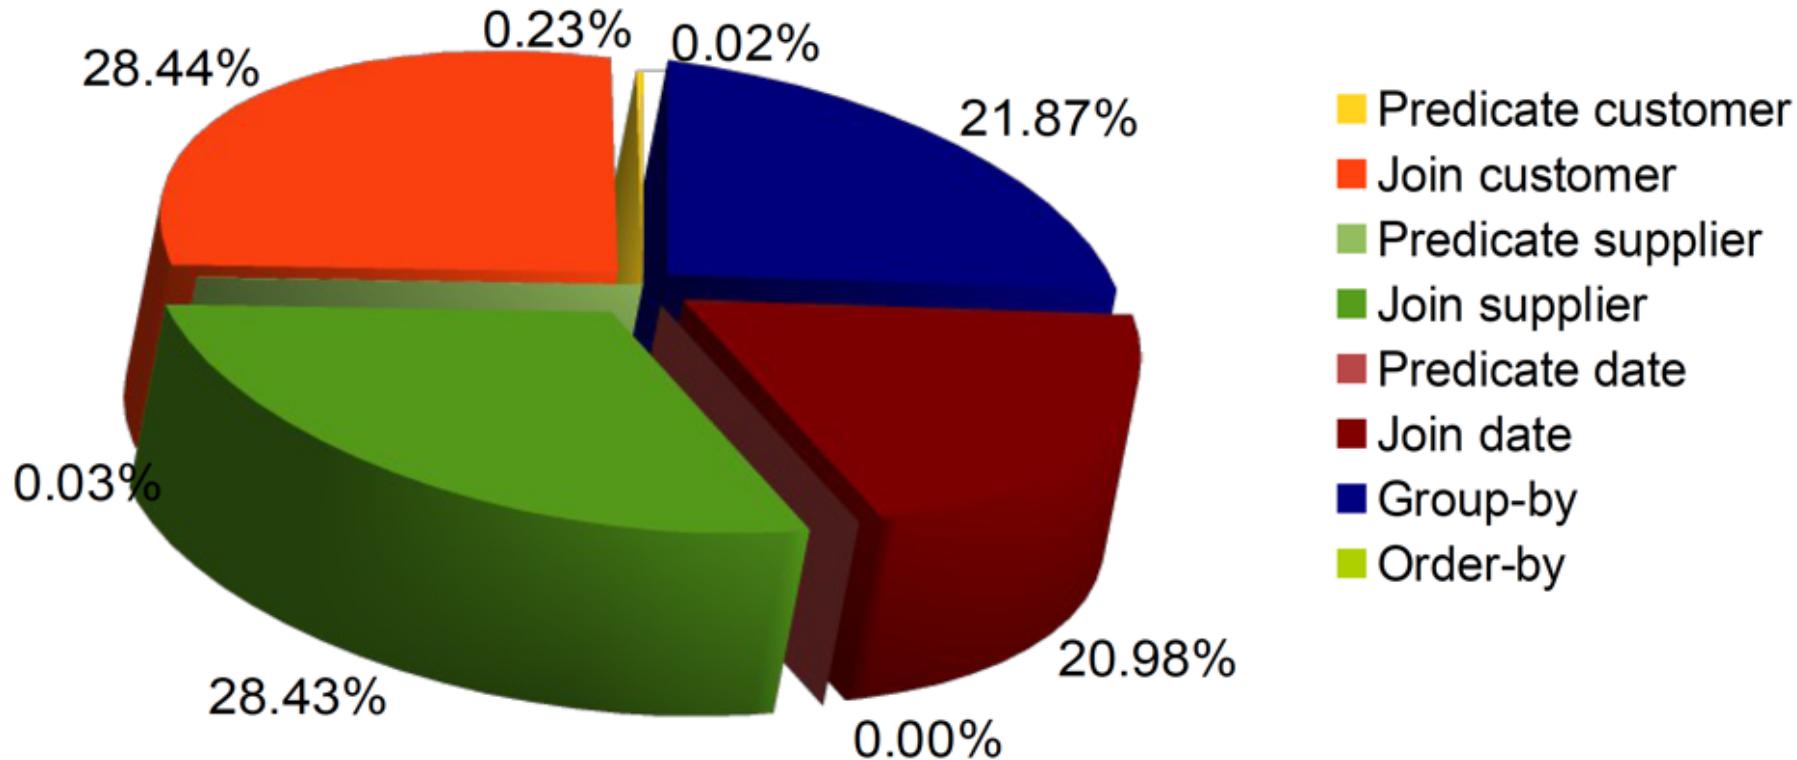
```
SELECT c.city, s.city, d.year, SUM(lo.revenue)
  FROM lineorder lo, customer c, supplier s, date d
 WHERE lo.custkey = c.custkey
   AND lo.suppkey = s.suppkey
   AND lo.orderdate = d.datekey
   AND c.nation = 'UNITED STATES'
   AND s.nation = 'UNITED STATES'
   AND d.year >= 1998 AND d.year <= 2012
 GROUP BY c.city, s.city, d.year
 ORDER BY d.year asc, revenue desc;
```

### Database primitives (operators):

- Predicate(s): **customer**, **supplier**, and **date** *direct filter (yes/no)*
- *Join(s)*: **lineorder** with **part**, **supplier**, and **date** *correlate tables & filter*
- Group By (aggregate): **city** and **date** *correlate tables & sum*
- Order By: **year** and **revenue** *sort*

What are the most time-consuming operations?

## Where does time go?

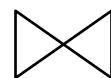


```
SELECT c.city, s.city, d.year, SUM(lo.revenue)
  FROM lineorder lo, customer c, supplier s, date d
 WHERE c.nation = 'UNITED STATES'          AND lo.custkey = c.custkey
       AND s.nation = 'UNITED STATES'          AND lo.suppkey = s.suppkey
       AND d.year >= 1998 AND d.year <= 2012 AND lo.orderdate = d.datekey
 GROUP BY c.city, s.city, d.year
 ORDER BY d.year asc, revenue desc;
```

## Relational Joins

**Sales (Fact Table)**

Revenue	Customer
\$10.99	23
\$49.00	14
\$11.00	56
\$103.00	11
\$84.50	39
\$60.10	27
\$7.60	23



**Customers (living in US)**

Key	Zip
11	95014
23	94303
27	95040
39	95134



Revenue	Zip
\$10.99	94303
\$103.00	95014
\$84.50	95134
\$60.10	95040
\$7.60	94303



Join  
Results

Measure (m)

Primary Key (k)      Payload (p)

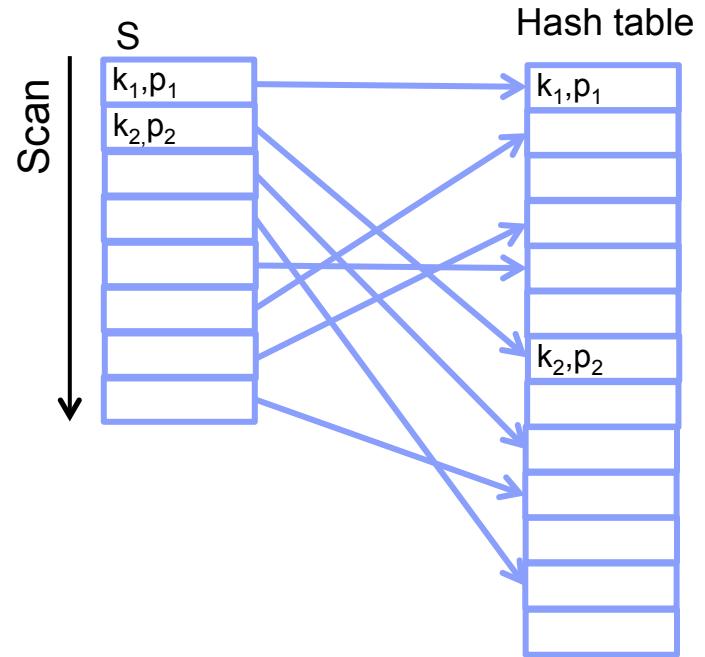
Foreign Key (fk)

## Hash Join

Join two tables ( $|S| < |R|$ ) in 2 steps

### 1. Build a hash table

- Scan S and compute a location (hash) based on a unique (primary) key
- Insert primary key **k** with payload **p** into the hash table
- If the location is occupied pick the next free one (open addressing)



## Hash Join

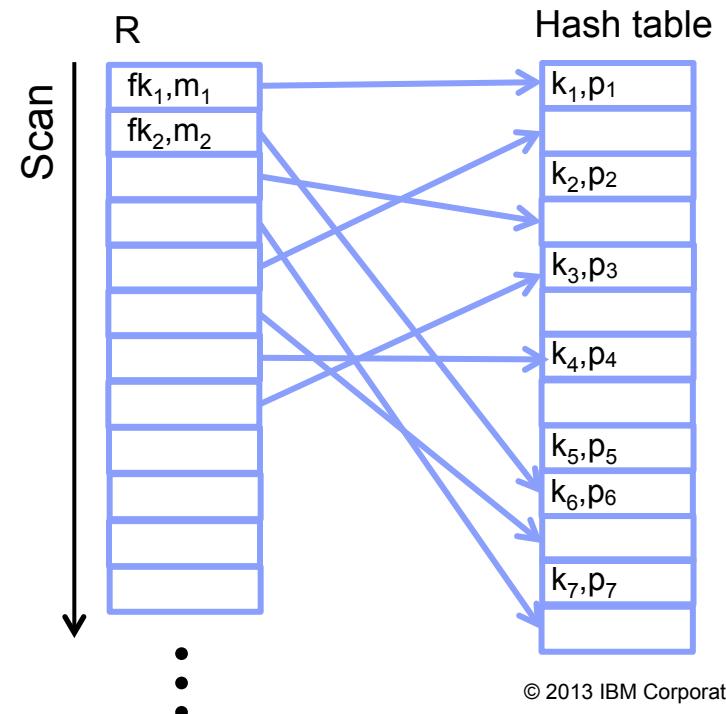
Join two tables ( $|S| < |R|$ ) in 2 steps

### 1. Build a hash table

- Scan S and compute a location (hash) based on a unique (primary) key
- Insert primary key **k** with payload **p** into the hash table
- If the location is occupied pick the next free one (open addressing)

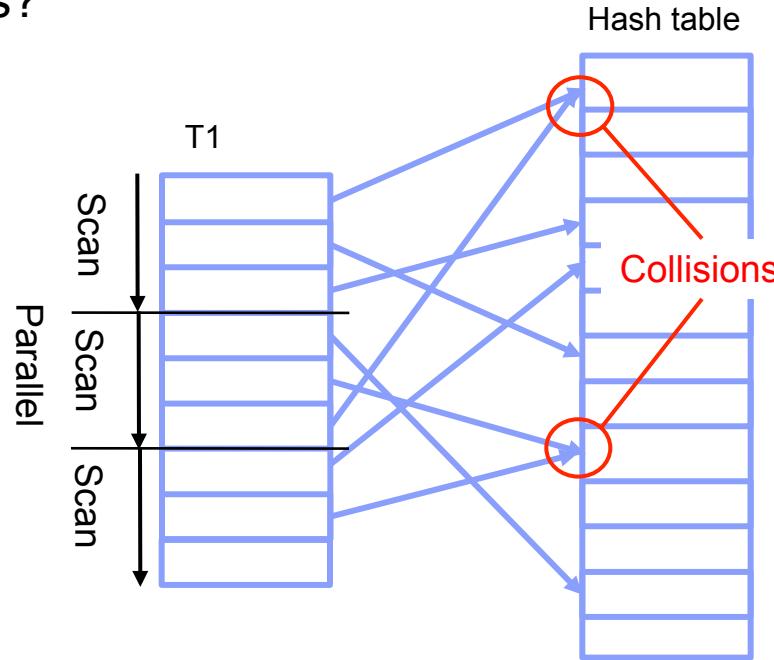
### 2. Probe the hash table

- Scan R and compute a location (hash) based on the reference to S (foreign key)
- Compare foreign key **fk** and key **k** in hash table
- If there is a match store the result (**m,p**)



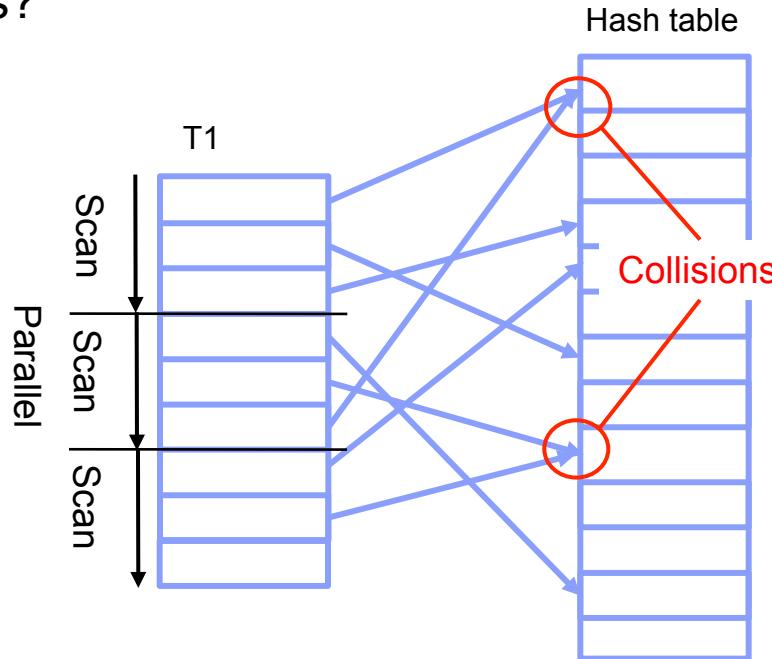
## Parallel Hash Join

- Multiple threads scan T1 and attempt to insert  $\langle \text{key}, \text{rid} \rangle$  pairs into the hash table
  - How to handle hash collisions?



## Parallel Hash Join

- Multiple threads scan T1 and attempt to insert  $\langle \text{key}, \text{rid} \rangle$  pairs into the hash table
  - How to handle hash collisions?



- Is this a good access pattern?
- Parallel probe is trivial as it requires read-only access

## Hash Join

**Sales (Fact Table)**

Revenue	Customer
\$10.99	23
\$49.00	14
\$11.00	56
\$103.00	11
\$84.50	39
\$60.10	27
\$7.60	23

Probe Inputs

**Hash Table (HT)**

**Customers (living in US)**

Key	Zip
11	95014
23	94303
27	95040
39	95134



Primary Key (k)

Payload (p)

Foreign Key (fk)

Revenue	Zip
\$10.99	94303
\$103.00	95014
\$84.50	95134
\$60.10	95040
\$7.60	94303

Join  
Results

How fast are hash probes ?  

- Computation
- Data (memory) access

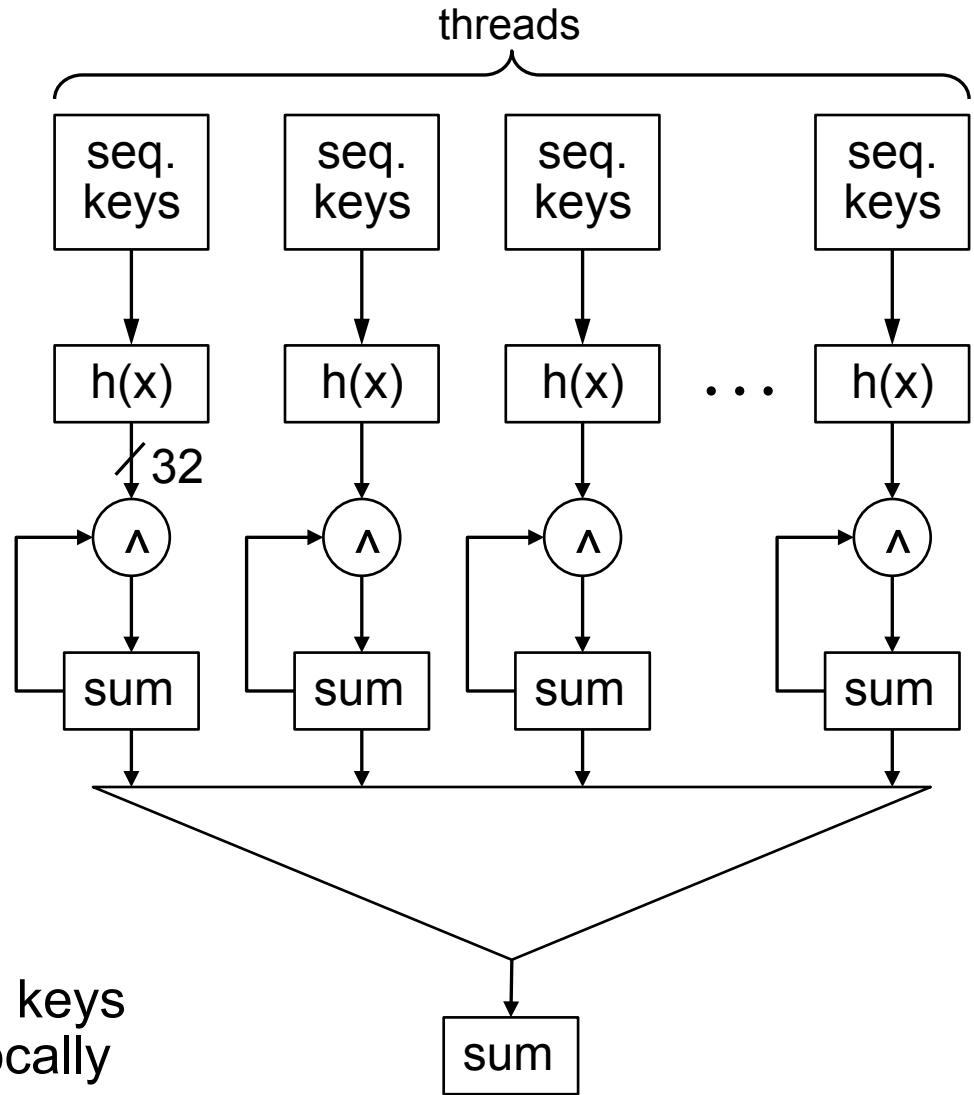
# Computing Hash Functions on GTX580 – Compute only \*

32-bit keys, 32-bit hashes

Hash Function/ Key Ingest GB/s	Seq keys+ Hash
LSB	338
Fowler-Noll-Vo 1a	129
Jenkins Lookup3	79
Murmur3	111
One-at-a-time	85
CRC32	78
MD5	4.5
SHA1	0.81

Cryptographic message  
digests

- Threads generate sequential keys
- Hashes are XOR-summed locally



## Hash Join – Data Access Patterns

- Primary data access patterns:
  - *Scan* the input table(s) for HT creation and probe
  - *Compare and swap* when inserting data into HT
  - *Random read* when probing the HT

## Hash Join – Data Access Patterns

- Primary data access patterns:
  - *Scan* the input table(s) for HT creation and probe
  - *Compare and swap* when inserting data into HT
  - *Random read* when probing the HT
- Data (memory) access on



	<b>GPU (GTX580)</b>	<b>CPU (i7-2600)</b>	
Peak memory bandwidth [spec] <sup>1)</sup>	179 GB/s	21 GB/s	Upper bound for:
Peak memory bandwidth [measured] <sup>2)</sup>	153 GB/s	18 GB/s	Scan R, S

(1) Nvidia:  $192.4 \times 10^6$  B/s  $\approx$  179.2 GB/s

(2) 64-bit accesses over 1 GB of device memory

## Hash Join – Data Access Patterns

- Primary data access patterns:
  - *Scan* the input table(s) for HT creation and probe
  - *Compare and swap* when inserting data into HT
  - *Random read* when probing the HT
- Data (memory) access on



	GPU (GTX580)	CPU (i7-2600)	
Peak memory bandwidth [spec] <sup>1)</sup>	179 GB/s	21 GB/s	Upper bound for:
Peak memory bandwidth [measured] <sup>2)</sup>	153 GB/s	18 GB/s	
Random access [measured] <sup>2)</sup>	6.6 GB/s	0.8 GB/s	Probe
Compare and swap [measured] <sup>3)</sup>	4.6 GB/s	0.4 GB/s	Build HT

(1) Nvidia:  $192.4 \times 10^6 \text{ B/s} \approx 179.2 \text{ GB/s}$

(2) 64-bit accesses over 1 GB of device memory

(3) 64-bit compare-and-swap to random locations over 1 GB device memory

## GPU Hash Join Implementation (Summary)

### 1. Pin input tables

- Required for Build and Probe table, done by the CPU
- Only pinned CPU memory is accessible by the GPU
- “GPU direct” now allows to read directly from network/storage devices ...

### 2. Allocate memory for HT

- CPU handles memory allocation of GPU memory
- This is supposed to change with the next GPU generation ...

### 3. Build HT

- GPU reads build table (T1) sequentially from pinned CPU memory
- GPU creates HT (open addressing) in GPU memory
- Collisions are handled using atomic compare-and-swap

### 4. Probe HT

- GPU reads probe table (T2) sequentially from CPU memory
- GPU probes hash table (in GPU memory) and writes results to CPU memory

### 5. Cleanup

- free GPU memory
- Unpin input tables

## GPU Hash Join – Build HT

- GPU reads build table ( $T_1$ ) sequentially from pinned CPU memory
- GPU creates HT (open addressing) in GPU memory
- Collisions are handled using atomic compare-and-swap

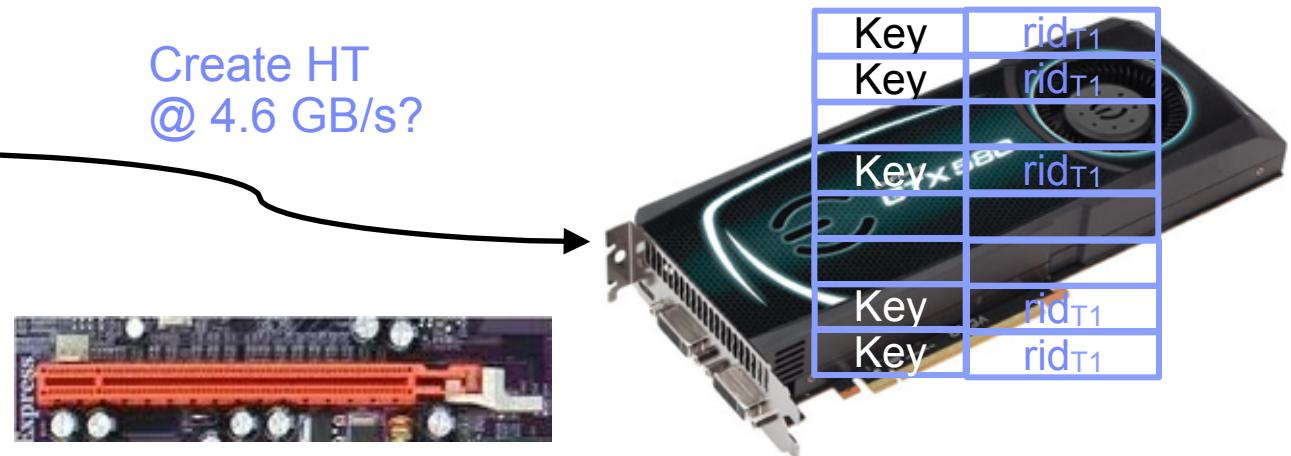
**Build table ( $T_1$ )**

key	$rid_{T_1}$
key	$rid_{T_1}$

Create HT  
@ 4.6 GB/s?

**Hash table**

Key	$rid_{T_1}$
Key	$rid_{T_1}$



## Build HT – Memory Management & Function call

```
// register input table
// 32-bit key + 32-bit rid are stored as a single 64-bit value
unsigned long long int* buildT;
cudaHostRegister(T1, num_tuples*2*sizeof(int), cudaHostRegisterMapped);
cudaHostGetDevicePointer(&buildT, T1, 0);

// make space for hash table
unsigned long long int* HT;
int HT_rows = 4 * num_tuples;
cudaMalloc(&HT, HT_rows * sizeof(int));
cudaMemSet(HT, 0, HT_rows * sizeof(int));

// call device function
dim3 Dg = dim3(16, 0, 0);
dim3 Db = dim3(512, 0, 0);
gpuCreateHashtable <<< Dg, Db >>>(builtT, num_tuples,
                                         HT, HT_rows);
```

## Build HT – Local variables

```
__global__ static void gpuCreateHashtable(unsigned long long int *buildT,
                                         int num_tuples,
                                         unsigned long long int *HT,
                                         int HT_rows) {

    int insert_loc;                                // insert location for tuple
    int tupleID;                                   // iterator for the build table
    int cas_result;                                // HT was initialized with 0, i.e.
                                                   // if insert was successful then
                                                   // cas_result = 0
    int hash_mask = HT_rows - 1;                   // LSB hash mask (for powers of 2!)
    unsigned long long int buildT_cache;           // register cache for a build table
    int key;                                       // key extracted from build table
```

## Build HT – Outline

```
// Iterate through the tuples of the build table and insert them into the
// hash table

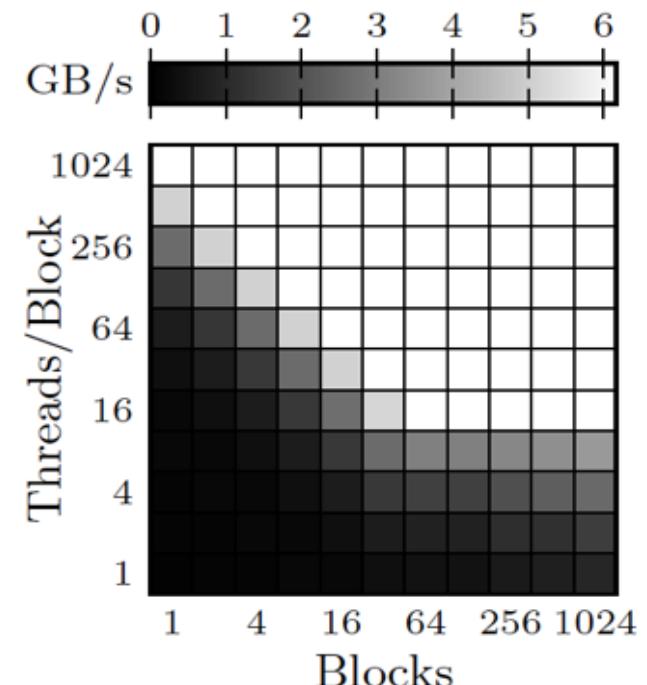
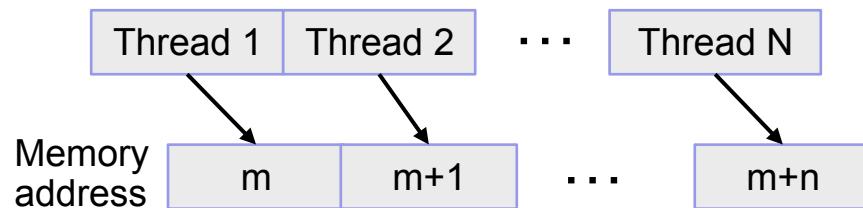
for (tupleID = blockIdx.x*blockDim.x+threadIdx.x;
      tupleID < num_tuples;
      tupleID += blockDim.x*gridDim.x) {
    /* 1) Cache the build table entry (key,rid) in a register
     * 2) Apply hash function (LSB) to key to determine insert position
     * 3) Starting from the insert position, scan for the next available
     *     slot
     * 4) Atomically insert the entry into the hash table
  */
```

## Build HT – Memory Access

Read build table from host memory

```
for (tupleID = blockIdx.x*blockDim.x+threadIdx.x;
     tupleID < num_tuples;
     tupleID += blockDim.x*gridDim.x) {
    buildT_cache = buildT[tupleID];
```

- Ideal memory access pattern is coalesced memory access
  - Threads of a block/warp access consecutive memory addresses



- Same applies to ZCA to host(main) memory
  - Coalesced access up to 6.2 GB/s
  - Random = faux pas !

## Build HT – Core Loop

```
for (tupleID = blockIdx.x*blockDim.x+threadIdx.x;
      tupleID < num_tuples;
      tupleID += blockDim.x*gridDim.x)
{
    cas_result = 42; // answer to everything ;-)
    // 1) Cache the build table entry (key,rid) in a register
    buildT_cache = buildT[tupleID];
```

## Build HT – Core Loop

```
for (tupleID = blockIdx.x*blockDim.x+threadIdx.x;
     tupleID < num_tuples;
     tupleID += blockDim.x*gridDim.x)
{
    cas_result = 42; // answer to everything ;-)
    // 1) Cache the build table entry (key,rid) in a register
    buildT_cache = buildT[tupleID];

    // 2) Apply LSB hash to key to determine insert position
    //     Little endian: <key,rid> becomes <rid,key> in the register
    key = (int)(buildT_cache & 0xFFFFFFFF); // key in the lower half
    insert_loc = key & hash_mask;
```

## Build HT – Core Loop

```
for (tupleID = blockIdx.x*blockDim.x+threadIdx.x;
     tupleID < num_tuples;
     tupleID += blockDim.x*gridDim.x)
{
    cas_result = 42; // answer to everything ;-)

    // 1) Cache the build table entry (key,rid) in a register
    buildT_cache = buildT[tupleID];

    // 2) Apply LSB hash to key to determine insert position
    //      Little endian: <key,rid> becomes <rid,key> in the register
    key = (int)(buildT_cache & 0xFFFFFFFF); // key in the lower half
    insert_loc = key & hash_mask;

    // 3) From insert position scan for the next available slot (0) to
    //      avoid repeated atomic compare-and-swap ($$$)

    while (HT[insert_loc] != 0)
        insert_loc = ++insert_loc & hash_mask;
```

## Build HT – Core Loop

```
// 1) Cache the build table entry (key,rid) in a register
buildT_cache = buildT[tupleID];

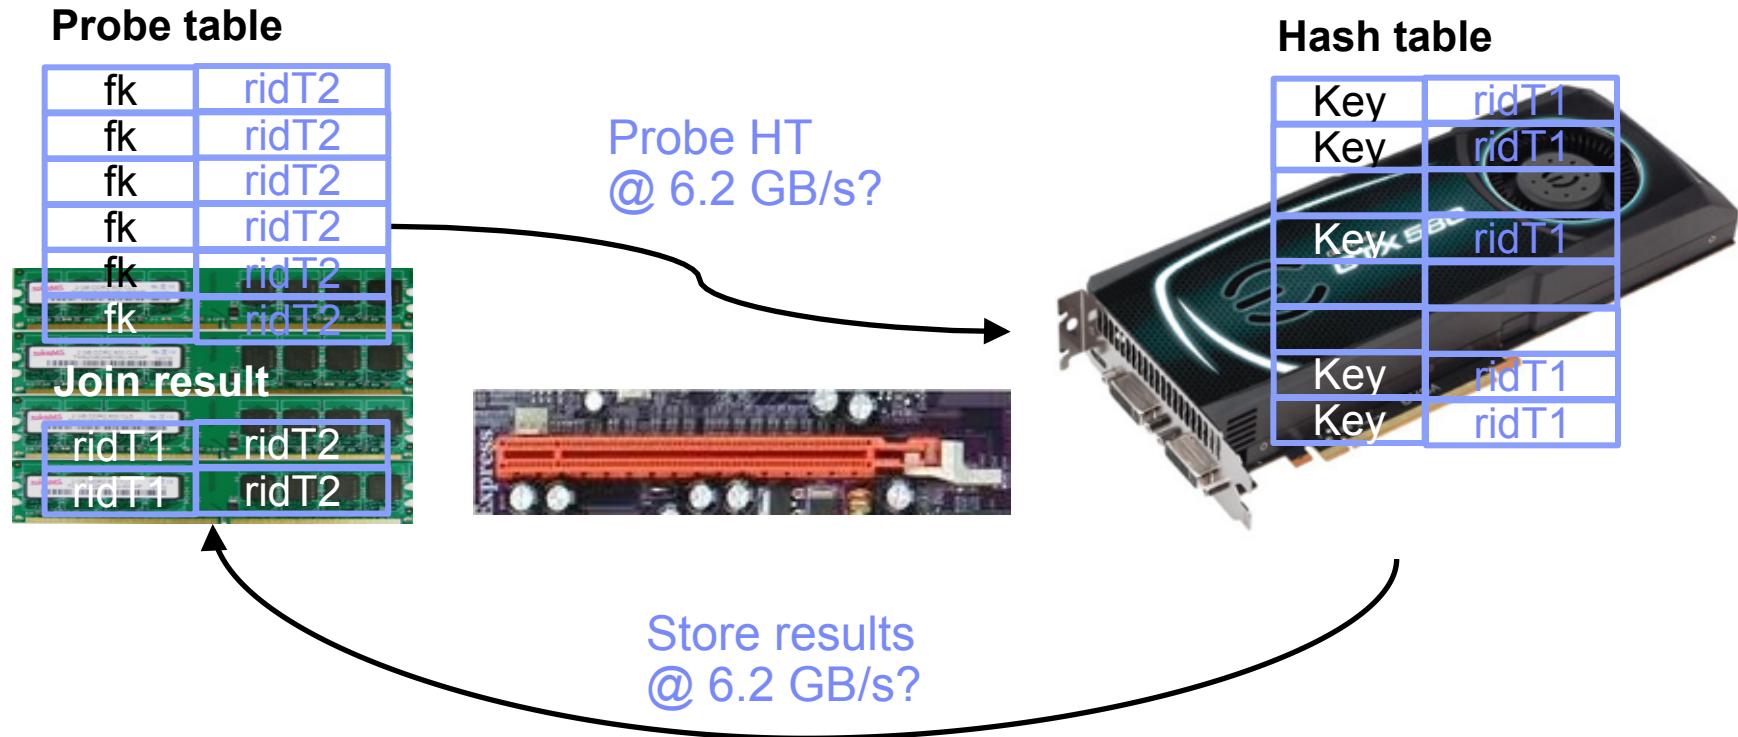
// 2) Apply LSB hash to key to determine insert position
//     Little endian: <key,rid> becomes <rid,key> in the register
key = (int) (buildT_cache & 0xFFFFFFFF); // key in the lower half
insert_loc = key & hash_mask;

// 3) From insert position scan for the next available slot (0) to
//     avoid repeated atomic compare-and-swap ($$$)
while (HT[insert_loc] != 0)
    insert_loc = ++insert_loc & hash_mask;

// 4) Atomically insert entry into the hash table
while(cas_result != 0) {
    cas_result = atomicCAS(&(HT[insert_loc]), 0, buildT_cache);
    insert_loc = ++insert_loc & hash_mask;
}
```

## GPU Hash Join – Probe HT

- GPU reads probe table (T2) sequentially from CPU memory
- GPU probes hash table (in GPU memory) and writes results to CPU memory



## Probe HT – Memory Management & Function call

```
// register input table
// 32-bit key + 32-bit rid are stored as a single 64-bit value
unsigned long long int* probeT;
cudaHostRegister(T2, num_tuples*2*sizeof(int), cudaHostRegisterMapped);
cudaHostGetDevicePointer(&probeT, T2, 0);

// make space for results
unsigned long long int* resG;
cudaHostAlloc(&resG, 2 * num_tuples * sizeof(int));

// result index
__device__ int gpu_result_index;
cudaMemcpyToSymbol(gpu_result_index, &null, sizeof(int));

// call device function
dim3 Dg = dim3(16,0,0);
dim3 Db = dim3(512,0,0);
gpuProbe <<< Dg, Db >>>(probeT, HT, resG, num_tuples, HT_rows);
```

## Probe HT – Local Variables

```
__global__ static void gpuProbe(unsigned long long int* probeT,  
                                unsigned long long int* HT,  
                                unsigned long long int* resG,  
                                int probeT_rows, int HT_rows)  
{  
int probeT_key;                      // the probe table key used for a probe  
int HT_idx;                          // hash table location the probe lead to  
int HT_key;                           // the key found at the hash table  
                                      // location of hashtable_idx  
int tupleID;                          // iterator for the probe table  
int hash_mask = HT_rows - 1;           // LSB hash mask  
int result_insert_position;          // index to the result, shared by ALL  
                                      // threads (atomic insert)  
unsigned long long int probeT_cache; // register cache for probe table  
unsigned long long int HT_cache;    // register cache for hash table
```

## Probe HT – Outline

```
// Iterate through the tuples of the probe table and
for (tupleID=blockIdx.x*blockDim.x+threadIdx.x;
    tupleID < probeT_rows;
    tupleID+=blockDim.x*gridDim.x) {
    /* 1) Cache the fact table entry (key,rid) in a register & extract
     *      the fact table key
     * 2) Apply the hash function to the key to determine the location
     *      in the hash table
     * 3) Probe the hash table and cache the entry (key,rid) in a
     *      register
     * 4) Scan the hash table for more matching keys until we hit an
     *      empty (0) position
    */
}
```

## Probe HT – Core Loop

```
for (tupleID=blockIdx.x*blockDim.x+threadIdx.x;
      tupleID < probeT_rows;
      tupleID+=blockDim.x*gridDim.x) {
    // 1) Cache the fact table entry (key,rid) in a register
    probeT_cache = probeT[tupleID];
    // Extract the fact table key
    // Little endian: <key,rid> becomes <rid,key> in the register
    probeT_key = (int)(probeT_cache & 0xFFFFFFFF); // key in lower half
```

## Probe HT – Core Loop

```
for (tupleID=blockIdx.x*blockDim.x+threadIdx.x;
     tupleID < probeT_rows;
     tupleID+=blockDim.x*gridDim.x) {
    // 1) Cache the fact table entry (key,rid) in a register
    probeT_cache = probeT[tupleID];
    // Extract the fact table key
    // Little endian: <key,rid> becomes <rid,key> in the register
    probeT_key = (int)(probeT_cache & 0xFFFFFFFF); // key in lower half

    // 2) Hash the key to determine the location in the hash table
    HT_idx = probT_key & hash_mask;
```

## Probe HT – Core Loop

```
for (tupleID=blockIdx.x*blockDim.x+threadIdx.x;
     tupleID < probet_rows;
     tupleID+=blockDim.x*gridDim.x) {
    // 1) Cache the fact table entry (key,rid) in a register
    probet_cache = probet[tupleID];
    // Extract the fact table key
    // Little endian: <key,rid> becomes <rid,key> in the register
    probet_key = (int)(probt_cache & 0xFFFFFFFF); // key in lower half

    // 2) Hash the key to determine the location in the hash table
    HT_idx = probt_key & hash_mask;

    // 3) Probe the hash table and cache the entry (key,rid) in a register
    HT_cache = HT[HT_idx];
```

## Probe HT – Core Loop

```
/* Scan open addressing hash table until we hit an empty(0) slot
 * 4.1) If keys match insert rids from the probe and hash table into
 *       the global result set
 * 4.2) Cache the next hash table entry and extract the key
 */
while (HT_cache != 0) {
    HT_key = (int) (HT_cache & 0xFFFFFFFF);
    if (probeT_key == HT_key) {
        // determine position in global result set
        result_insert_position = atomicAdd(&gpu_result_index, 1);
        // insert result=<rid,rid>
        // rids are both in the upper half of the register caches,
        // so we need to shift one of them (hashtable cache) down
        resG[result_insert_position] = (probeT_cache & 0xFFFFFFFF00000000)
                                       | (HT_cache >> 32);
    }
    HT_idx = ++HT_idx & hash_mask;
    HT_cache = HT[HT_idx];
}
```

## Retrieving result count & cleanup

```
...
// After GPU function completes
cudaDeviceSynchronize();
cudaMemcpyFromSymbol(rescount, gpu_result_index, sizeof(int)) ;

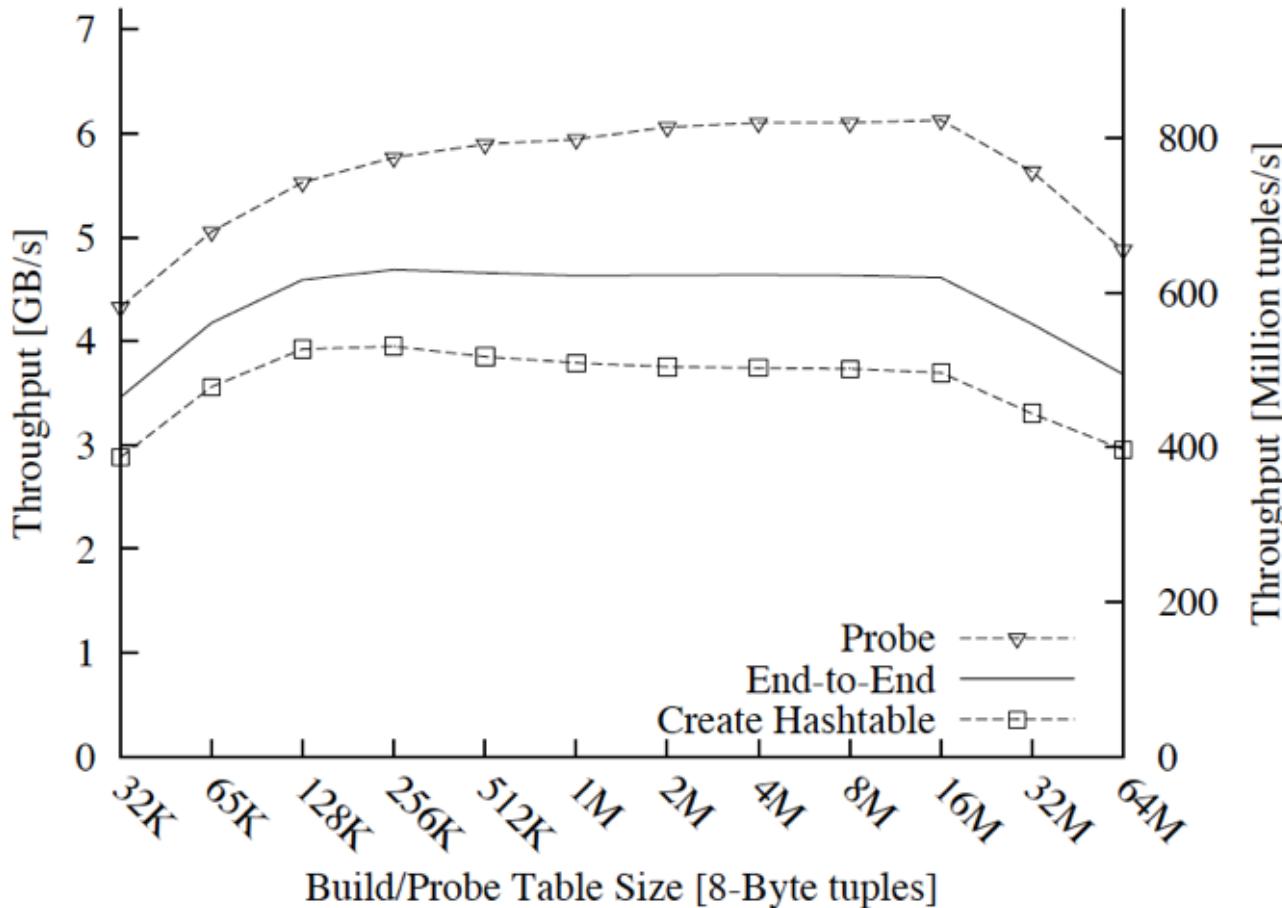
// clean up memory
cudaHostUnregister(T1) ;
cudaHostUnregister(T1) ;
cudaFree(HT) ;
...
```

## Throughput

- Join 2 equal size tables (16M rows) of 32-bit <key, row-ID> pairs (4+4 Byte)  
– Uniformly distributed randomly generated keys  
– 3% of the keys in the probe table have a match in the build table  
– Measuring End-to-End throughput, i.e. input tables & results in host memory

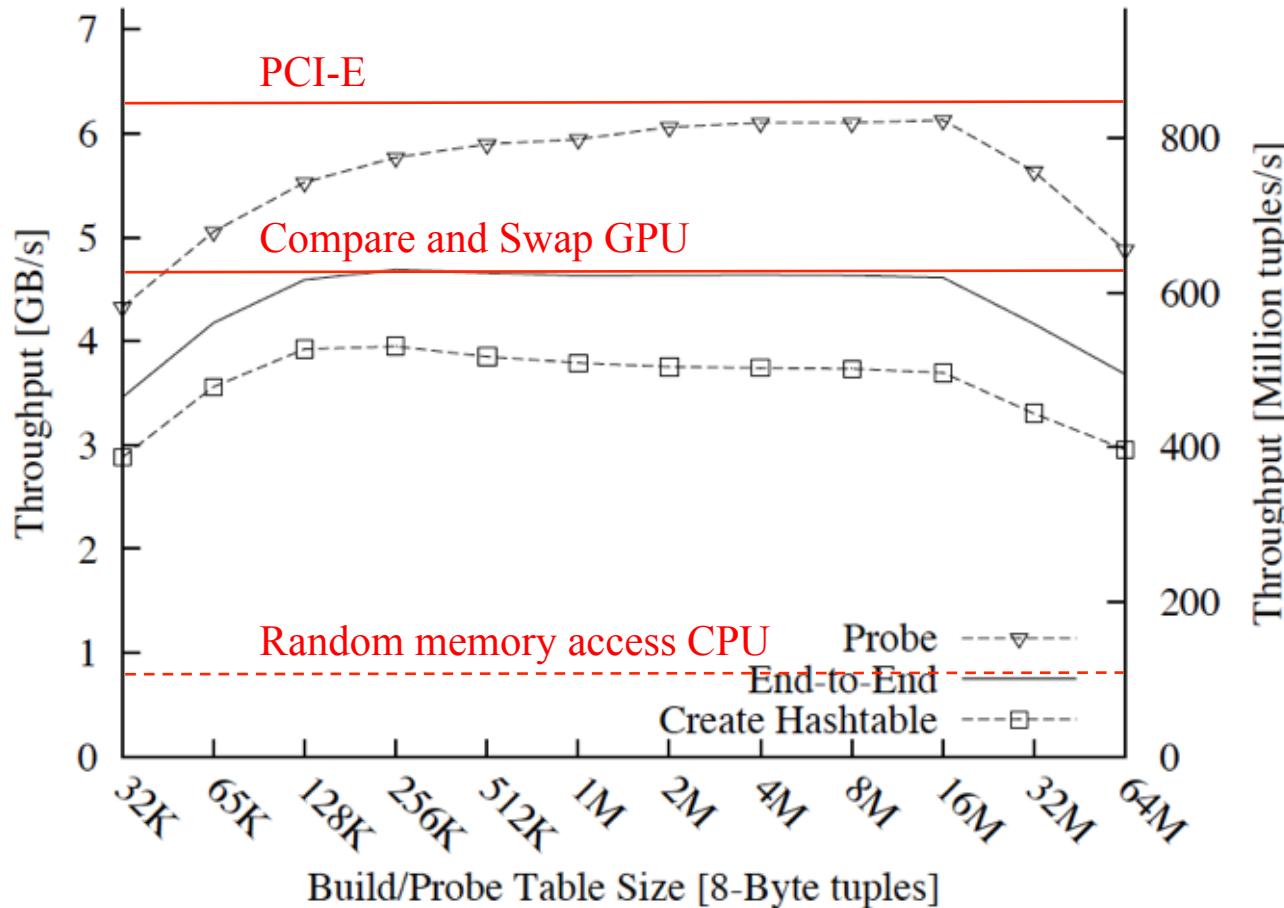
## Throughput

- Join 2 equal size tables (16M rows) of 32-bit <key, row-ID> pairs (4+4 Byte)
  - Uniformly distributed randomly generated keys
  - 3% of the keys in the probe table have a match in the build table
  - Measuring End-to-End throughput, i.e. input tables & results in host memory

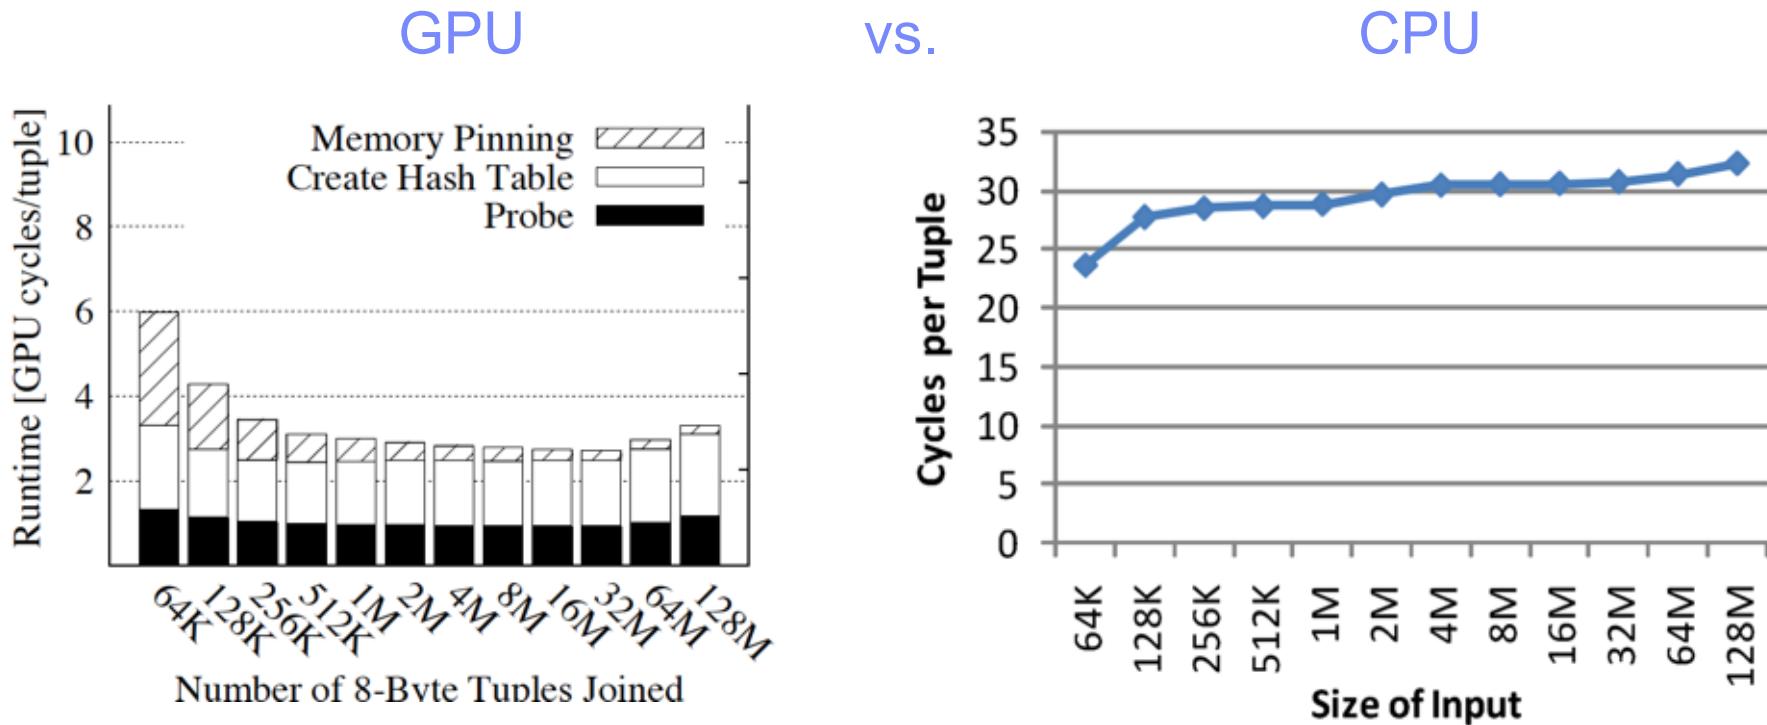


## Throughput

- Join 2 equal size tables (16M rows) of 32-bit <key, row-ID> pairs (4+4 Byte)
  - Uniformly distributed randomly generated keys
  - 3% of the keys in the probe table have a match in the build table
  - Measuring End-to-End throughput, i.e. input tables & results in host memory

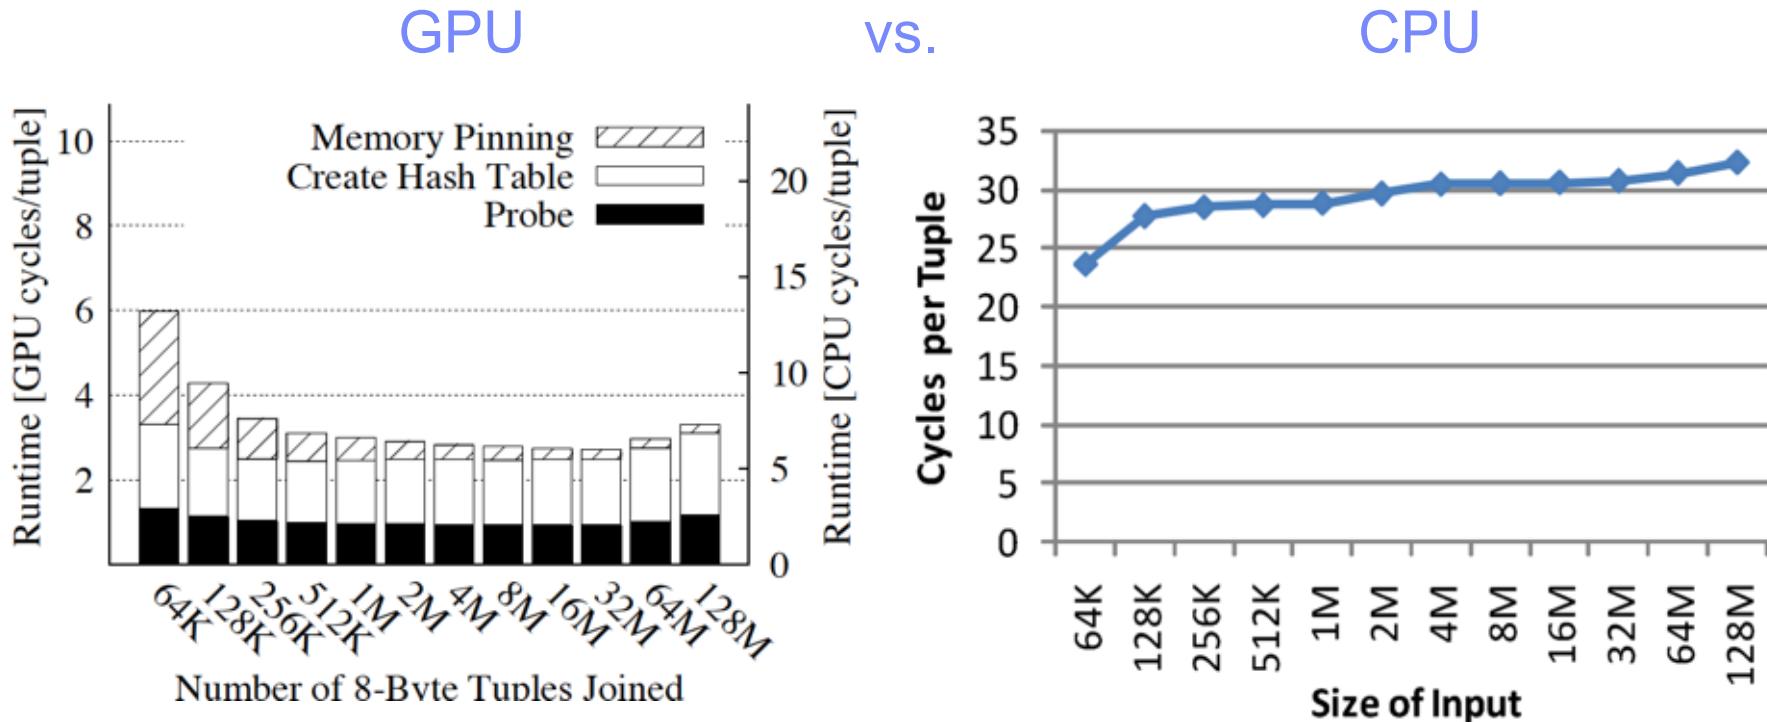


- Join 2 equal size tables (512K to 128M) of 32-bit <key, row-ID> pairs (4 + 4 Byte)
- Uniformly distributed randomly generated keys
- 3% of the probe keys have a match in the build table
- CPU implementation does not materialize results



<sup>2</sup> C. Kim, E. Sedlar, J. Chhugani, T. Kaldewey, A. Nguyen, A. Di Blas, V. Lee, N. Satish, P. Dubey. Sort vs. Hash revisited: fast join implementation on modern multi-core CPUs. VLDB 2009

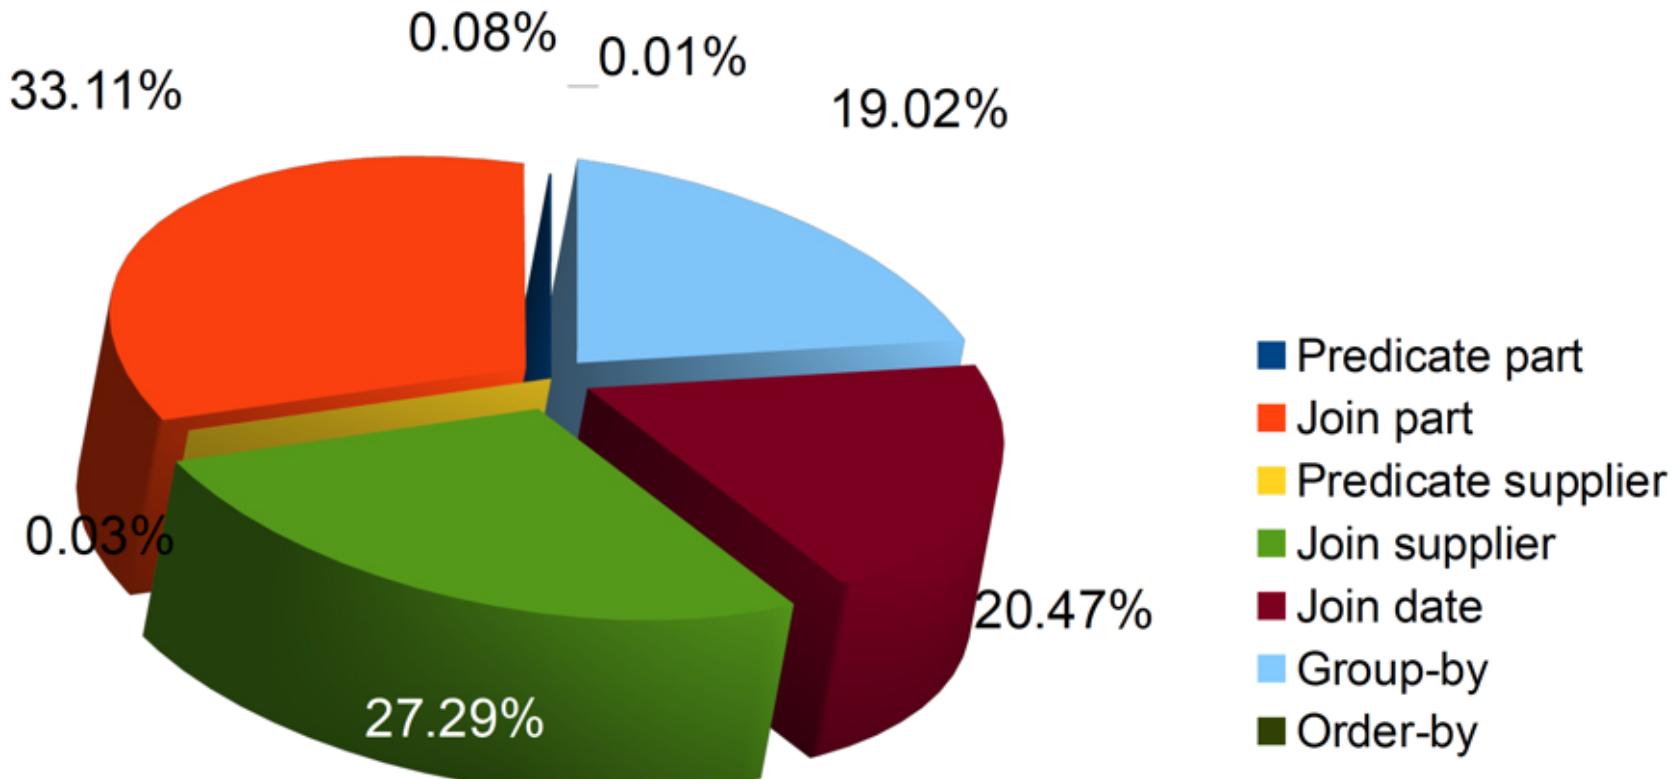
- Join 2 equal size tables (512K to 128M) of 32-bit <key, row-ID> pairs (4 + 4 Byte)
- Uniformly distributed randomly generated keys
- 3% of the probe keys have a match in the build table
- CPU implementation does not materialize results
- Cycles/tuple not a meaningful metric
  - depends on processor frequency, tuple size, ...



<sup>2</sup> C. Kim, E. Sedlar, J. Chhugani, T. Kaldewey, A. Nguyen, A. Di Blas, V. Lee, N. Satish, P. Dubey. Sort vs. Hash revisited: fast join implementation on modern multi-core CPUs. VLDB 2009

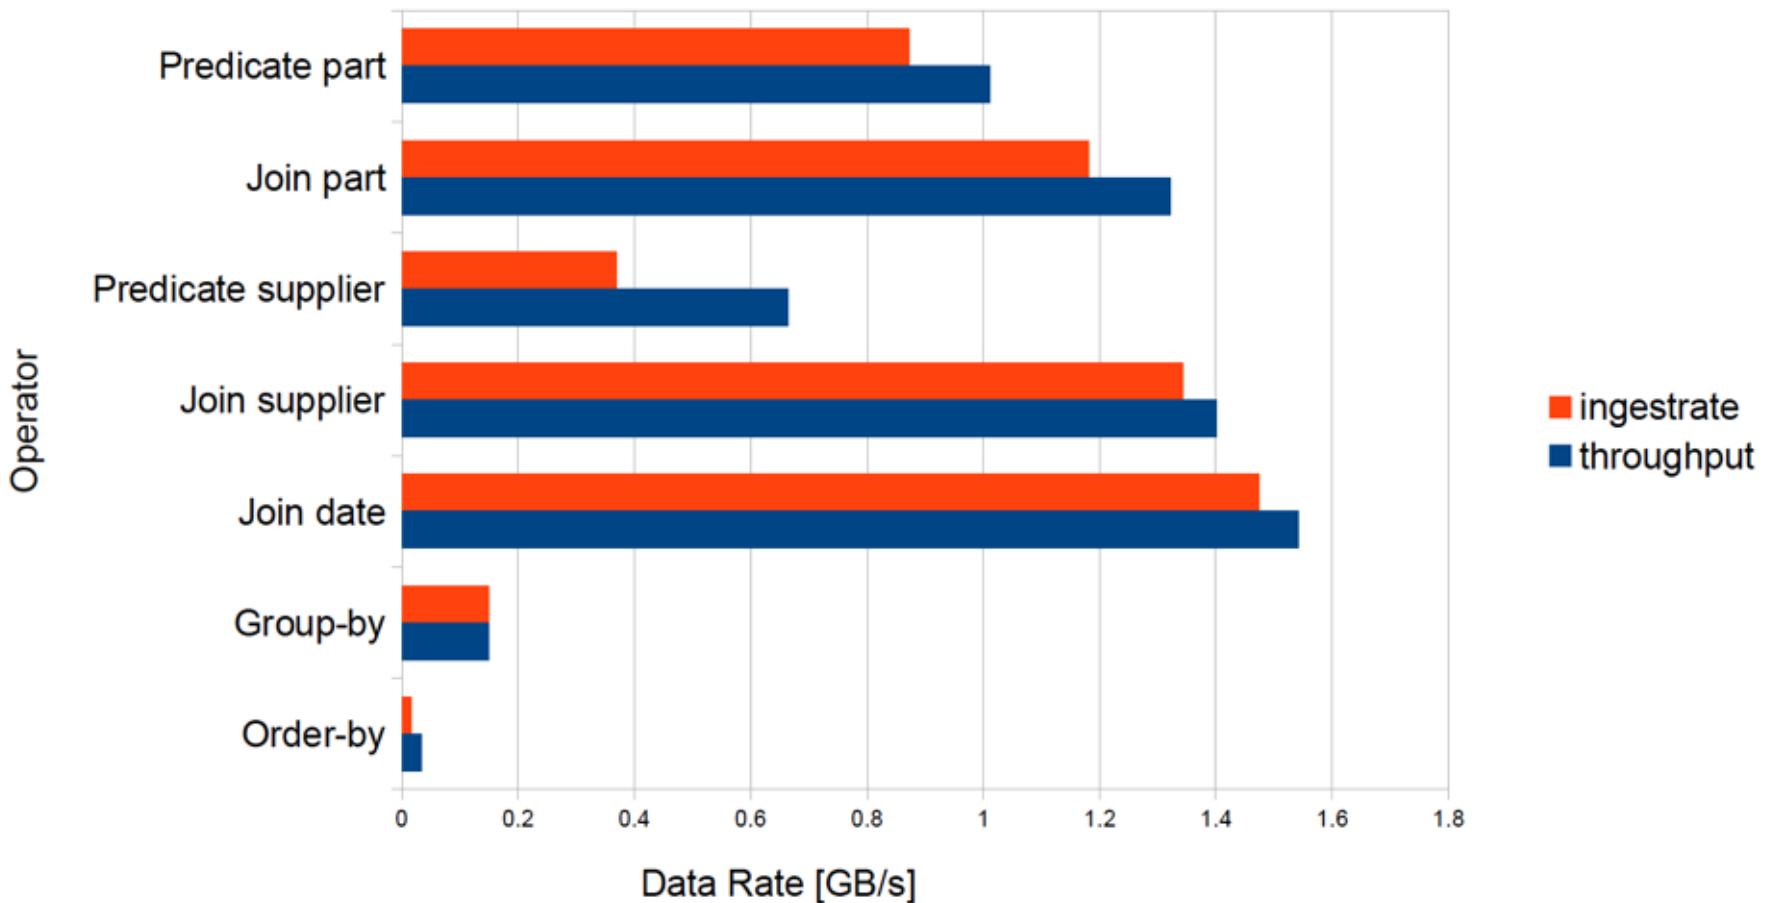
What's there beside join?

## Where does time go?



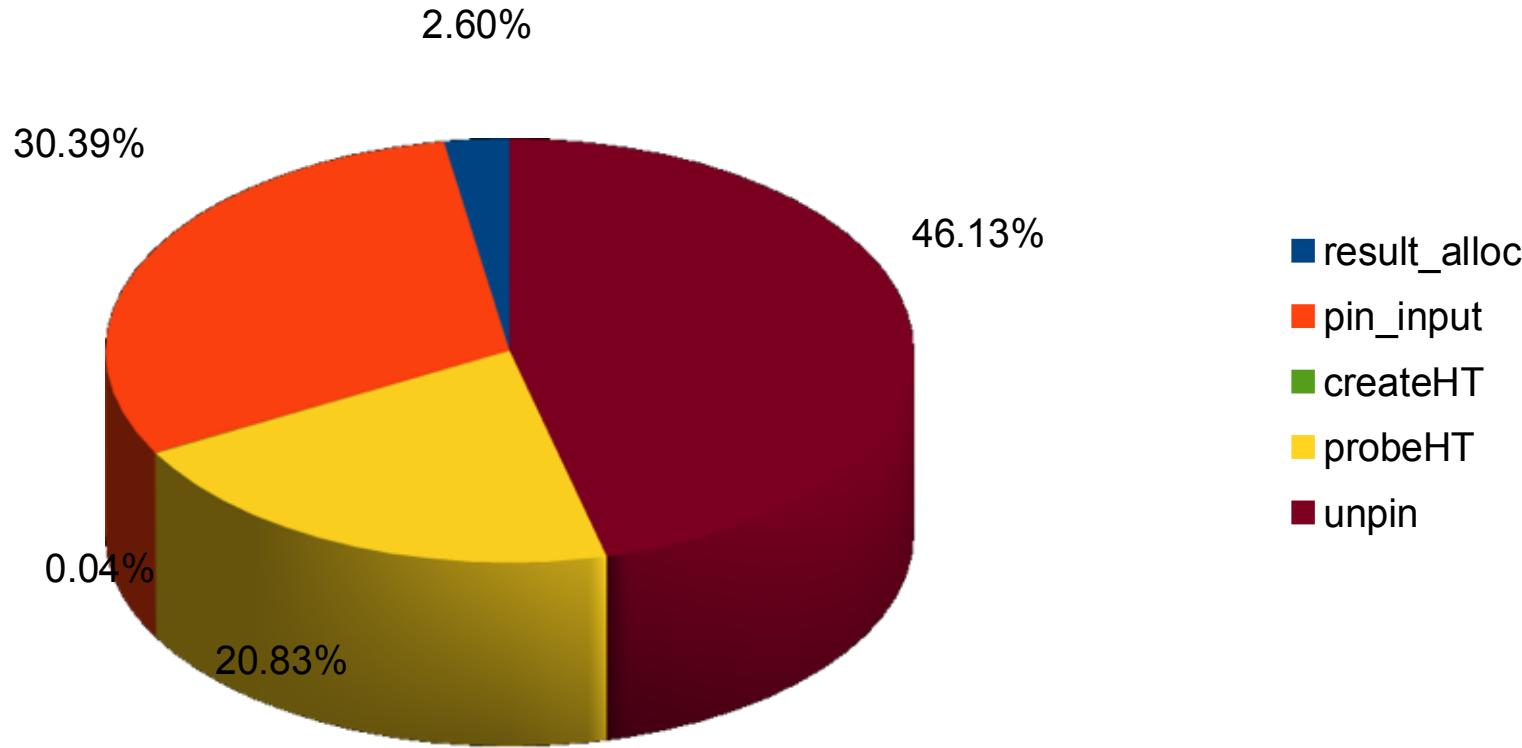
```
SELECT SUM(lo.revenue), d.year, p.brand
  FROM lineorder lo, part p, supplier s, date d
 WHERE p.category = 'MFGR#12' AND lo.partkey = p.partkey
   AND s.region = 'AMERICA' AND lo.supkey = s.supkey
   AND lo.orderdate = d.datekey
 GROUP BY d.year, p.brand
 ORDER BY d.year, p.brand;
```

## Operator throughput



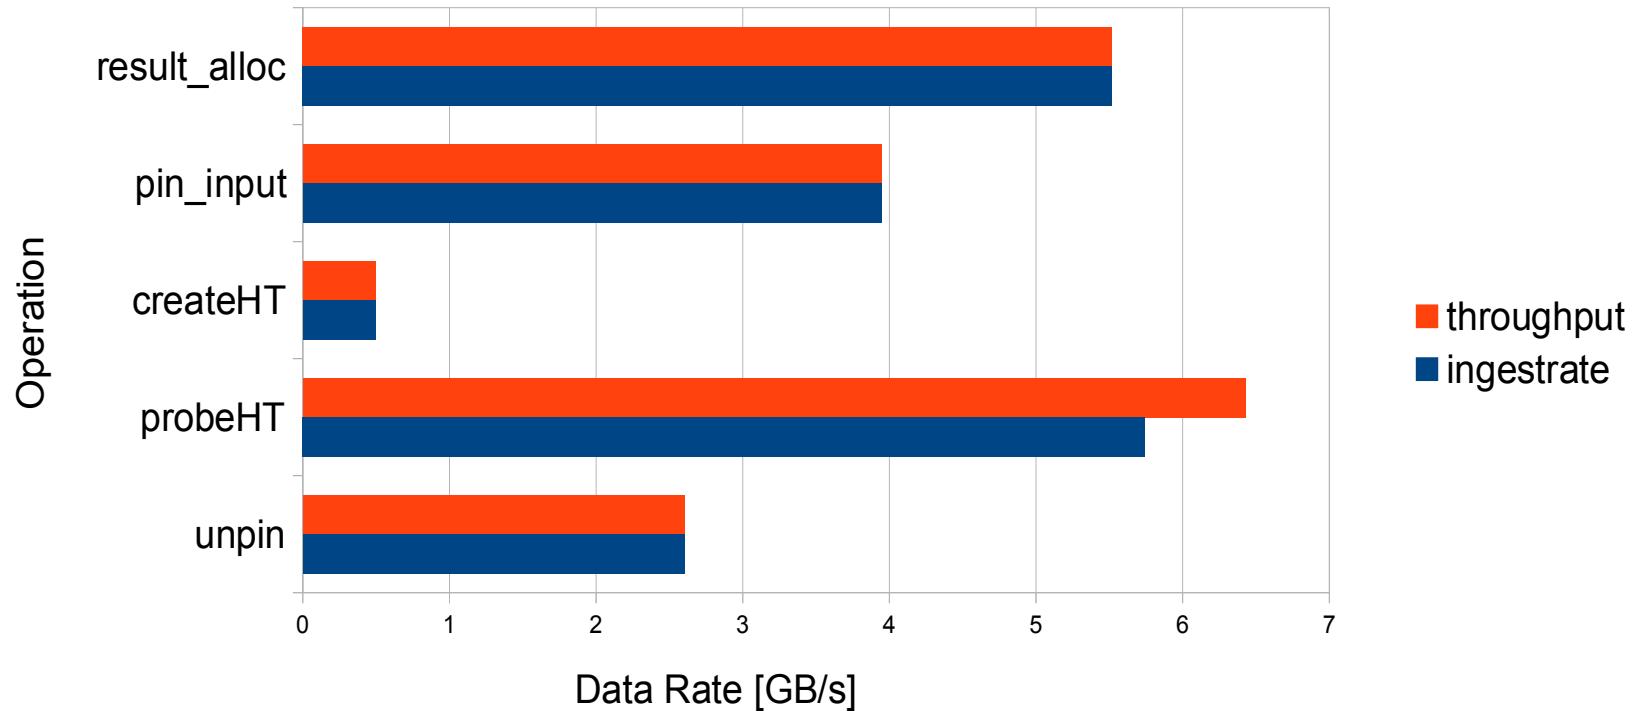
- Using a straight forward GPU implementation
  - Joins are running at < 1.5GB/s,  $\frac{1}{4}$  of the expected speed!
  - Where does time go?

## GPU Join – Where does time go?



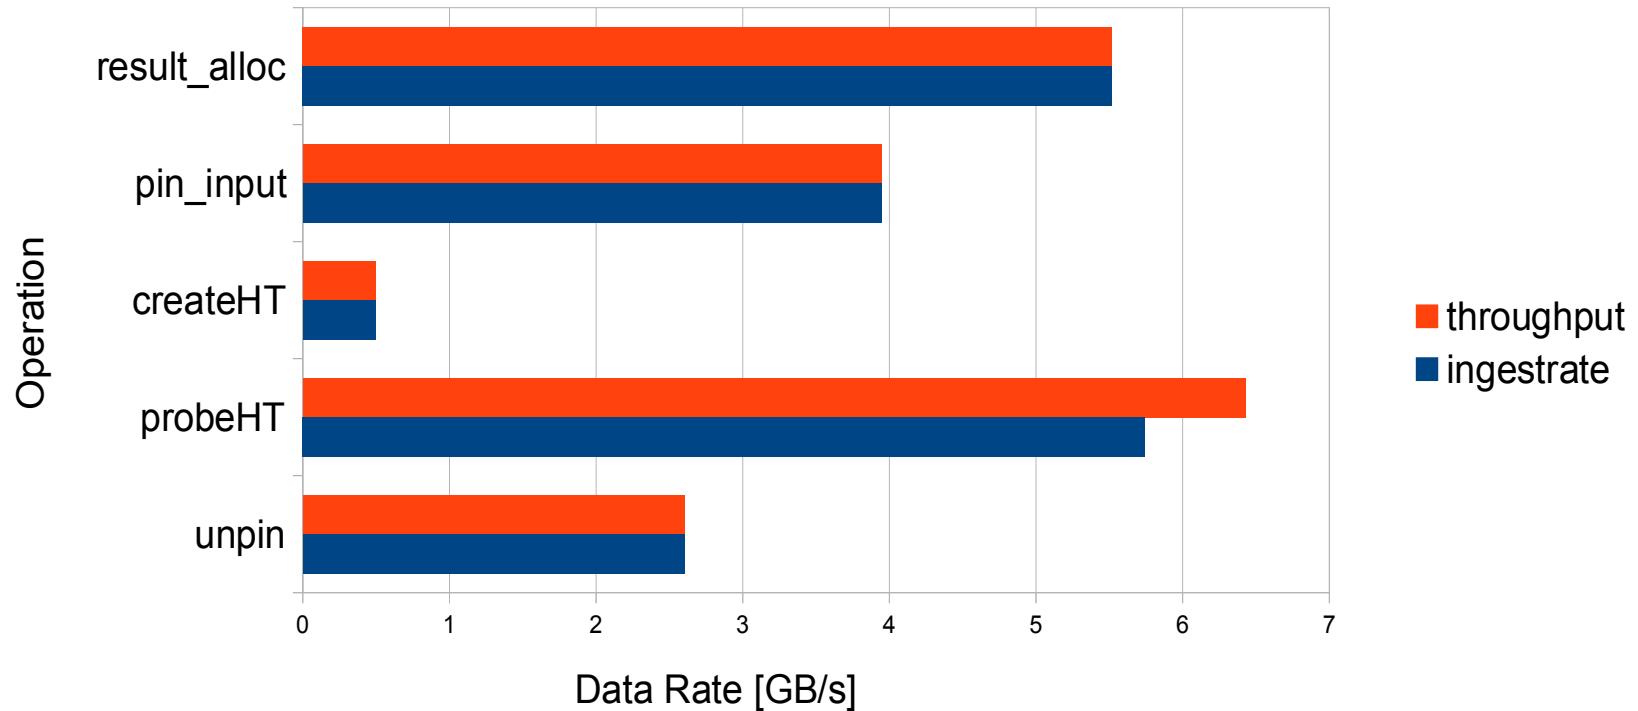
- < 21% of the runtime is spent on the actual join!
- Join lineorder & part has < 4% selectivity
- At SF 100 (100GB database) p.partey is 5.4 MB, lo.partkey is 2.3 GB
  - Need to pin & unpin 2.3 GB of lineorder data

## GPU Join – Where does time go?



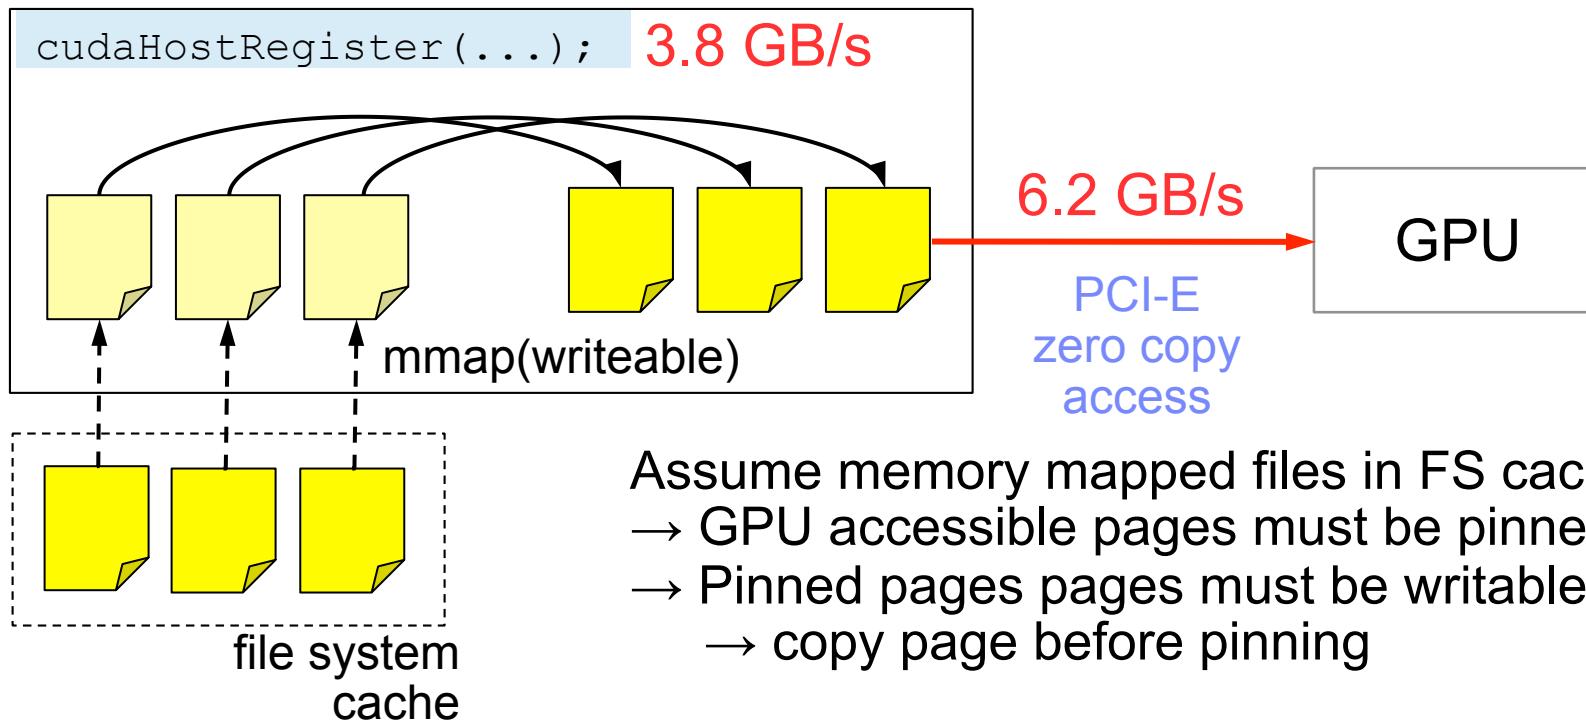
- Pinning/Unpinning large amounts of memory is inefficient and time consuming!

## GPU Join – Where does time go?

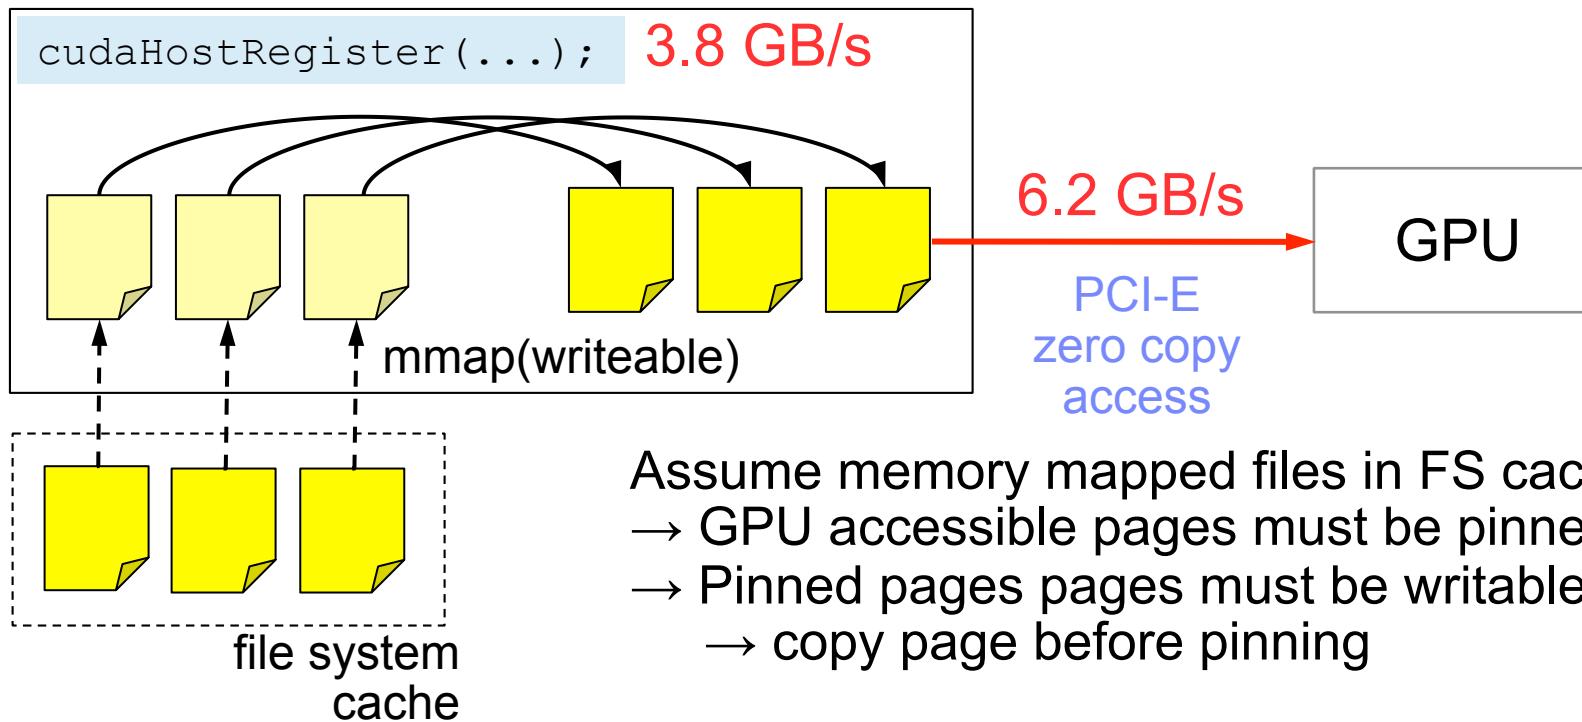


- Pinning/Unpinning large amounts of memory is inefficient and time consuming!
- All steps are sequential ... overlapping across operators is messy =(
- We could copy data chunk-wise into a (smaller) pinned buffer ...
- Since we are already at it, how do we get the data from the file system (cache)?

## Data flow – Current approach

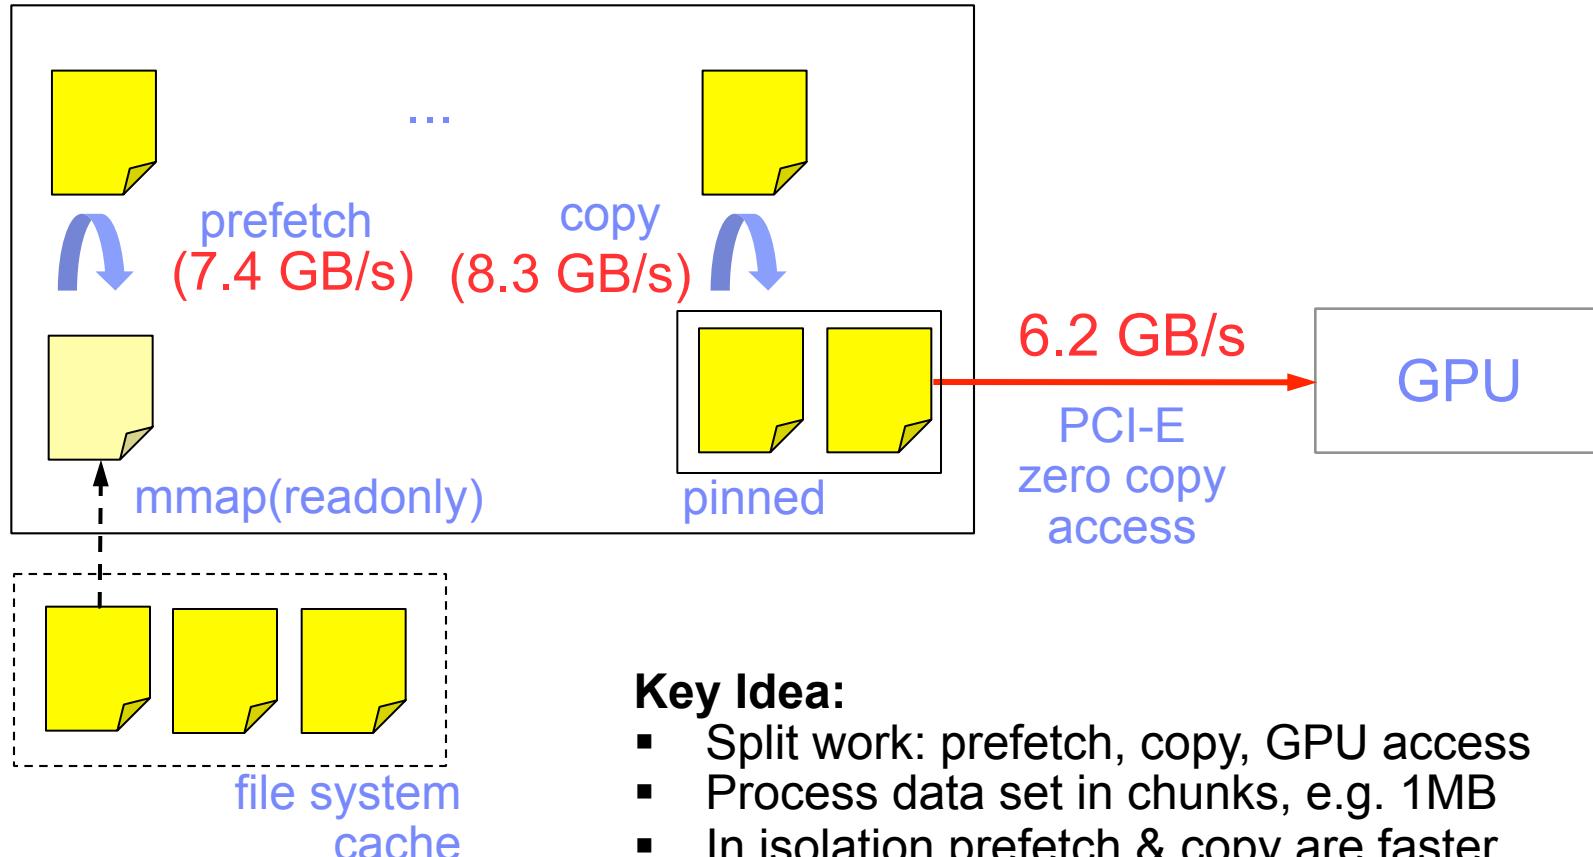


## Data flow – Current approach



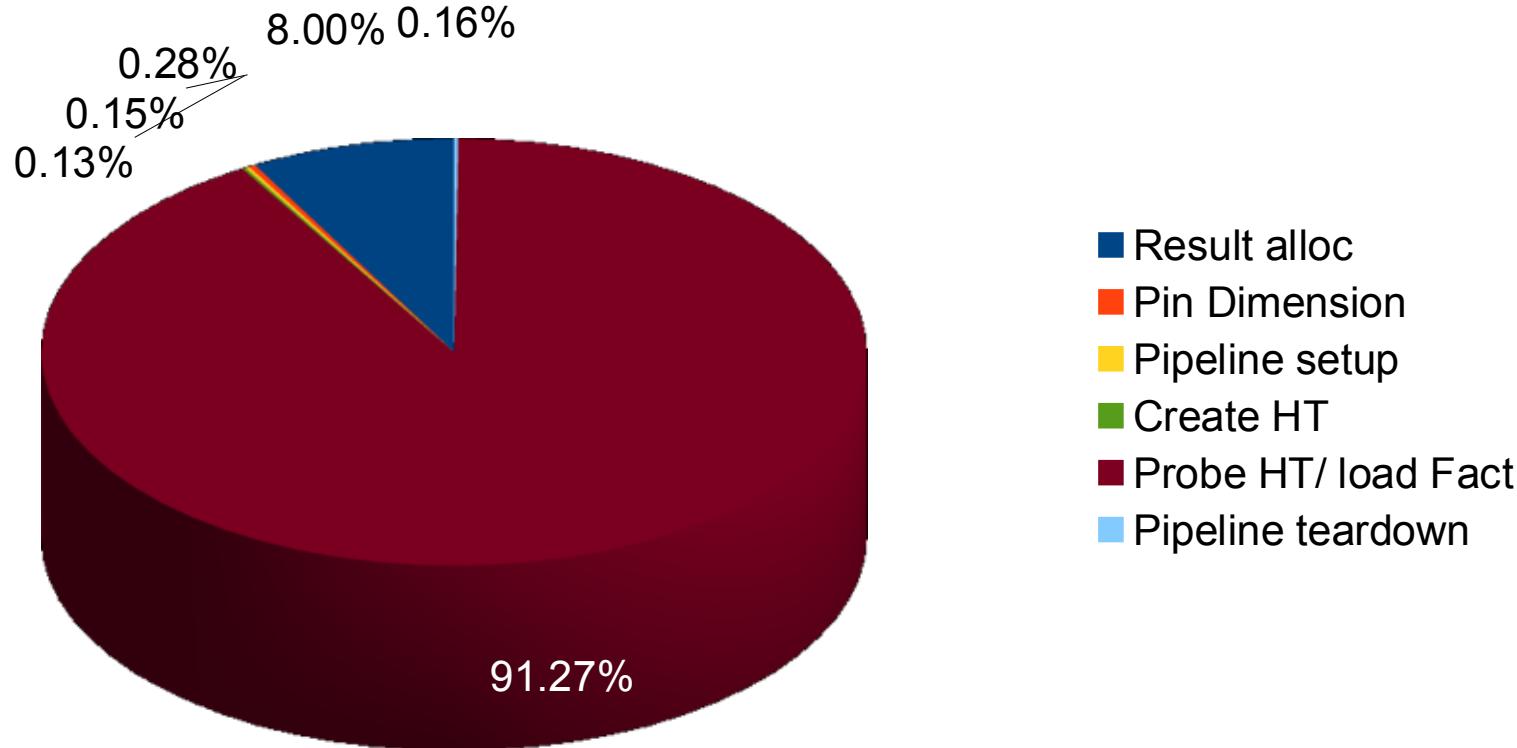
- Even overlapping query execution with pinning pages for next operator (join) leaves **pinning as a bottleneck!**
- What if we use 2 pre-allocated buffers of pinned memory:
  - Copy data into one of the pinned buffers
  - Meanwhile the GPU can work on the data in the other buffer

## Data flow: prefetch → memcpy → GPU access



## Join with 3-stage pipeline

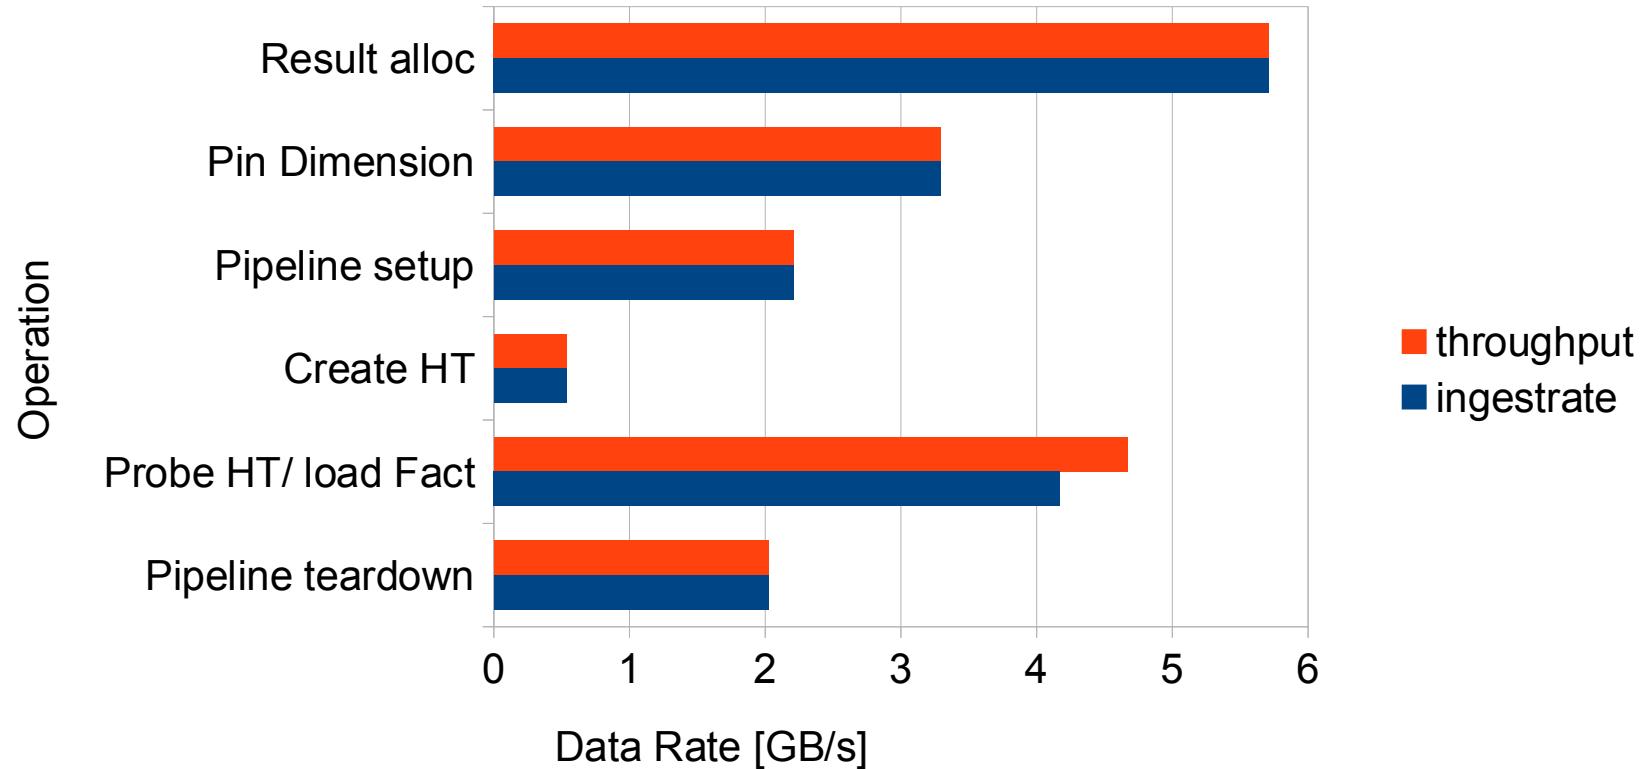
- 2x 1MB buffers, ~2300 Kernel invocations



```
SELECT SUM(lo.revenue), d.year, p.brand FROM lineorder lo, date d, part p, supplier s  
WHERE lo.orderdate = d.datekey AND lo.partkey = p.partkey AND lo.supkey = s.supkey  
AND p.category = 'MFGR#12' AND s.region = 'AMERICA'  
GROUP BY d.year, p.brand ORDER BY d.year, p.brand
```

## Join with 3-stage pipeline

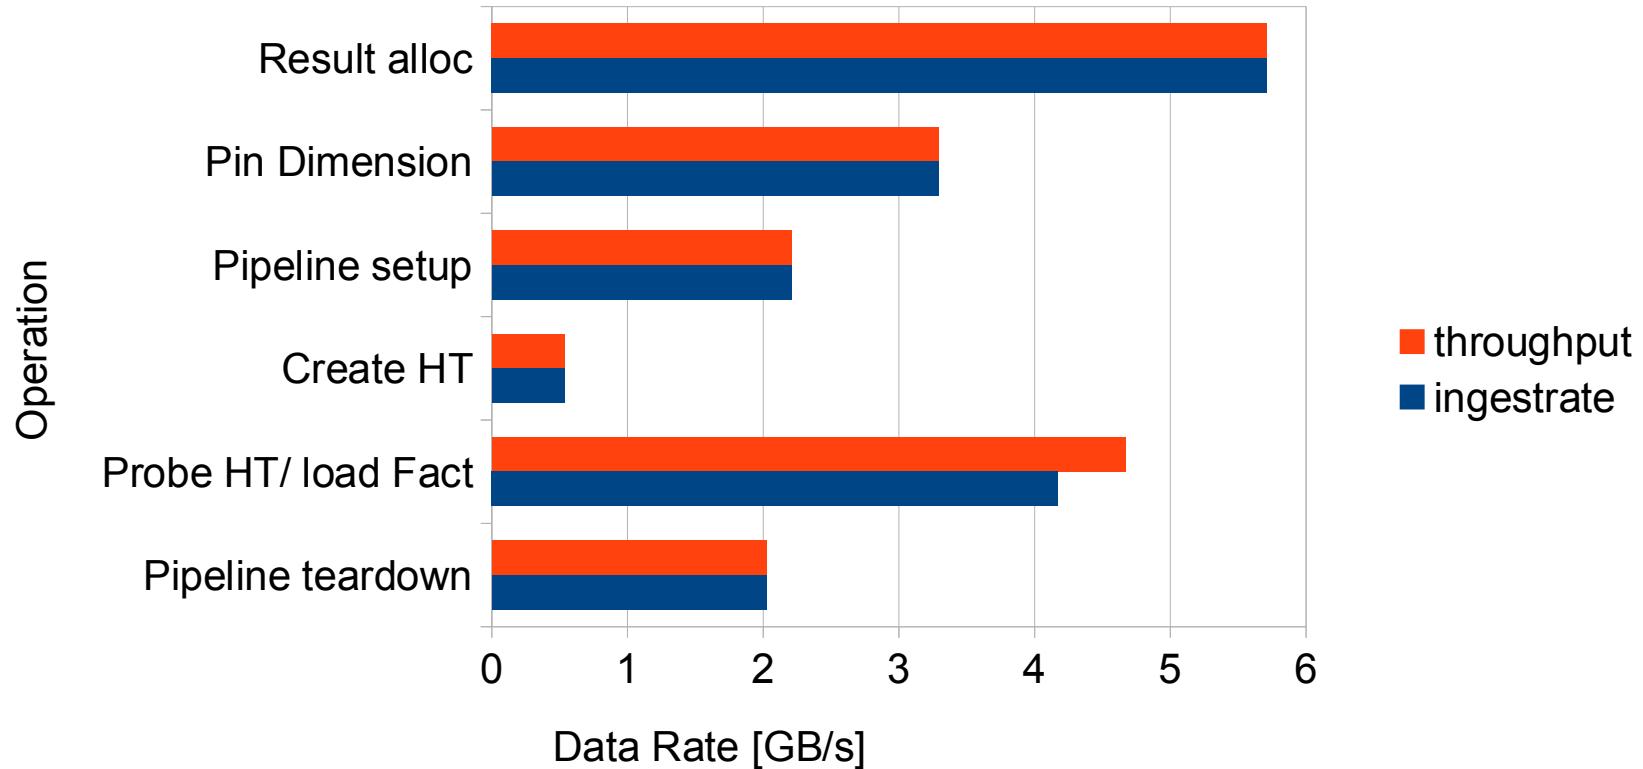
- Join lineorder & part using 2x 1MB buffers yields > 4GB/s overall throughput



- Other joins (with supplier and date) exhibit similar performance

## Join with 3-stage pipeline

- Join lineorder & part using 2x 1MB buffers yields > 4GB/s overall throughput



- Other joins (with supplier and date) exhibit similar performance
- Group-by operator is quite similar to join, i.e. requires a hash table and an atomic add and also achieves similar performance
- Accelerating other operators is not worthwhile ...

## Agenda

- GPU search
  - Reminder: Porting CPU search
  - Back to the drawing board:
    - P-ary search
    - Experimental evaluation
    - Why it works
- Building a complete data warehouse runtime with GPU support
  - From a query to operators – what to accelerate?
  - What are the bottlenecks/limitations
- Maximizing data path efficiency
  - Extremely fast storage solution
  - Storage to host to device
- Putting it all together
  - Prototype demo

## 6GB/s Storage Subsystem ?



- According to OCZ spec a Revodrive3 x2 can deliver 1.5 GB/s per card

## 6GB/s Storage Subsystem ?



- According to OCZ spec a Revodrive3 x2 can deliver 1.5 GB/s per card
- Can we “just stripe” the data across 4 cards and achieve 6 GB/s?

## 6GB/s Storage Subsystem ?

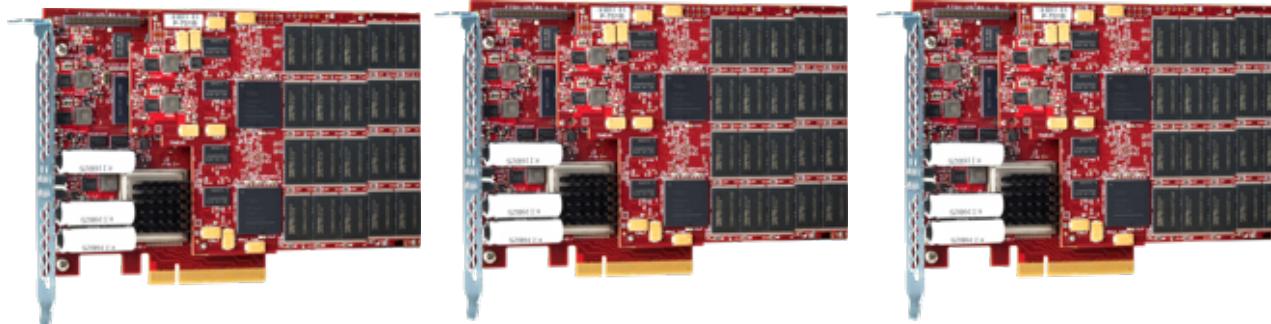


- According to OCZ spec a Revodrive3 x2 can deliver 1.5 GB/s per card
- Can we “just stripe” the data across 4 cards and achieve 6 GB/s?
  - Linux tools, i.e. mdraid + ext, max 2 GB/s
  - IBM GPFS with striping and heavy prefetching(72 threads) achieves 3 GB/s
  - SSD controllers on commodity SSDs use compression to improve throughput

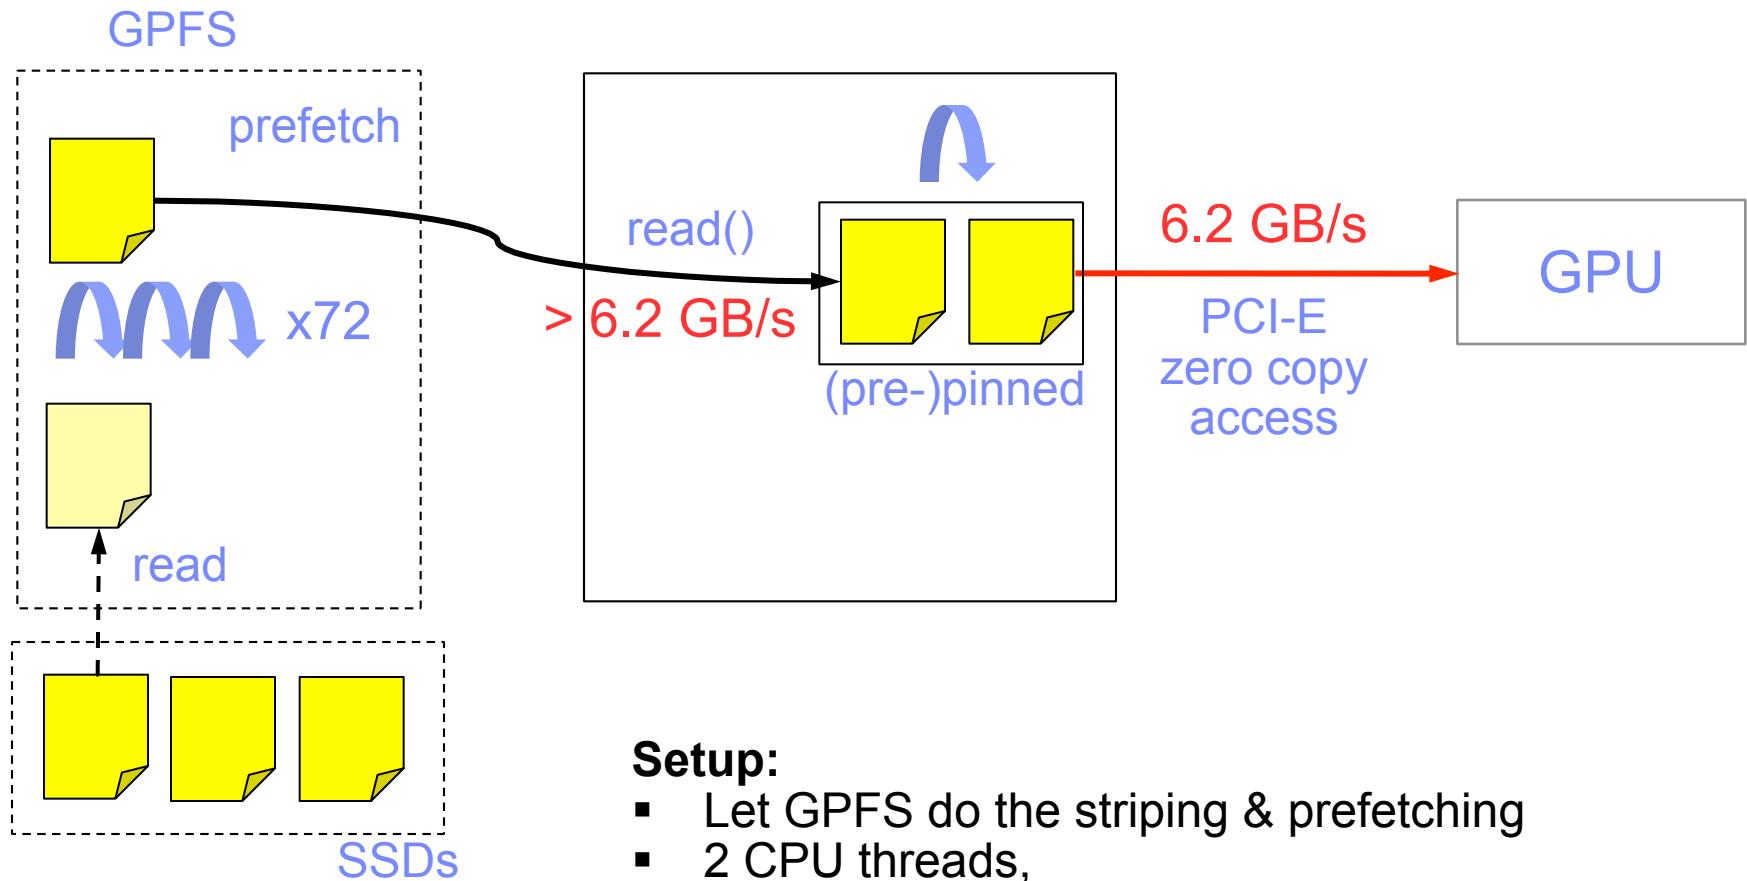
## 6GB/s Storage Subsystem ?



- According to OCZ spec a Revodrive3 x2 can deliver 1.5 GB/s per card
- Can we “just stripe” the data across 4 cards and achieve 6 GB/s?
  - Linux tools, i.e. mdraid + ext, max 2 GB/s
  - IBM GPFS with striping and heavy prefetching(72 threads) achieves 3 GB/s
  - SSD controllers on commodity SSDs use compression to improve throughput
- Using 3 Texas Memory RamSan-70 and GPFS we get up to 7.5 GB/s =)



## Data flow: read → memcpy → GPU access

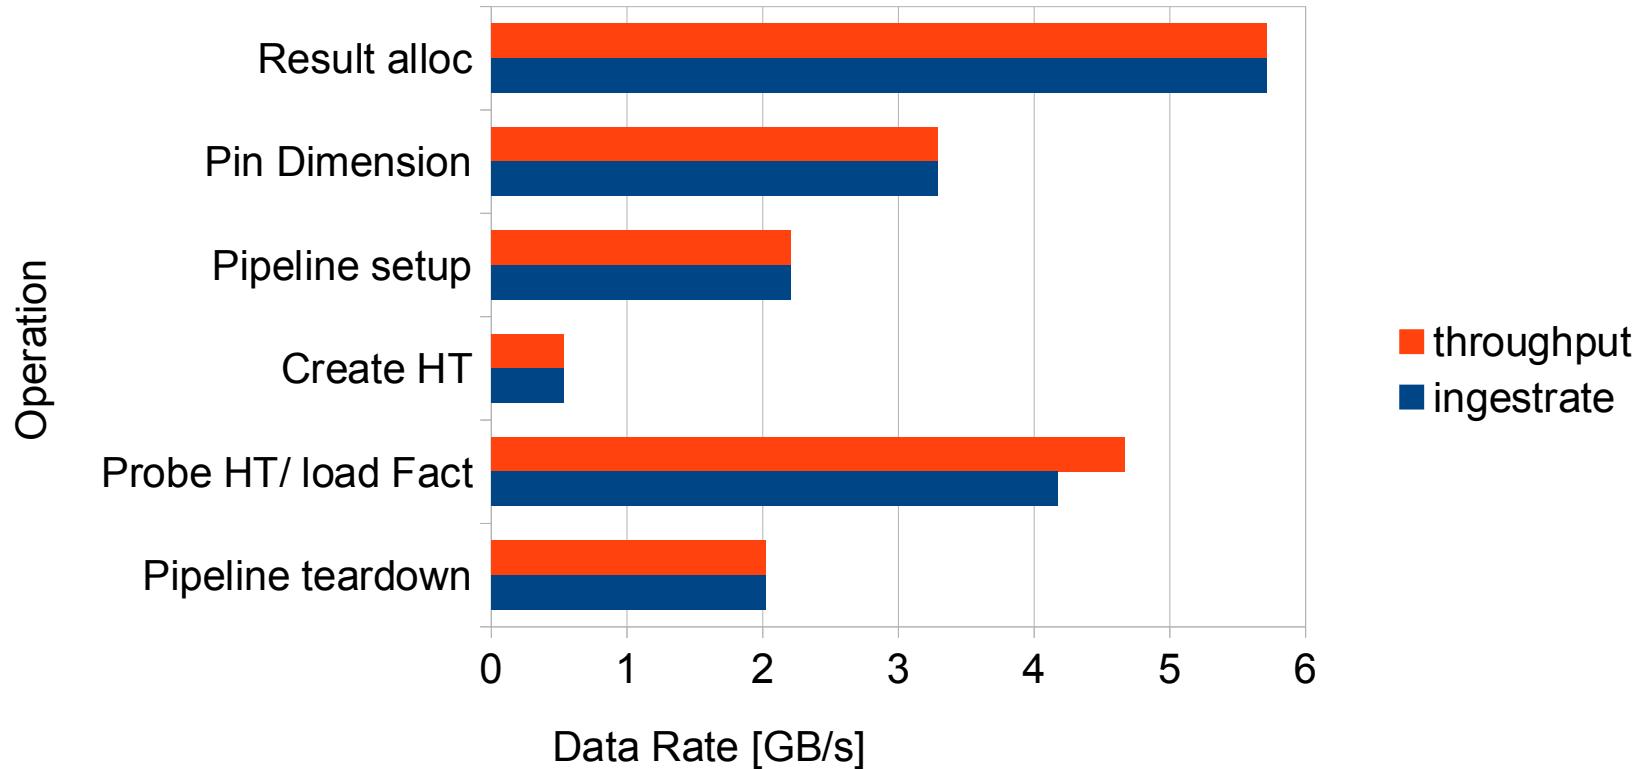


### Setup:

- Let GPFS do the striping & prefetching
- 2 CPU threads,
  - 1 for filling a pinned buffer from FS
  - 1 for controlling GPU execution
- GPU reads data from pinned buffer(s)

## Join with 3-stage pipeline from SSD

- Join lineorder & part using 2x 2MB buffers yields > 4GB/s overall throughput



- Virtually no performance difference to in-memory solution =)

## Agenda

- GPU search
  - Reminder: Porting CPU search
  - Back to the drawing board:
    - P-ary search
    - Experimental evaluation
    - Why it works
- Building a complete data warehouse runtime with GPU support
  - From a query to operators – what to accelerate?
  - What are the bottlenecks/limitations
- Maximizing data path efficiency
  - Extremely fast storage solution
  - Storage to host to device
- Putting it all together
  - Prototype demo

---

## Questions?